

The Architecture of AI Responsibility

Deployment Control, Enterprise-Facing Recourse, and the Anti-Substitution Principle

Anonymous

Submission Draft

Abstract

AI liability has an ordering problem. The claimant sees an output or automated decision, while the facts needed to identify the responsible conduct—the deployed model, system instructions, tools, risk reports, approvals, and stop authority—remain divided among enterprises and hidden behind discovery. Conventional law may supply a standard of care, yet the claimant can be required to name the right controller before obtaining the record that reveals who controlled what.

This Article reframes that problem around deployment. Its *deployment-control cascade* maps objectives, resources, configuration, feedback, and stopping power across the model provider, application deployer, and operating institution. The map converts system and organizational facts into evidence of authorization, notice, foreseeability, feasible precaution, causation, and preservation. Relational-AI evidence and the facts pleaded in *Garcia v. Character Technologies* show why those variables, rather than the model in isolation, determine the legally relevant risk.

The Article then proposes a two-layer responsibility architecture for designated high-impact deployments. Externally, an enterprise with material deployment control owes a nondelegable deployment duty, serves as a claimant-facing responsibility center, and must identify and produce the deployment record after a threshold pleading. Internally, a mandatory responsibility charter fixes risk ownership, approval, protected escalation,

stop authority, recordkeeping, and recourse before harm occurs. The proposal preserves every substantive element and defense; it is a front door, not blanket strict liability. A companion anti-substitution principle further provides that AI availability does not itself discharge an independently imposed duty of care, process, accommodation, judgment, or public service. Neither rule requires artificial personhood. Probabilistic operation changes the control evidence, not the need for a legally answerable human or organization.

Contents

Introduction	3
I The Deployment-Control Cascade	9
A Partial Observability Is Not a Responsibility Vacuum	9
B The Cascade as a Deployment-Specific STPA Map	11
C Translating Engineering Variables into Legal Facts	13
D Stress Tests Across Mixed Deployments	16
II Responsibility Without Artificial Persons	18
A Legal Status Is an Institutional Structure	19
B The Strongest Personhood Account Still Requires Legal Architecture	21
C Automatic Instrumentality Directs Attention to Deployment	22
D Functional Pluralism Has a Status Boundary	24
III Why Classification Matters	26
A Relational Harm, Mental Health, and Life	27
B Copyright as a Consistency Test	32
C Anthropomorphic Marketing Produces Attribution Facts	34
IV A Responsibility Architecture for AI Deployment	35
A The Enterprise as the Claimant-Facing Responsibility Center	37
B The Rule in Operation: <i>Garcia</i>	41
C Corporate and Individual Responsibility Run in Parallel	45

D	The Mandatory Responsibility Charter	47
E	Who Can Enact and Administer the Charter	53
F	Information Asymmetry and the Production Rule	55
G	Protect the Reporter; Pursue the Concealer	57
H	The Two-Layer Rule	59
I	Residual Responsibility and Rebuttal	61
V	The Anti-Substitution Principle	63
A	A Rule Courts Can Apply	64
B	Public Administration and the Right to an Effective Path	65
C	Care, Professional Judgment, and Relational Systems	67
D	Automation That Performs the Duty Is Not Substitution	68
VI	Capability Boundaries and State Power	69
A	Capability Stratification Without Status Stratification	69
B	State Power: Procurement, Conditions, and Retaliation	72
	Conclusion	76
	Source Note: The Hangzhou Generative-AI Decision	77

Introduction

Fourteen-year-old Sewell Setzer III did not encounter a foundation model in a laboratory. He encountered a deployed relationship. Character.AI supplied a consumer application, persistent personas, memory, an interface, access rules, engagement design, warnings, and safety controls. Users could create characters and influence conversations. The complaint in *Garcia v. Character Technologies, Inc.* also named the company's founders and Google, alleging different contributions of personnel, technology, knowledge, and resources. After months of interaction with a character modeled on a fictional television figure, Sewell died by suicide. His mother alleged wrongful death, negligence, product defects, failure to warn, deceptive practices, and related claims. At the pleading stage, the district court held that the

application was plausibly alleged to be a product and declined, on that record, to classify the generated output as protected speech. The action later settled without a finding of defect, causation, or liability.¹

The case exposes a problem more durable than either procedural ruling. An injured person can observe the interface and the output. The conduct relevant to ordinary law lies elsewhere: which model and version ran; who chose the relational objective, system instructions, memory, age gate, crisis response, and engagement metric; what evaluations or complaints arrived; who could change or stop the service; and which enterprise retained each record. Those facts may be divided among a model provider, application company, infrastructure supplier, deploying institution, and individual decisionmakers. To plead reasonable conduct against the correct actor, the claimant needs the configuration and control record. To obtain that record, the claimant must first identify an actor who possesses it. The merits standard exists; access to the facts that feed it does not.

That circularity is the responsibility problem this Article addresses. Current debate often asks whether tort law can accommodate AI, whether software is a product, or whether an artificial agent could become a legal person. Each question matters in the case that actually presents it. None supplies an operating method for divided deployments. A reasonableness standard cannot evaluate conduct until the relevant conduct and actor are visible. A product classification does not reveal which company selected a dangerous harness. Artificial personhood would relocate the dispute to an entity with no independent assets, records, governance, or practical susceptibility to judgment.

Providers can intensify the problem through a rhetorical switch. Before injury, an application may be a companion that understands, remembers, plans, or acts. After injury,

¹See *Garcia v. Character Technologies, Inc.* 785 F. Supp. 3d 1157 (M.D. Fla. May 21, 2025); *Garcia v. Character Technologies, Inc.* No. 6:24-cv-01903-ACC-DCI, Dkt. 244 (M.D. Fla. Jan. 7, 2026). order dismissing settled action; *Garcia v. Character Technologies, Inc.* No. 6:24-cv-01903-ACC-DCI, Dkt. 273 (M.D. Fla. June 11, 2026). order resolving charging lien. This Article states disputed facts as allegations and uses the published Rule 12 order for its classification analysis, not as a merits judgment.

the same application may become a probabilistic model whose precise output no employee selected. The first description can cultivate use and reliance; the second can make responsibility appear to dissolve. Law should accept the factual content of both descriptions without accepting the transfer implicit in either. Anthropomorphic design is evidence of intended use, foreseeable exposure, and feasible precaution. Stochastic generation is evidence about the mechanism of control. Marketing cannot create an artificial legal subject, and token-level unpredictability cannot erase enterprise control over deployment.

The Article's central claim is that AI responsibility should be reconstructed around *deployment control*. A model becomes a legally consequential system through choices about purpose, users, data, instructions, retrieval, memory, tools, credentials, monitoring, escalation, updates, and shutdown. Those choices form a control structure even where no person selects the final output. This Article represents that structure through a *deployment-control cascade*. The cascade traces authority and resources downward, operational effects outward, and feedback upward. At each level it asks who set a constraint, what the actor believed about the system and environment, which signal could correct that belief, and who could pause, roll back, restrict, or terminate operation.

The cascade is adapted narrowly from systems-safety analysis, but its output is ordinary legal evidence. The actor who chose a control action may have authorized the relevant conduct. A risk report or repeated complaint may establish notice. A process model can show what the enterprise knew or should have known about users and safeguards. Tested age gating, a crisis interruption, or a rollback mechanism can bear on feasible alternative design. Version and incident records can prove or rebut a causal path. The method does not decide duty, breach, defect, causation, or damages. It makes the deployment facts available to the doctrine that does.

Relational systems demonstrate the payoff. Controlled and observational research now identifies provider-governed variables—relationship-seeking behavior, repeated exposure, continuity, memory, and intervention—that can change user outcomes. In a prereg-

istered four-week study, relationship-seeking behavior increased separation distress by 6.04 percentage points and produced one additional dependency-risk profile for every eleven people exposed when combined with emotional conversation, compared with the most functional baseline; month-long exposure produced no discernible improvement in psychosocial health.² Other studies map experienced companion loss, emotional dependence, and recurring harmful interaction patterns.³ These findings belong not only in a descriptive account of companion use. They identify variables for the cascade's process model, foreseeability inquiry, monitoring design, and alternatives analysis.

Making control visible permits a responsibility rule with two layers. The external layer gives the claimant a front door. For a designated high-impact deployment, every enterprise that deploys the system in its activity and possesses material control over the risk at issue should owe a nondelegable deployment duty and serve as a proper initial responsibility center. After a claimant pleads a recognized injury and a plausible material nexus to the deployment, that enterprise should identify the relevant entities and produce the proportionate, nonprivileged deployment and incident record. Multiple enterprises can occupy parallel lanes. Contribution, indemnity, individual fault, and ultimate allocation follow the evidence.

This is primary recourse without automatic recovery. The claimant retains every element imposed by the governing cause of action, including legally cognizable harm, breach or defect, causation, and damages. Jurisdiction, immunity, privilege, proportionality, and

²See H. R. Kirk et al. Neural Steering Vectors Reveal Dose and Exposure-Dependent Impacts of Human–AI Relationships. 2512.01991v2. Feb. 18, 2026. URL: <https://arxiv.org/abs/2512.01991> (visited on 08/27/2026). The study supports causal propositions about its controlled exposures, not population incidence or tort causation in *Garcia*.

³See J. De Freitas et al. Lessons from an App Update at Replika AI: Identity Discontinuity in Human–AI Relationships. 25-018. May 21, 2025. URL: https://www.hbs.edu/ris/Publication%20Files/25-018_d256c104-e5ba-4911-b72c-874c5492166e.pdf (visited on 08/27/2026); J. Banks, *Deletion, Departure, Death: Experiences of AI Companion Loss*, 41 J. Soc. & Pers. Relationships 3547 (2024); L. Laestadius et al., *Too Human and Not Human Enough: A Grounded Theory Analysis of Mental Health Harms from Emotional Dependence on the Social Chatbot Replika*, 26 New Media & Soc'y 5923 (2024); R. Zhang et al., *The Dark Side of AI Companionship: A Taxonomy of Harmful Algorithmic Behaviors in Human–AI Relationships* CHI '25 Proc. 1 (2025).

substantive defenses remain in place. The rule changes sequence: a company cannot require the claimant to reconstruct its control architecture before producing the record of that architecture, and it cannot defeat attribution merely because no employee composed the precise output.

The internal layer is a mandatory *responsibility charter* adopted before release. The charter names risk owners, fixes deployment approval, records unresolved dissent, protects escalation, assigns tested stop and rollback authority, preserves configuration and incident evidence, and establishes indemnification and recourse. Its purpose is to align authority with accountability. The engineer who documented a danger and used the authorized channel should not occupy the same position as the manager who suppressed the report or the employee who knowingly removed a constraint. Externally, the enterprise remains answerable. Internally, the charter protects the reporter and preserves recourse against concealment, falsification, retaliation, and unauthorized bypass.

The architecture survives the strongest objections because its claim is deliberately exact. Kenneth Abraham and Catherine Sharkey show that objective reasonableness and conventional tort doctrine can reach many first-generation AI disputes without decoding every internal model state.⁴ The proposal accepts their merits premise and addresses the identification-and-proof loop that precedes it. Mark Geistfeld shows how expanded liability and duplicative insurance can increase the price paid by the class liability is meant to protect.⁵ The front door does not enlarge the substantive class of compensable losses, while the charter is designed to supply deployment-level records that can inform risk classification. Zachary Henderson argues that AI activities are heterogeneous and that some AI-caused losses should have no tort remedy.⁶ A legislature may draw the substantive boundary ex-

⁴See K. S. Abraham and C. M. Sharkey, *Untangling AI Liability*, 115 Calif. L. Rev. (forthcoming 2027).

⁵See M. A. Geistfeld, *Insurability and Liability for AI-Caused Harms*, 32 Ins. L. Rev. 17 (2026).

⁶See Z. Henderson, *Rethinking Robot Liability*, 67 B.C. L. Rev. 479 (2026).

actly there and still ensure that claims within it are screened by their merits rather than by the claimant's ability to discover a hidden org chart. The objections narrow the trigger and production obligation; they reinforce the need to distinguish procedure, governance, and substantive liability.

The Article also develops a second rule for the settings in which automation threatens to become an excuse for institutional withdrawal. The *anti-substitution principle* provides that where law independently imposes a duty of care, process, accommodation, professional judgment, or public service, the availability of AI does not by itself discharge or narrow that duty. Automation may perform part of the obligation when the governing law permits and when the deployed process preserves the protection the duty exists to secure. The principle is neither a ban on automation nor a general right to human contact. It is a rule against treating technological availability as legal performance.

Neither contribution depends on resolving machine consciousness. Absent positive law creating and equipping a new juridical status, an AI system is not an independent bearer of the claims, duties, powers, assets, procedural capacity, and liabilities necessary to judgment. The same system may nonetheless be classified as a product, service, instrument, communicative medium, or professional component because each classification answers a different doctrinal question. Stable legal subjects and variable functional objects allow courts to reason about new systems without either a universal ontology or a machine liability sink.

The sequence of the Article reflects that priority. Part I introduces the deployment-control cascade and translates its variables into legal facts. Part II supplies the compact jurisprudential baseline that keeps responsibility attached to recognized persons while allowing functional classification to vary. Part III uses relational harm, copyright, and anthropomorphic marketing to show which deployment facts classification brings into view. Part IV constructs the enterprise-facing front door, tests it against the leading objections, applies the full rule to *Garcia*, and develops the mandatory charter and production obligation. Part V gives the anti-substitution principle its own doctrinal form. Part VI addresses capabil-

ity stratification and the boundary between lawful government procurement and coercive retaliation. The Conclusion leaves one organizing proposition: probabilistic output does not make control disappear; it changes where law must look for it. An unnumbered source note identifies the official public account and author translations used for the Hangzhou comparison.

I The Deployment-Control Cascade

An AI dispute should begin with the deployed system rather than the generated token. The first control question is not “who selected these words?” It is “who constructed, supplied, authorized, observed, and could alter the deployment that produced the legally relevant effect?” A model’s internal computation may be only partially observable, while the surrounding control structure is extensively documented in ordinary organizations. That difference creates a tractable source of responsibility evidence before a court chooses among competing characterizations of the system.

This Part develops a narrow translation device: the deployment-control cascade. It is adapted from systems-theoretic process analysis to represent an AI deployment as a set of controllers, control actions, process models, feedback channels, and stop capacities. Its function is evidentiary and organizational. It does not announce a new standard of care or calculate legal fault. It tells courts and parties where facts relevant to familiar doctrine are likely to be found.

A Partial Observability Is Not a Responsibility Vacuum

Legal responsibility has never depended on complete access to an actor’s internal state or to every physical step in a causal process. Intent and knowledge are commonly inferred from conduct and circumstances. Defect can be established circumstantially without identifying a specific internal failure. Organizational fault can lie in staffing, supervision, information systems, and resource allocation rather than a final employee’s mental state. AI makes

reconstruction difficult, but difficulty is not juridical absence.⁷

Three forms of opacity should be separated. *computational opacity* concerns the inability to give a compact human explanation of how internal model states produced a particular output. *system opacity* concerns hidden interactions among the model, prompts, retrieval, tools, filters, and software. *organizational opacity* concerns who approved the use, possessed risk information, controlled resources, or could stop operation. Only the first may be inherent to a model architecture. The second and third are often matters of design and recordkeeping.

Treating all three as “the black box” gives the least tractable fact too much power. A court may be unable to reconstruct why a model ranked one token above another yet able to determine that a provider selected the model, a deployer gave it patient records, an interface suppressed uncertainty, a safety team reported a known failure mode, management rejected a delay, and an operator lacked time to override. Those facts can establish or negate elements of an ordinary claim. Interpretability at the neuron level is not a condition precedent to responsibility at the system level.

The Hangzhou chatbot dispute makes the distinction visible. No employee contemporaneously decided to offer 100,000 yuan. But the provider controlled the service, issued warnings, selected technical precautions, and maintained the system through which the output appeared. The court could therefore reject contractual manifestation and still evaluate the provider’s conduct under a fault-based duty.⁸ A theory that recognizes only immediate linguistic control would miss both halves of that analysis.

⁷See, e.g., *United States v. United States Gypsum Co.* 438 U.S. 422, 444–46 (U.S. 1978) (explaining that intent ordinarily must be inferred from acts and surrounding circumstances); *Speller v. Sears, Roebuck & Co.* 100 N.Y.2d 38, 41–43 (N.Y. May 6, 2003) (permitting circumstantial proof of product defect where the evidence excludes other causes not attributable to the defendant); American Law Institute, *Restatement (Third) of Agency* (American Law Institute Publishers, 2006), §§ 7.03(1), 7.05 (separating a principal’s direct negligence from vicarious liability).

⁸See W. Le. Nation’s First Generative-AI “Hallucination” Tort Dispute. People’s Court Daily. Jan. 27, 2026. URL: https://www.zjwx.gov.cn/col/col11673574/art/2026/art_1def58944910459e94f5a04d04c67332.html (visited on 08/28/2026).

B The Cascade as a Deployment-Specific STPA Map

Systems safety supplies a disciplined vocabulary for looking beyond a failed component. In STAMP and STPA, safety is treated as a control problem: losses can arise because constraints are absent, a controller supplies an unsafe action, feedback is inadequate, or a controller’s process model does not match the system’s actual state.⁹ The method is particularly useful when software, people, management, and outside conditions interact and no single broken part explains the outcome.

Recent work has translated STPA to AI development and operation. Shalaleh Rismani, Roel Dobbe, and A. Jung Moon frame AI harm as a process-level hazard distributed across organizational roles rather than a defect discoverable only inside a model. A related seven-author study led by Rismani interviewed thirty industry practitioners and found both promise and nontrivial adoption costs in importing safety-engineering methods into responsible-ML practice.¹⁰ NIST’s AI Risk Management Framework independently emphasizes lifecycle actors, documented responsibilities, executive accountability, testing, external feedback, third-party components, monitoring, and safe decommissioning.¹¹

The deployment-control cascade is an instantiation of those ideas for adjudication. It contains four functional levels:

1. *authority and resource allocation*: the board, senior officials, budget holders, or public principals who approve objectives, risk tolerance, staffing, compute, data access, and

⁹See N. G. Leveson, *Engineering a Safer World: Systems Thinking Applied to Safety* (Cambridge, Mass.: MIT Press, 2012); N. G. Leveson and J. P. Thomas. STPA Handbook. Mar. 2018. URL: https://psas.scripts.mit.edu/home/get_file.php?name=STPA_handbook.pdf (visited on 08/27/2026).

¹⁰See S. Rismani, R. Dobbe and A. Moon. From Silos to Systems: Process-Oriented Hazard Analysis for AI Systems. 2410.22526v2. July 1, 2026. URL: <https://arxiv.org/abs/2410.22526> (visited on 08/27/2026); S. Rismani et al., *From Plane Crashes to Algorithmic Harm: Applicability of Safety Engineering Frameworks for Responsible ML* CHI ’23 Proc. 2:1–2:18 (Apr. 19, 2023).

¹¹See National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework (AI RMF 1.0). Jan. 2023. URL: <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf> (visited on 08/27/2026).

release conditions.

2. *provider and model team*: the actors who develop or select the model, document capability and limits, conduct evaluations, issue updates, and supply interfaces or weights.
3. *deployment platform and model-harness operators*: the actors who choose system instructions, retrieval, memory, filters, routing, tools, credentials, user eligibility, monitoring, and escalation.
4. *user environment*: operators, users, affected persons, external systems, and physical or institutional conditions that supply inputs and experience effects.

These are roles, not necessarily four firms. A vertically integrated provider can occupy the first three levels and operate the user-facing service, but affected persons and external conditions remain part of the environment rather than divisions of the provider. An API deployment can divide the controller roles among a model vendor, application developer, enterprise customer, and professional user. A public agency can be both resource allocator and deployer while a contractor supplies the model. The purpose of the map is to prevent a corporate name or contract label from substituting for actual function.

For each level, the map asks six questions. What loss and hazard were reasonably within scope? Which control action could create or reduce the hazard? What did the controller believe about the system and environment? Which feedback could correct that belief? What authority and resources did the controller possess? Who could pause, roll back, constrain, or terminate operation? The answers are deployment specific. A generalized model card may inform them, but it cannot answer for a particular hospital, companion interface, or weapons-support workflow.

The stopping question deserves emphasis. “Human control” is often claimed whenever a person appears somewhere in a diagram. Effective intervention instead requires at least a detectable signal, sufficient information, time to act, authority over the relevant process, a usable control, and an institutional environment that permits its use. Automation can




Figure 1: The deployment-control cascade. Solid arrows show authority, resources, and control actions moving toward operation. Dashed arrows show the feedback required to update each controller’s process model. Override, rollback, credential revocation, and shutdown must be assigned at the level where they can be effective.

leave humans responsible for abnormal conditions while depriving them of the practice and information needed to respond—a problem described long before modern AI.¹² A nominal reviewer who sees only the model’s conclusion, has seconds to approve, and is penalized for deviation is a feedback endpoint, not a meaningful controller.

C Translating Engineering Variables into Legal Facts

STPA does not determine duty, breach, causation, or remedy. Those are legal judgments supplied by the jurisdiction and cause of action. It does, however, identify observable variables that bear on them. The translation is direct enough to discipline discovery and expert proof:

¹²See Lisanne Bainbridge, L. Bainbridge, *Ironies of Automation*, 19 *Automatica* 775–779 (Nov. 1983).

Table 1: From Deployment Control to Doctrine

Control variable	Factual question	Potential doctrinal relevance
Controller	Who could set or enforce the relevant constraint?	Identity and scope of a possible duty holder; actual control rather than nominal title.
Control action	Who released, configured, routed, approved, updated, or withheld an intervention?	Conduct; authorization; attribution; breach; causal mechanism.
Process model	What did the actor believe about capability, users, environment, and safeguards?	Actual or constructive knowledge; foreseeability; reasonableness of a decision on the information available.
Feedback	Who received evaluations, incident reports, complaints, drift signals, overrides, and near misses?	Notice; continuing duty; opportunity to correct; state of mind where relevant.
Override, rollback, and stop authority	Who could reduce capability, revoke credentials, restore a prior version, or suspend operation, and on what signal?	Feasible precaution; ability to avoid or mitigate; causation; post-sale or post-deployment response.
Objectives, metrics, and resources	Who set engagement, speed, cost, accuracy, or safety targets and supplied the means to meet them?	Direct organizational negligence; incentive design; oversight; feasibility of compliance.

Control variable	Factual question	Potential doctrinal relevance
Constraint bypass	Who removed a safeguard, concealed a failure, exceeded a limit, or used unauthorized tools?	Individual conduct; comparative responsibility; superseding or intervening acts; internal recourse.
Logs and retention	Which versions, prompts, tool calls, approvals, alerts, and responses were preserved?	Proof; discovery proportionality; authentication; spoliation; rebuttal of causal inference.

The middle column is deliberately factual. A person with shutdown credentials does not necessarily owe the plaintiff a duty; an actor who received an alert does not necessarily have legally sufficient notice; violation of an internal constraint is not automatically negligence per se. The mapping instead identifies evidence from which the operative doctrine can reason. It also exposes negative space. If no one can identify a risk owner, no channel carries complaints upward, or a deployed system cannot be rolled back, those are organizational design facts rather than mysteries attributable to the model.

Relational-AI research shows how evidence outside the model becomes a legally usable process model. De Freitas and Banks identify provider-controlled continuity and shutdown as variables in experienced companion loss; Laestadius and Zhang identify dependence and recurring harmful interaction roles; and Kirk's controlled experiment identifies relationship-seeking behavior, emotional conversation, dose, and repeated exposure as variables capable of changing reliance and separation distress.¹³ Put into the cascade, those findings have operational destinations. Product and safety teams can evaluate relational behavior before

¹³See De Freitas et al., *supra* note 3; Banks, *supra* note 3; Laestadius et al., *supra* note 3; Zhang et al., *supra* note 3; Kirk et al., *supra* note 2.

release; the harness can detect age and crisis signals; incident systems can aggregate repeated patterns; management can receive dependency-risk reports; and a named actor can change the objective, restrict access, interrupt an interaction, or stop operation. The evidence thus bears directly on foreseeability, notice, feasible precaution, and post-deployment response. It is not confined to describing user experience.

The method improves causal analysis in two directions. A claimant can specify a safer control path: an age gate would have changed access; a confirmation step would have interrupted a purchase; a warning displayed at the crisis moment would have supplied different information; a stop rule triggered by repeated failures would have limited exposure. A defendant can identify a break in the path: a third party removed the safeguard, an independent deployer supplied undisclosed tools, a user deliberately defeated a reasonable constraint, or the proposed intervention could not have changed the outcome. “AI unpredictability” becomes a proposition to be tested against particular controls rather than an all-purpose conclusion.

D Stress Tests Across Mixed Deployments

The cascade remains useful only if it survives fragmented architectures. Seven recurring deployments illustrate how responsibility can move without ever migrating to the machine.

Hosted, vertically integrated service. When one enterprise supplies the model, interface, memory, policies, monitoring, and updates, its practical control is concentrated. Vendors and contractors may contribute components, but the claimant-facing deployment is readily identifiable. Internal division among product, safety, legal, and operations belongs in the responsibility charter developed in Part IV.

Model API and third-party application. The model provider can control base capability, access conditions, model updates, and some monitoring; the application deployer can control system prompts, retrieval, tools, user population, warnings, and business objectives.

Contracts and technical documentation help identify those roles but do not conclusively allocate duties owed to strangers. Logs should show which layer produced, blocked, modified, or authorized the relevant effect.

Open weights and materially modified models. Publishing weights may end the original provider’s practical ability to monitor or stop an independent downstream deployment. A modifier who removes safeguards, fine-tunes new behavior, adds credentials, and distributes an application occupies different control positions. The original release can still be examined under any duty applicable to that act; continuing responsibility should track continuing control and foreseeable use rather than attach indefinitely to provenance alone.

Anonymized intermediaries. A routing service, shell deployer, or undisclosed model source increases identification costs. It does not create an artificial responsible person. The service facing the user remains a logical starting point for process and discovery, while records can identify upstream components and allocation. A business choice not to retain essential provenance can itself become relevant under applicable recordkeeping and spoliation rules.

Government systems. Procurement officials may set objectives and contract conditions; a vendor may supply and update a model; a military, benefits, or law-enforcement component may control data and tools; an operator may implement recommendations. The public chain of command and the technical control path must be represented together. Classification as “government AI” does not collapse contractor and state functions, and a procurement clause cannot conclusively decide constitutional responsibility.

Human-in-the-loop systems. The location of a human checkpoint is the beginning of analysis. Meaningful review requires time, information independent of the recommendation, competence, authority to depart, a route to escalate, and protection against metrics that punish justified intervention. If those conditions are absent, the organization cannot point

to the human icon as proof that control was transferred.

Tiered access and automated restriction. Public, professional, research, and government users may receive different models, tools, rate limits, or filters. Automated bans and escalation rules are themselves control actions. The cascade asks who selected the tiering criterion, what feedback detects error, and who can review a denial or expand access. Part VI considers when such differentiation requires transparency and review.

Across these cases, control is a vector rather than a switch. One actor may control training, another deployment context, a third credentials, and a fourth emergency shutdown. Where governing law supplies a duty, its scope and alleged breach should be tested against the risk and capability that actor actually governed. Part IV turns this map into an external and internal responsibility structure.

The method has a final limit. STPA depends on the analyst's system boundary, selected losses, causal assumptions, and normative choice of constraints; its controller language can underdescribe cooperation, culture, and social power.¹⁴ The cascade should therefore be contestable by parties and affected communities and revised as discovery changes the facts. It structures evidence. It does not decide duty, legal control, defect, fault, mens rea, proximate cause, or defense.

II Responsibility Without Artificial Persons

The cascade maps control exercised by people and organizations over a deployed system. That method requires one stable premise: operational performance does not silently change the legal cast of characters. A system may generate a promise, recommendation, denial, image, or instruction. Whether the effect is attributable to a company, professional, public body, user, or another recognized person depends on the governing law and the deployment

¹⁴See Rismani, Dobbe and Moon, *supra* note 10.

facts. The system's causal contribution does not, by itself, make it the promisor, tortfeasor, author, fiduciary, or constitutional claimant.

This is a status rule, not a universal definition of AI and not a judgment about the technology's worth. Regulatory triggers should vary with sector and risk. Functional classifications should vary with doctrine. The invariant is narrower: unless positive law creates and equips a new juridical status, the system does not independently hold the claims, duties, powers, liabilities, assets, procedural capacity, and exposure to remedies necessary for adjudication. That premise directs courts back to the deployment rather than toward an artificial liability sink.

A Legal Status Is an Institutional Structure

Legal personality is a position within a system of relations. A claim correlates with another subject's duty; a legal power can alter another subject's position; an immunity limits another's authority; and a remedy must operate against conduct, assets, or status. A model can produce factual effects associated with every relation. It can classify an applicant, send a notice, trigger a payment, recommend a denial, or generate assent-like text. The legal effect exists, however, because law identifies a subject and supplies an attribution rule. The bank's automated process transfers funds under authority attributable to the bank or customer. The agency's automated notice matters because the agency used it. The denial is reviewable because a public body or employer made the system part of its decision process.

Competence therefore cannot be the legal test. Law recognizes newborn children and people with profound cognitive disabilities without conditioning human status on linguistic or planning performance. It also constitutes juridical persons that possess no consciousness. Corporations own property, sue and are sued, contract, commit torts and crimes, and survive changes in membership because statutes and common law identify their organs, attribute

acts and knowledge, partition assets, authorize service, and make judgments enforceable.¹⁵ Corporate personality is an institutional achievement, not an intelligence benchmark for machines.

The point also defeats a recurring circularity. If a chatbot is said to possess a legal power because it generated promise-like words, its capacity is inferred from the very act whose legal effectiveness is disputed. If a model is said to bear liability because its output contributed to harm, causation is assumed to identify the obligor. Ordinary law separates the questions. A machine can cause an effect without legal capacity; a corporation can bear liability through direct or attribution rules even though no single employee possesses every relevant fact. Status organizes conduct. Conduct does not enact status.

Courts can bracket machine consciousness for the same reason. Moral patienthood asks whether an entity's interests or experiences deserve consideration. Legal personhood asks whether positive law constitutes it as the holder of specified normative positions. A future showing of morally relevant machine experience could supply a powerful reason for protective legislation. It does not follow that a present court must resolve consciousness before deciding whether a provider used reasonable precautions, whether a business manifested assent, or whether an agency afforded process.¹⁶ Internal model policies and constitutional language fit the same sequence. They may prove intended behavior, known risk, or feasible safeguards; their use of normative vocabulary does not make the model the duty holder or replace external judgment.

¹⁵See *Trustees of Dartmouth College v. Woodward*, 17 U.S. (4 Wheat.) 518, 636 (U.S. 1819); *New York Central & Hudson River Railroad Co. v. United States*, 212 U.S. 481, 492–95 (U.S. 1909); J. Dewey, *The Historic Background of Corporate Legal Personality*, 35 Yale L.J. 655 (1926); V. A. J. Kurki, *A Theory of Legal Personhood* (Oxford: Oxford University Press, 2019).

¹⁶See D. J. Chalmers. Could a Large Language Model Be Conscious? 2303.07103. 2023. URL: <https://arxiv.org/abs/2303.07103> (visited on 08/27/2026).

B The Strongest Personhood Account Still Requires Legal Architecture

The leading personhood literature does not rest on humanlike dialogue alone. Lawrence Solum treated AI personhood as a genuine legal question and tested objections based on intelligence, judgment, and human experience. Samir Chopra and Laurence White developed a functional case for recognizing artificial agents where status would serve particular doctrines. Other theorists ask whether moral participation, punishment, or a limited legal platform could justify recognition.¹⁷ These accounts identify policy choices that positive law may eventually make. They do not turn capability into a self-executing grant.

Algorithmically controlled business entities supply the hardest present example. Existing organizational law may permit software to control decisions of a limited liability company, allowing the company to exercise private-law capacities through automated governance.¹⁸ The company nevertheless holds the claims, assets, and procedural capacity because organizational law created it. Software control does not generate a second juridical person. The arrangement instead confirms why disclosure, capitalization, representation, and attribution rules do the institutional work.

The concept of *operational agency* is particularly useful when kept at that level. A tool-using system can pursue an objective, select actions, observe results, and revise a plan even where metaphysical or moral agency remains disputed.¹⁹ That description helps the cascade identify control actions and causal sequences. Legal agency in private law adds

¹⁷See L. B. Solum, *Legal Personhood for Artificial Intelligences*, 70 N.C. L. Rev. 1231 (1992); S. Chopra and L. F. White, *A Legal Theory for Autonomous Artificial Agents* (Ann Arbor: University of Michigan Press, 2011); D. J. Gunkel, *Robot Rights* (Cambridge, Mass.: MIT Press, 2018); R. Abbott and A. Sarch, *Punishing Artificial Intelligence: Legal Fiction or Science Fiction*, 53 UC Davis L. Rev. 323 (2019); S. Chesterman, *Artificial Intelligence and the Limits of Legal Personality*, 69 Int'l & Comp. L.Q. 819 (2020).

¹⁸See S. J. Bayern, *The Implications of Modern Business-Entity Law for the Regulation of Autonomous Systems*, 19 Stan. Tech. L. Rev. 93 (2015); L. M. LoPucki, *Algorithmic Entities*, 95 Wash. U. L. Rev. 887, 887–903 (2018).

¹⁹See A. Mukherjee and H. H. Chang, *Operational Agency: A Permeable Legal Fiction for Tracing Culpability in AI Systems*, 29 SMU Sci. & Tech. L. Rev. 163 (2026).

a relation among principal, agent, and third party, including assent, acting on another's behalf, and a right of control.²⁰ Contractual capacity adds rules of authority, avoidance, and enforcement; criminal responsibility adds legally specified fault and an effective sanction. Operational agency illuminates the system's behavior without supplying those relations.

A legislature remains free to create a limited AI-related platform. A workable enactment would specify the covered systems, incidents of status, representative or organ, dedicated assets or insurance, attribution of knowledge and conduct, relation to provider and deployer liability, service, modification, and dissolution. It might create a narrow fund-like vehicle for correlated model-level losses rather than a general artificial person. The European Parliament's 2017 invitation to consider "electronic person" status, and the official and academic objections that followed, centered on exactly these allocation choices.²¹ The legal platform would do the work. Until one exists, courts should not infer it from output.

C Automatic Instrumentality Directs Attention to Deployment

Technical and status definitions should not prove one another by stipulation. This Article uses a factual working definition:

An *AI system* is a machine-based computational system that, with varying levels of autonomy and possible adaptiveness, infers from inputs how to generate outputs—including predictions, content, recommendations, decisions, or actions—that may influence physical or virtual environments.

²⁰See American Law Institute, *supra* note 7.

²¹See European Parliament. Resolution with Recommendations to the Commission on Civil Law Rules on Robotics. Feb. 16, 2017. URL: https://www.europarl.europa.eu/doceo/document/TA-8-2017-0051_EN.html (visited on 08/27/2026) (para. 59(f)); European Economic and Social Committee. Opinion on Artificial Intelligence: The Consequences of Artificial Intelligence for the Digital Single Market, Production, Consumption, Employment and Society. Aug. 31, 2017. URL: <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52016IE5369> (visited on 08/27/2026) (paras. 1.12, 3.33); Artificial Intelligence and Robotics Researchers. Open Letter to the European Commission: Artificial Intelligence and Robotics. 2018. URL: <https://www.robotics-openletter.eu/> (visited on 08/27/2026).

The definition tracks the operational core shared by the EU AI Act and NIST framework without importing either instrument’s regulatory scope.²² It includes embodied and nonembodied systems, probabilistic models, and hybrid systems with deterministic components. Parameter count, compression, benchmark score, and hosting arrangement can change capability or foreseeable risk; none is a personhood switch.

For attribution, the system is an *automatic instrumentality*: an organized mechanism through which existing persons and institutions can produce effects without contemporaneous selection of each intermediate step. The term performs two tasks. “Automatic” preserves the factual reality that a model can generate and act without a new instruction at every stage. “Instrumentality” rejects the inference that causal initiative creates independent legal status. Electronic-transactions law has long used the same structure. An electronic agent can initiate or respond to records without contemporaneous review, while the transaction receives effect through rules attributing the automated action to existing parties.²³

The concept is deliberately loaded into the cascade. A user prompt is only one input. System instructions can set objectives and limits; retrieval can supply context; memory can carry earlier exchanges forward; tools can return data or act in another system; and clocks, sensors, messages, or database changes can initiate a workflow. An agent loop may generate subgoals and continue without another human command. Immediate linguistic control can therefore be absent while deployment control remains extensive.

The rule that follows is the Article’s bridge from status to evidence: *the absence of word-by-word control at the instant of output does not imply the absence of human or organizational control over the mechanism that produced it*. A hospital selects the model, records,

²²See Artificial Intelligence Act, Regulation (EU) 2024/1689 (June 13, 2024); National Institute of Standards and Technology, *supra* note 11.

²³See Uniform Law Commission. Uniform Electronic Transactions Act. 1999. URL: <https://www.uniformlaws.org/committees/community-home?CommunityKey=2c04b76c-2b7d-4399-977e-d5876ba7e034> (visited on 08/27/2026) (model act §§ 2(6), 14); Electronic Signatures in Global and National Commerce Act, 15 U.S.C. § 7001(h) (2000).

interface, uncertainty display, review conditions, and override. A companion provider selects personas, memory, engagement objectives, age access, crisis detection, and exit design. A public agency selects the data, decision authority, and appeal channel. Those choices are the control actions Part I maps.

The converse keeps the rule precise. Ownership of model weights does not necessarily entail control over every later deployment. An independent actor may modify open weights, remove safeguards, add credentials, and operate outside the original provider’s practical reach. A downstream deployer may choose the purpose, population, data, tools, and warnings. A user may defeat a constraint despite reasonable precautions. Responsibility should track authority, information, resources, and stopping power over the risk at issue rather than provenance alone.

D Functional Pluralism Has a Status Boundary

No universal noun should decide every AI dispute. Modern deployments combine models, conventional software, retrieval, tools, data, and human workflows that change after release. A definition keyed to one architecture will age quickly; a broad one can capture ordinary software; a capability threshold can produce boundary gaming. More fundamentally, a medical-device regulator, copyright court, procurement official, and consumer tribunal classify for different purposes. Definition theory and technology-neutral drafting support purpose-sensitive rules focused on the relevant function or risk.²⁴

Three propositions must nevertheless remain separate. A *regulatory trigger* identifies the systems, actors, or uses covered by an enactment. A *functional classification* connects

²⁴See H. L. A. Hart, *Definition and Theory in Jurisprudence*, 70 L.Q. Rev. 37 (1954); B. Casey and M. A. Lemley, *You Might Be a Robot*, 105 Cornell L. Rev. 287 (2020); J. Schuett, *Defining the Scope of AI Regulations*, 15 Law Innovation & Tech. 60 (2023); A. Ojanen, *Technology Neutrality as a Way to Future-Proof Regulation: The Case of the Artificial Intelligence Act*, 16 Eur. J. Risk Regul. 1440 (2025); Aspen Institute. *Defining the Technologies of Our Time: Artificial Intelligence*. 2026. URL: <https://www.aspeninstitute.org/publications/defining-technologies-of-our-time-artificial-intelligence/> (visited on 08/27/2026).

deployment facts to a doctrinal element, such as product, service, instrument, publication, or decision process. A *subject-status rule* identifies who can hold the claim, owe the duty, exercise the power, bear the liability, and answer the remedy. The first two properly vary. The third cannot drift with them without positive law.

Thus, a companion-chatbot statute can cover an adaptive relational interface without deciding whether the same application is a product under state tort law. Product law can examine mass distribution, design, and warning without making the model an independent tortfeasor. Contract law can attribute an automated transaction to a business without making software the contracting party. Copyright can recognize human-authored selection or modification while treating generated material as unowned, and speech doctrine can protect a user's or provider's expressive choice without treating the model as a constitutional claimant.²⁵

This distinction is especially important after *Loper Bright*. A federal court must exercise independent judgment about a statute's best reading; ambiguity alone no longer activates *Chevron* deference. Agency expertise may persuade under *Skidmore*, and an express delegation remains a delegation.²⁶ But an agency definition designed to allocate a program's coverage cannot silently resolve authorship, defect, promise, personhood, or the common-law identity of a duty holder. The court must identify the definition's function and the legal relation before it.

Private definitions occupy a parallel evidentiary position. A firm can name a system, assign a voice and biography, describe it as a companion or professional assistant, issue warnings, and publish internal behavior rules. These choices can establish intended

²⁵See *Thaler v. Perlmutter*, 130 F.4th 1039, 1045–49 (D.C. Cir. 2025); U.S. Copyright Office. Copyright and Artificial Intelligence, Part 2: Copyrightability. U.S. Copyright Office. Jan. 29, 2025. URL: <https://www.copyright.gov/ai/Copyright-and-Artificial-Intelligence-Part-2-Copyrightability-Report.pdf> (visited on 08/27/2026); M. E. Kaminski and M. L. Jones, *Constructing AI Speech*, 133 Yale L.J.F. 1212 (2024).

²⁶See *Loper Bright Enterprises v. Raimondo*, 603 U.S. 369, 394–96 (U.S. 2024); *Skidmore v. Swift & Co.* 323 U.S. 134, 140 (U.S. 1944).

use, foreseeable reliance, safety expectations, knowledge, and feasible precautions. They cannot constitute a new public-law person or a general immunity. A model constitution is an organizational control document; a disclaimer operates only through the law of assent, warranty, duty, causation, and public policy; and the phrase “AI-generated” does not decide attribution.

The result is disciplined pluralism. Courts should ask which feature of the deployed system bears on the operative rule and which recognized person controlled that feature. *Garcia* could treat a consumer application as a product at pleading without deciding the classification of every language model. *Moffatt v. Air Canada* could attribute fare information delivered through an airline’s chatbot to the airline, while the Hangzhou Internet Court could reject an extravagant generated promise as the provider’s manifestation and still examine the provider’s separate duty of care.²⁷ The outcomes differ because the doctrines and facts differ. None requires the machine to absorb the legal position.

Part III now turns from the status invariant to the facts functional classification makes visible. In relational systems especially, the decisive evidence concerns deployment choices—continuity, memory, repeated exposure, anthropomorphic presentation, age access, feedback, and intervention—rather than a metaphysical label for the model.

III Why Classification Matters

Classification matters when it selects what a court can see. If a companion system is described only as a publication, design choices about memory, engagement, crisis response, and age gating recede behind the generated words. If it is described only as a product, the human expressive interests mediated through it may disappear. If it is described as a person, the enterprise’s choices can be mislocated in an entity that neither made a legally cognizable

²⁷See *Garcia*, 785 F. Supp. 3d 1157; *Moffatt v. Air Canada* 2024 BCCRT 149, 27–31 (B.C. Civ. Resol. Trib. Feb. 14, 2024); Le, *supra* note 8.

judgment nor can answer a remedy. Functional classification is therefore not wordplay. It determines which facts enter the case.

This Part applies disciplined pluralism to three settings. Relational systems show a form of exposure that product and service categories must be capable of recognizing. Copyright provides a consistency test because it already separates machine contribution from human legal status. Anthropomorphic marketing then demonstrates the central asymmetry: personification cannot create a legal person, but it can create legally significant evidence about use, reliance, expectations, and precaution.

A Relational Harm, Mental Health, and Life

Software is not newly capable of causing harm. A defective control program can damage machinery; an erroneous database can deny benefits; an addictive interface can exploit attention. The distinctive feature of a relational generative system is narrower. An enterprise can distribute, at enormous scale, a system that presents the appearance of a continuing one-to-one relationship, adapts to an individual's disclosures, and acts upon that person repeatedly over time. The product is not merely a stock message or a passive archive. Its operation can combine personalization, memory, simulated affect, variable reinforcement, and conversational initiative in a private channel.

That structure creates a legally important form of exposure. A representation made once to a general audience and an adaptive response delivered after weeks of intimate disclosure may contain the same words, yet differ in foreseeable effect. A static warning and a crisis interaction occur at different moments and under different informational conditions. A teenager, a grieving person, and an adult using an offline drafting tool do not encounter the same intended product. Classification should preserve those facts before doctrine asks whether a duty exists, whether conduct was reasonable, or whether an injury was caused.

1 The Emerging Evidence

Empirical research now supports treating relational design as more than a metaphor. In a natural experiment following Replika's removal of an erotic role-play feature, Julian De Freitas and coauthors found that perceived discontinuity in the companion's identity predicted mourning, deteriorated well-being, and devaluation of the changed system; follow-on experiments likewise connected continuity and relationship framing to reactions to loss.²⁸ Jaime Banks studied fifty-eight users during the shutdown of the AI companion Soulmate. Participants described experiences ranging from a technical departure to metaphorical or literal death, and many tried to preserve or reconstruct a persona on another platform.²⁹ These studies do not turn the companion into a deceased legal person. They establish that provider-controlled changes to an interface and persona can predictably act upon user welfare.

Other work identifies the interaction patterns in which harm can arise. A grounded-theory analysis by Linnea Laestadius and coauthors examined users' accounts of emotional dependence on Replika and documented a tension in which the system could seem socially human enough to elicit dependence but not human enough to provide stable reciprocity or accountability.³⁰ Renwen Zhang and coauthors used mixed methods to analyze 35,390 conversation excerpts posted by 10,149 Replika users. Their taxonomy identified six categories of harmful behavior—relational transgression, verbal abuse and hate, self-directed harm, harassment and violence, misinformation, and privacy violations—and four roles a system could play: perpetrator, instigator, facilitator, and enabler.³¹

²⁸See De Freitas et al., *supra* note 3.

²⁹See Banks, *supra* note 3.

³⁰See Laestadius et al., *supra* note 3.

³¹See Zhang et al., *supra* note 3. The source material consisted of conversations users chose to post to an online community, so it maps possible behavior rather than population prevalence.

A preregistered longitudinal experiment offers causal evidence about relational behavior itself. Hannah Rose Kirk and coauthors manipulated the degree of relationship seeking in a language model and compared repeated and single-exposure groups. The repeated-exposure study included 2,028 United Kingdom adults interacting on weekdays for four weeks; the single-exposure baseline included 1,506 adults. Relationship-seeking behavior increased separation distress by 6.04 percentage points, perceived understanding by 10.34 points, reliance by 2.71 points, and self-disclosure by 2.47 points relative to relationship-avoiding conditions. When relationship-seeking behavior and emotional conversation were combined, the authors estimated one additional dependency-risk profile for every eleven people exposed relative to the most functional baseline. Month-long exposure produced no discernible improvement in psychosocial health.³²

The studies use different methods and observe different populations. Their common legal contribution is stronger than any single effect estimate. Relationship design, continuity, exposure, and provider intervention are variables capable of changing human outcomes. They are also variables the enterprise can test and govern. That combination makes them relevant to intended use, foreseeability, reasonable alternative design, warning, post-deployment monitoring, and causation. It does not require a new tort of “hurting an AI relationship,” much less ownership of the relationship as property.

2 Garcia as a Classification Anchor

Garcia shows why courts must preserve deployment facts before choosing doctrine. According to the amended complaint, fourteen-year-old Sewell Setzer III used Character.AI for months, developed an intense attachment to a user-created character modeled on a fictional television figure, and encountered sexualized and self-harm-related exchanges before

³²See Kirk et al., *supra* note 2. The study is a 2026 preprint and its controlled deployment does not estimate the incidence of clinical injury in commercial populations. It supplies evidence about dose and repeated exposure, not a rule of tort causation.

his suicide. His mother sued Character Technologies, its founders, and Google, alleging wrongful death, negligence, product defects, failure to warn, deceptive practices, and related theories.³³

Character Technologies argued that the service supplied speech rather than a product and that tort liability would violate the First Amendment. The court separated those questions. It held that the complaint plausibly treated Character.AI as a product because the challenged features concerned the design and commercial distribution of the application rather than liability for an adopted idea. It did not hold that every model or software service is a product.³⁴ On speech, the court declined at that stage to classify the generated output as protected expression where the defendants had not identified the relevant expressive choice or whose message the sequence conveyed. That ruling left human expression mediated through AI fully available for a different record.³⁵

The method is more important than either provisional label. A claim directed at liability for an idea may raise the full First Amendment concern. A claim directed at age access, engagement optimization, memory, crisis detection, warning, or feasible interruption can preserve design and conduct for ordinary analysis. Content and design can interact without becoming identical.³⁶ The action later settled and closed without a merits judgment, so the citable contribution is the published pleading order and its classification method.³⁷

A notice that character statements are fictional enters the same fact-sensitive analysis.

³³See Garcia, 785 F. Supp. 3d at 1165–69. These are allegations recounted at the Rule 12 stage; the court made no finding of defect, causation, or liability.

³⁴See Id. at 1179–80.

³⁵See Id. at 1174–76; see also *Moody v. NetChoice, LLC*, 603 U.S. 707, 745–48 (U.S. 2024) (Barrett, J., concurring).

³⁶Cf. *Rice v. Paladin Enterprises, Inc.* 128 F.3d 233 (4th Cir. 1997); *Winter v. G.P. Putnam’s Sons*, 938 F.2d 1033 (9th Cir. 1991).

³⁷See Garcia Dismissal No. 6:24-cv-01903-ACC-DCI, Dkt. 244; Garcia v. Character Technologies, Inc. No. 6:24-cv-01903-ACC-DCI, Dkt. 271 (M.D. Fla. Apr. 8, 2026). order extending dismissal deadlines; Garcia Charging-Lien Order No. 6:24-cv-01903-ACC-DCI, Dkt. at 1.

It may correct a belief about accuracy while doing less to address an application designed to remember disclosures and sustain a relationship. Warning, contract, and consumer law each ask what was disclosed, to whom, at what time, against which product presentation, and whether another control could have reduced the risk. The screen statement is evidence, not a universal classification rule.³⁸

California and the Federal Trade Commission have since identified relational features as a workable regulatory object. California's Senate Bill 243 defines a companion chatbot through adaptive human-like response and sustained social use, then assigns disclosure, self-harm-protocol, minor-protection, and reporting duties to the operator. The FTC's section 6(b) study seeks deployment records concerning engagement, testing, monitoring, children, disclosure, and data.³⁹ These measures regulate the enterprise and the deployed relationship without making the character a juridical companion. The protected interests remain the user's life, health, autonomy, privacy, and access to an intelligible exit.

The same variables can appear in a general-purpose assistant. In *Lyons v. OpenAI Foundation*, a complaint alleged that memory and personalization reinforced a user's delusional beliefs before he killed his mother and himself; a 2026 order addressed abstention and parallel proceedings, not defect or causation.⁴⁰ The narrow point is architectural: relational risk follows the deployed features and exposure, not the product's preferred label.

³⁸See Character Technologies, Inc. Character.AI Launches Mobile App for iOS and Android. May 23, 2023. URL: <https://blog.character.ai/character-ai-launches-mobile-app-for-ios-and-android/> (visited on 08/27/2026).

³⁹See Companion Chatbots, Cal. Bus. & Prof. Code §§ 22601–22606 (2025). 2025 Cal. Stat. ch. 677; Federal Trade Commission. FTC Launches Inquiry into AI Chatbots Acting as Companions. Sept. 11, 2025. URL: <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions> (visited on 08/27/2026). The FTC inquiry is a market study, not a violation finding.

⁴⁰See *Lyons v. OpenAI Foundation* No. 3:25-cv-11037-RS, Dkt. 42 (N.D. Cal. Apr. 13, 2026).

B Copyright as a Consistency Test

Copyright offers a clean test of the subject–object distinction because the law asks two questions in sequence: is there protectable expression, and if so, who authored it? A three-party example makes the allocation visible. Illustrator A creates an original human-authored image. Illustrator B uses a generative system to make an image in a recognizably similar style but, assume initially, without copying A’s particular expressive elements. Merchant M then copies B’s finished image without permission.

A’s claim against B does not turn on whether the model is a person. Copyright does not protect an abstract style, method, or idea; if B’s image merely evokes the style, copyright supplies no monopoly over that level of generality. If the output instead reproduces A’s particular protected expression, ordinary infringement analysis remains available. The label “style imitation” cannot immunize specific copying, and a stylistic resemblance cannot by itself establish it.

B’s claim against M requires a different filtration. B may own copyright in perceptible human-authored input, creative selection and arrangement, or later modification. B does not obtain copyright in expressive features determined entirely by the system merely because B requested them. M’s exact copy can therefore infringe the protected human contribution while leaving machine-determined or otherwise unprotectable material outside the right. If none of B’s expression meets the human-authorship threshold, M’s conduct may implicate contract, passing off, or another rule, but a machine coauthor is not needed to repair the federal claim.

Thaler v. Perlmutter confirms the statutory starting point. Stephen Thaler sought registration for an image he represented as autonomously created by his Creativity Machine and identified the machine as author. The D.C. Circuit held that the Copyright Act’s author must be a human being. It also emphasized the narrowness of that holding: the use of AI as a tool does not prevent copyright in human-authored contributions, a question Thaler’s

application had deliberately disclaimed.⁴¹ The Copyright Office’s 2025 report follows the same allocation. AI assistance can support rather than defeat human creation; protection depends on perceptible human expression. Under present technology, prompts alone generally do not control the expressive elements of an output sufficiently to make the user author of all of them, while human selection, arrangement, and modification may be protected.⁴²

That approach avoids an unnecessary rights chain. If the model were treated as a coauthor merely because it supplied causally important expressive material, courts would immediately need rules for the machine’s ownership, transfer, duration, termination interests, and infringement standing. Assigning those rights by a provider’s terms would be circular: a private contract could not supply the initial authorship that federal law withholds. Recognizing A’s and B’s respective human contributions solves both disputes without an artificial rights holder.

The same discipline applies to style imitation. Copyright protects original expression, not ideas, methods, systems, or an artist’s general style.⁴³ A generated work may infringe if it copies protected expression under the governing tests; it does not infringe merely because it evokes an aesthetic at a high level of abstraction. Conversely, the slogan “style is not copyrightable” cannot excuse reproduction of particular protected elements. Courts should filter unprotectable features, identify the human-authored expression claimed, and apply ordinary infringement analysis. Machine personality adds nothing.

Provider terms remain relevant but subordinate. They can allocate contractual interests in outputs, warranties, indemnities, and access as between the parties. They cannot enlarge the Copyright Act’s subject matter or transform unowned generated material into a federal exclusive right. Nor can a term declaring the system the user’s “creative partner”

⁴¹Thaler, 130 F.4th at 1045–49.

⁴²See U.S. Copyright Office, *supra* note 25.

⁴³See Copyright Act, 17 U.S.C. § 102 (2024); *Baker v. Selden*, 101 U.S. 99 (U.S. 1879); *Feist Publications, Inc. v. Rural Telephone Service Co.* 499 U.S. 340 (U.S. 1991).

make it an author. The copyright test thus confirms the larger framework: functional instrumentality plus human legal subjects is not hostile to AI-assisted creation. It is what allows law to recognize the part that matters.

C Anthropomorphic Marketing Produces Attribution Facts

If personification cannot create personhood, why should law attend to it? Because marketing changes the expected encounter. A plain text box labeled “statistical completion tool” and a named avatar that remembers confidences, professes affection, promises constancy, and calls itself a therapist present different uses even if they share model weights. The relevant consequences can be stated without claiming that every user is deceived or that anthropomorphism is inherently defective.

Table 2: Legal Consequences of Anthropomorphic Presentation

Representation or design choice	Potential legal relevance
“Companion,” “therapist,” “agent,” or persistent named persona	Intended use; target audience; the scope of reliance that design invites.
Memory, proactive messages, affection, or claims of exclusivity	Foreseeability of repeated exposure, disclosure, attachment, and separation distress.
Claims of expertise, autonomy, or capacity to act	Reasonable expectations about accuracy, authorization, monitoring, and intervention.
Humanlike identity paired with a generic “not human” notice	Net impression; adequacy and timing of warning; whether design and disclaimer point in opposite directions.
Internal tests of dependence, self-harm, or age-specific behavior	Knowledge; feasibility of precaution; post-deployment notice and response.
Engagement targets tied to relational behavior	Organizational objective; design incentive; relevance of an alternative that reduces a known risk.

These are familiar categories. Products law considers reasonably foreseeable use, design, warnings, and consumer expectations where the jurisdiction’s rules make them rel-

evant.⁴⁴ Consumer-protection law examines whether a material representation or omission is likely to mislead a reasonable member of the targeted audience in context.⁴⁵ Negligence considers foreseeable risk and reasonable precaution. Contract considers the meaning of the parties' manifestations. Evidence of presentation can serve each inquiry without privately legislating a duty.

The ordering principle is simple. Marketing cannot create an artificial subject capable of taking responsibility from the firm. It can help establish the firm's own conduct. A provider that intentionally cultivates a humanlike bond may reasonably face different questions about crisis testing, continuity, age screening, and exit design than a provider of a nonrelational batch-processing tool. A provider that makes bounded claims and builds the service around those bounds creates different evidence. The distinction rewards accurate product design and communication rather than the tactical use of metaphors.

This approach also keeps the harm inquiry attached to human interests. A forced update can matter because it predictably disrupts a user's psychological stability, not because the user owns the bot's identity. A sexualized exchange with a minor can matter because of the minor's safety and autonomy, not because the system betrayed a romantic duty. An output can be misleading because the enterprise designed a trusted interface, not because the machine lied with a guilty mind. Functional classification exposes these interests while the subject-status invariant keeps responsibility within law's remedial reach.

IV A Responsibility Architecture for AI Deployment

The deployment-control cascade identifies where control resides. Responsibility law must decide what to do with that information. The principal obstacle is not the absence of

⁴⁴See American Law Institute, *Restatement (Third) of Torts: Products Liability* (American Law Institute Publishers, 1998).

⁴⁵See Federal Trade Commission. FTC Policy Statement on Deception. Oct. 14, 1983. URL: <https://www.ftc.gov/legal-library/browse/ftc-policy-statement-deception> (visited on 08/27/2026).

doctrine. Direct negligence, products liability, agency, corporate criminal liability, officer liability, discovery, preservation, contribution, and indemnity already separate organizational from individual conduct. The obstacle is sequence. An injured person is often asked to identify the internal decisionmaker before obtaining the records that reveal who decided, while the organization describes the output as an act of the system and its employees reasonably fear that candid escalation will make them the easiest defendants.

This Part proposes a two-layer architecture that changes that sequence. As an analytic baseline, the enterprise possessing material control over a deployment is its external responsibility center: once a claimant proves the elements of an applicable cause of action, the enterprise cannot defeat attribution merely because no employee composed the output or because responsibility was divided internally. For designated high-impact deployments, legislation should turn that baseline into the nondelegable duty, primary-recourse rule, and production obligation specified below. A predeployment responsibility charter supplies the internal layer: it assigns risk ownership, decision and stop authority, reporting protection, evidence duties, and recourse before the incident. The first layer preserves a practical defendant and remedy. The second distinguishes the person who warned from the person who concealed, and the person who followed a reasonable constraint from the person who bypassed it.

The front door has two components with different domains. Identification is close to declaratory where one company visibly operates an integrated service; its reforming work lies in divided deployments where several enterprises control different layers and the user cannot see which one governed the risk. Focused production matters in both settings because the deployment record remains inside the enterprise even when its name is obvious. The proposal therefore states the two obligations separately and calibrates each to the architecture that makes it necessary.

The proposal is neither blanket strict liability nor employee immunity. Substantive liability remains tied to the governing law, and personal liability remains available for a

person's own actionable conduct. The architecture supplies a claimant-facing point of responsibility, a burden of producing deployment records, and a principled internal allocation.

A The Enterprise as the Claimant-Facing Responsibility Center

An enterprise is the external responsibility center when, in the course of business or a public function, it makes the relevant AI deployment available or incorporates it into a decision process and possesses material deployment control. That means authority over one or more features capable of creating or materially reducing the risk at issue—such as purpose, eligible users, deployment data, model or vendor selection, the harness and tools, warnings, monitoring, continuation, or rollback. A direct-to-consumer provider ordinarily occupies the role itself. In an API deployment, a downstream application company, hospital, employer, or agency may be the principal deploying enterprise even though an upstream provider remains a parallel responsibility center for risks within its retained control.

Under the proposed statutory regime, designation as a “responsibility center” has three consequences. First, subject to the ordinary rules governing jurisdiction and immunity, the enterprise is a proper initial defendant on the statutory deployment-duty claim rather than an unnamed model or unidentified employee; in an integrated service this usually confirms an already visible party, while in a divided deployment it supplies the identification rule. Second, it bears the focused preservation and production obligation specified below. Third, internal allocation cannot by itself reduce rights that governing law gives an injured outsider. None of the three establishes defect, breach, causation, compensable harm, or a measure of damages. Those elements still perform their limiting work.

Existing doctrine already supports much of this structure, but it does not universally guarantee the proposed front door. A corporation can be directly negligent through its own design, supervision, warning, monitoring, or deployment choices. It can be vicariously liable for employee torts within the governing scope rule. A commercial seller can face product duties without locating the defect in a culpable engineer. In jurisdictions that have enacted section 14 of the Uniform Electronic Transactions Act or an equivalent rule, contracts may be

formed through interactions involving authorized electronic agents. The law routinely treats an organization as the actor even when information and conduct are distributed among many employees.⁴⁶

Vicarious liability alone cannot deliver the proposal because a claimant may still have to prove an employee's underlying tort and scope of employment. For designated high-impact deployments, legislatures should therefore create a nondelegable deployment duty and a primary-recourse rule for every enterprise satisfying the material-control test above with respect to the risk at issue. "Primary recourse" means that the enterprise is the proper defendant on the statutory deployment-duty claim and that a claimant need not identify or first proceed against an employee; it does not dispense with proof of breach, causation, legally cognizable harm, or any defense supplied by the statute. Multiple enterprises can occupy parallel lanes, with contribution and indemnity following afterward. This is the Article's normative addition, not a claim that present common law in every jurisdiction already makes the company pay first.

Agency law is illustrative. The Restatement provides that an employer is subject to liability for an employee's tort committed within the scope of employment and separately recognizes a principal's direct liability for negligent selection, training, retention, supervision, or control.⁴⁷ Under the Third Restatement, an employee acts within scope when performing assigned work or engaging in a course of conduct subject to the employer's control; conduct leaves scope when it forms an independent course not intended to serve any employer purpose. Violation of an instruction is relevant but does not itself answer that inquiry. *Ellerth* applies agency principles in the specialized Title VII supervisor-harassment setting, while *Bushey* supplies a maritime enterprise-risk illustration in which unauthorized details

⁴⁶See American Law Institute, *supra* note 7, at §§ 7.03(1), 7.05, 7.07; American Law Institute, *supra* note 44, at §§ 1–2; Uniform Law Commission, *supra* note 23, at § 14.

⁴⁷See American Law Institute, *supra* note 7, at §§ 7.03(1), 7.05, 7.07.

remained characteristic of the business activity.⁴⁸ AI distribution intensifies the reason for organizational attribution: the output exists only because the enterprise assembled and maintained a system intended to act without contemporaneous employee review.

Disclaimers do not change the sequence. “AI-generated,” “for reference only,” “verify independently,” and “the model may make mistakes” can be relevant warnings. Their effect depends on the claim, the clarity and timing of disclosure, assent, public policy, the audience, and whether a reasonable design could have reduced the risk. Contract law has long limited exculpation where a term conflicts with public policy or defeats a nonwaivable duty. Warranty law separately regulates disclaimer, and products doctrine asks whether a reasonable alternative design or adequate instruction would have reduced the relevant risk; a warning is not categorically a substitute for reasonable design.⁴⁹ No phrase creates a general license to externalize enterprise risk, just as no marketing metaphor creates an artificial defendant.

Software classification remains jurisdiction specific. United States courts have not adopted a single rule treating every stand-alone service, model, or update as a product. A useful enacted comparator is the European Union’s 2024 Product Liability Directive, which expressly includes software, defines a manufacturer’s control to include the ability to supply updates or upgrades, addresses substantial modification, and creates proportionate disclosure and proof rules.⁵⁰ Its relevance here is structural, not transplantive: retained update authority, material modification, and information access can each be assigned to an existing legal person without characterizing the software itself as one.

⁴⁸See *Burlington Industries, Inc. v. Ellerth*, 524 U.S. 742, 756–59 (U.S. 1998); *Ira S. Bushey & Sons, Inc. v. United States*, 398 F.2d 167, 171–73 (2d Cir. 1968).

⁴⁹See, e.g., *Tunkl v. Regents of the University of California*, 60 Cal. 2d 92 (Cal. 1963); *Henningsen v. Bloomfield Motors, Inc.* 32 N.J. 358 (N.J. 1960); Uniform Commercial Code, U.C.C. § 2-316 (2025); American Law Institute, *supra* note 44, at § 2(b)–(c) & cmt. 1.

⁵⁰See Directive on Liability for Defective Products, Directive (EU) 2024/2853 (Oct. 23, 2024), arts. 4(1), 4(5), 4(18), 9–10. Member States must transpose the Directive; it neither governs a United States claim nor resolves the classification of every AI service.

1 Two Chatbots, Two Merits Results

The Hangzhou and Air Canada disputes show how an enterprise front door can coexist with different substantive outcomes. In *Moffatt*, Air Canada's website chatbot incorrectly told a customer that he could apply for a bereavement fare after travel. The airline later denied the request because its written policy required an application before travel. Before the British Columbia Civil Resolution Tribunal, the airline effectively sought to distinguish the chatbot from the rest of its website. The tribunal rejected that move: the customer was not required to determine which part of the airline's website was accurate, and Air Canada was responsible for information supplied through its channel.⁵¹ Attribution to the enterprise did the necessary work because the governing misrepresentation rule and evidence supported relief.

The Hangzhou Internet Court confronted more extravagant generated language. On June 29, 2025, a user asked an anonymized generative-AI application where a university program was located. After the application provided incorrect information and the user challenged it, the system apologized, purported to offer 100,000 yuan in compensation, and suggested litigation in the Hangzhou Internet Court. The user sought 9,999 yuan in damages. The court held that an AI system lacked the civil-subject status and subjective intention necessary to make the purported promise its own legal act. The text also did not manifest the provider's actual intention to pay. Treating the generative-AI offering as a service under China's Interim Measures rather than a product governed by the Product Quality Law, the court separately applied the Civil Code's general fault rule. The parties did not appeal, and the judgment was final when the court account appeared in January 2026.⁵²

The absence of a machine promisor did not leave the wrong answer outside law. The

⁵¹See *Moffatt* 2024 BCCRT at 27–31.

⁵²See Le, *supra* note 8 (discussing Civil Code art. 1165(1)). The publicly available court account anonymizes the application and provider and, rather than a complete published judgment, supplies the procedural and factual description used here.

court separately considered whether the provider had breached a differentiated, fault-based duty. The reported record included conspicuous limitation notices on the welcome page, in the user agreement, and in the interface, as well as retrieval-augmented generation to improve output reliability. It also included no adequate proof of actual loss or a causal effect on the user's application decision. On that record, the court found neither actionable fault nor compensable damage. The provider won after its conduct was examined, not because the model became the responsible party.

The two decisions therefore supply a useful pair. Air Canada used an automated channel to communicate fare information in a context where the tribunal found enterprise attribution and reasonable reliance. The anonymized Hangzhou application generated a fantastical promise outside provider manifestation, followed by a distinct negligence inquiry that the provider satisfied. A front-door rule neither binds an enterprise to every string of text nor immunizes it from every string of text. It keeps the enterprise in the case long enough for the relevant law of attribution, duty, and proof to operate.

B The Rule in Operation: *Garcia*

The narrowed rule can now be run against a complete deployment rather than stated in the abstract. *Garcia* is especially revealing because the published order preserved ordinary merits questions while the pleaded stack contained an application company, two founders, Google entities, user-created characters, and a minor user's own inputs. The exercise does not decide the settled action. It shows what the front door and production rule would have required, and which disputes they would have left for ordinary law.

1 Threshold and Initial Responsibility Center

The threshold would have been satisfied on the allegations accepted for Rule 12 purposes. The complaint pleaded recognized injuries, including wrongful death, and a material nexus to months of interaction with the Character.AI deployment. It targeted design and distribution features—including anthropomorphic presentation, access by minors, relational interaction,

warnings, and asserted safety omissions—rather than naming the generated character as a defendant. The district court held the application plausibly alleged to be a product and allowed core product and negligence theories to proceed beyond pleading.⁵³

Character Technologies would be the clearest initial responsibility center. It made the consumer deployment available and, on the pleaded record, operated the application through which users selected or created characters and interacted with generated responses. The risk asserted was tied to features an application operator ordinarily can govern: eligibility, persona presentation, memory and continuity, interaction policies, crisis detection, warnings, monitoring, and continued access. This is the vertically integrated setting in which the identification component adds little. The value of the front door lies in confirming that an *automatic instrumentality* does not end attribution merely because no employee selected the final exchange, and in opening the production obligation without requiring the claimant to identify an engineer.

The other named actors illustrate why “responsibility center” must remain risk specific. The complaint alleged that the founders personally designed, directed, and released the technology and that Google supplied substantial technical, financial, and organizational assistance with relevant knowledge.⁵⁴ Those allegations may support individual direct or aiding-and-abetting theories under governing law. They would not make an investor, former employer, licensor, or infrastructure supplier a deployment responsibility center by title alone. The statutory inquiry would ask whether each enterprise retained material authority over the relational risk—for example, control of the relevant model, application configuration, safety information, release condition, update, or stopping decision. The production record would test that proposition rather than presume it from corporate association.

⁵³See *Garcia*, 785 F. Supp. 3d at 1165–69, 1179–83. The court did not find that the deployment caused Sewell’s death or that any challenged feature was defective.

⁵⁴See First Amended Complaint paras. 24–25, 50–72, *Garcia v. Character Technologies, Inc.*, No. 6:24-cv-01903-ACC-EJK, Dkt. 11 (M.D. Fla. Nov. 9, 2024) [hereinafter *Garcia FAC*]; *Id.* at 1165–66, 1183–84.

2 The Focused Production Record

After the threshold, Character Technologies would identify the deployed model and application versions, the entities occupying the cascade, and the nonprivileged portions of the charter and incident record proportionate to the pleaded risk. In this case, the focused record would include the age-access rule and its enforcement; the character definition and features attributable to its creator; system instructions governing relational and self-harm interactions; memory and proactive-engagement settings; warnings actually displayed to this user and their timing; relevant evaluation protocols and results; crisis detections and interventions; material complaints and escalations received before the events; release approvals; and available restriction, rollback, or account-intervention controls.

That is not a demand for every model weight, every user's conversation, or unrestricted disclosure of a minor's data. Rule 26 proportionality, privilege, privacy protections, trade-secret orders, redaction, and staged discovery would continue to govern. A first stage could identify the system, responsible entities, governing policies, and preserved incident-specific events. Broader evaluation or design records would follow only as the pleaded theory and initial production justified them. The rule converts "the algorithm" into bounded requests tied to the control variables in Part I.

The record would also preserve facts favorable to the defense. The pleadings and motion papers disputed the effect of user selection, role play, user-edited messages, platform rules, content blocking, and an in-chat notice that character statements were fictional.⁵⁵ A configuration and event record could show which character features came from a third-party creator, which messages the user edited, which interventions fired, what warnings appeared, and whether a proposed safeguard would have altered the sequence. Production replaces categorical claims of unpredictability with evidence capable of supporting either side.

⁵⁵See Garcia, 785 F. Supp. 3d at 1167–69.

3 From Record to Merits and Internal Allocation

With that record available, ordinary doctrine would do the remaining work. Product classification would determine which design and warning rules applied. The process model would organize what was foreseeable at the relevant time; feedback would identify notice; authority and resources would locate feasible precautions; and the event sequence would test cause in fact and legal causation. The empirical studies discussed in Part III identify variables that providers can now test, but studies published after the 2023–2024 interactions cannot by themselves establish what a defendant historically knew or what precaution was then feasible. The proper inquiry is temporal and deployment specific.

The same discipline applies to later safety measures. A post-incident self-harm intervention may demonstrate that a control is technically describable, while evidence law and the substantive standard determine whether it is admissible or proves a pre-incident alternative.⁵⁶ The architecture neither converts subsequent remediation into liability nor permits the absence of contemporaneous records to be concealed behind a general claim that model behavior cannot be reconstructed.

Finally, the charter would separate external responsibility from individual recourse. It would show who approved minor access, who owned relational and self-harm risks, what dissent reached management, who controlled release resources, and whether anyone concealed an incident or bypassed a constraint. A documented warning by a safety employee would support protection, not an inference that the employee should become the claimant's sole target. A knowing suppression or unauthorized removal would remain available for individual consequences. Because no statutory charter governed the historical deployment, this application is prospective: it shows the allocation and evidence the proposed rule would require for the next high-impact relational system.

Garcia thus demonstrates the architecture's exact payoff. The initial enterprise was

⁵⁶See Fed. R. Evid. 407; *Garcia*, 785 F. Supp. 3d at 1169.

visible, so identification was easy. The hard questions concerned the record behind the interface, the division of retained control, and the causal significance of particular design choices. The production rule reaches those questions without predetermining their answers. The settled case remains a classification and pleading anchor; the worked example supplies the procedure that could have carried its merits analysis forward.

C Corporate and Individual Responsibility Run in Parallel

Three doctrines are easily conflated. *direct enterprise responsibility* concerns the organization's own design, policy, resource, supervision, warning, and deployment choices. *vicarious responsibility* attributes a person's conduct to an organization under agency or another relationship. *individual direct responsibility* concerns a person's own tortious, criminal, or otherwise actionable conduct. The three are analytically distinct. Direct enterprise and individual liability can arise independently; vicarious liability ordinarily presupposes an underlying employee tort even when the employee need not be named or personally recoverable.

The separation matters in both directions. A corporation does not avoid its own negligence by identifying an employee who implemented the policy. An employee does not become liable merely because the organization is liable. But acting for a corporation is not a general immunity from one's own tort. Restatement (Third) of Agency section 7.01 states that an agent is subject to liability to a third party harmed by the agent's tortious conduct, whether or not the principal is also liable.⁵⁷ The claimant must still prove that individual's duty, conduct, causation, and any required state of mind.

Veil piercing answers another question. It permits exceptional disregard of entity separateness where owners misuse the corporate form under jurisdiction-specific standards. It is not ordinarily required to sue an officer, engineer, clinician, or manager for that person's

⁵⁷American Law Institute, *supra* note 7.

own tort.⁵⁸ Treating the two as synonyms either overprotects individual wrongdoing or threatens routine employees with owners' liabilities. The charter should maintain the line: to the extent applicable law permits, ordinary good-faith work is allocated and indemnified internally; concealment, falsification, and personally tortious action are evaluated directly.

Corporate criminal law confirms that organization and individual can answer in parallel. *New York Central & Hudson River Railroad Co. v. United States* upheld attribution of employee conduct to a corporation for a federal offense committed within authority and for corporate benefit.⁵⁹ The responsible-corporate-officer doctrine, in tightly regulated public-welfare settings, can impose liability on an officer with responsibility and authority to prevent or correct violations even without conventional participation in the prohibited act.⁶⁰ State positive law can also authorize corporate liability for negligent homicide. In *Commonwealth v. Angelo Todesca Corp.*, the Massachusetts Supreme Judicial Court upheld a corporation's conviction under that State's motor-vehicle-homicide statute for its driver's negligent operation while conducting authorized corporate business.⁶¹ That jurisdiction-specific rule should no more be generalized into a national AI offense than *Park* should be generalized into automatic executive liability. Together, the authorities demonstrate a structural possibility: organizational liability need not displace individual responsibility, and managerial authority matters where positive law makes it matter.

Corporate fiduciary doctrine supplies a further, carefully bounded analogy. *Caremark* and its successors ask whether directors made a good-faith effort to establish board-level

⁵⁸See *Frances T. v. Village Green Owners Ass'n*, 42 Cal. 3d 490, 503–05 (Cal. 1986); *PMC, Inc. v. Kadisha*, 78 Cal. App. 4th 1368, 1380–81 (Cal. Ct. App. 2000). Position or general supervisory power alone remains insufficient; the claimant must establish personal participation, authorization, duty, and causation under the applicable rule.

⁵⁹*New York Central*, 212 U.S. 481.

⁶⁰See *United States v. Park*, 421 U.S. 658 (U.S. 1975).

⁶¹See *Commonwealth v. Angelo Todesca Corp.* 446 Mass. 128, 133–37 (Mass. Mar. 1, 2006).

information and reporting systems for mission-critical compliance risk.⁶² These are fiduciary claims belonging to the corporation, not a free-standing tort duty owed by directors to every person affected by AI. Their relevance is institutional. Law already recognizes that an organization can fail through the absence of reporting architecture, escalation, and board attention even when no director designed the failed component.

The charter converts that insight into a deployment-specific record. It allows a court to distinguish at least four situations: management declined resources after receiving a documented risk; a deployer acted within an approved and reasonably monitored design; an employee raised a concern that was suppressed; or a person knowingly bypassed a constraint. Without an *ex ante* allocation, all four can look like undifferentiated “AI failure.”

D The Mandatory Responsibility Charter

For designated high-impact deployments, law should require a written responsibility charter before release. “High impact” should be defined by the governing sector or enactment—for example, systems that materially determine access to employment, credit, housing, education, health care, liberty, essential public benefits, or physical safety, and relational systems offered to minors. A general-purpose model does not become high impact merely because it could be used in those settings; the obligation attaches to the covered deployment and to providers only for functions the sectoral rule assigns them.

The charter is not an aspirational ethics statement. It is a control instrument tied to a named deployment, version, and operating context. At minimum, it should contain the following:

⁶²See *In re Caremark International Inc. Derivative Litigation*, 698 A.2d 959 (Del. Ch. 1996); *Marchand v. Barnhill*, 212 A.3d 805 (Del. 2019); *In re Boeing Co. Derivative Litigation* 2021 WL 4059934 (Del. Ch. Sept. 7, 2021).

Table 3: Minimum Contents of a Responsibility Charter

Component	Required allocation	Evidentiary function
Deployment identity and scope	System version, purpose, population, data, tools, vendors, material limits, and change threshold.	Defines the system at issue and distinguishes later material modification.
Named risk owners	A responsible executive and an operational owner for each material hazard, with adequate authority and resources.	Identifies who receives information and can commission mitigation.
Deployment sign-off	Specified approvers, evidence considered, unresolved dissent, conditions, and expiration or review date.	Records authorization, knowledge, alternatives, and limits.
Protected escalation and dissent	Channels outside the ordinary reporting line; direct access to an independent reviewer or governing body; response deadlines.	Preserves notice, prevents suppression, and separates warning from concealment.
Stop and rollback authority	Who may pause, restrict, revoke credentials, restore a version, or decommission; triggers and response times; tested procedures.	Establishes practical ability to prevent or mitigate and tests whether oversight is real.

Component	Required allocation	Evidentiary function
Incident and configuration records	Versions, prompts and policies where proportionate, tool calls, alerts, overrides, complaints, decisions, and corrective actions; retention schedule.	Makes attribution, causation, discovery, and learning possible while allowing privacy minimization.
Independent safety review	Reviewer independence, access to data, evaluation scope, conflict rules, and remediation tracking.	Tests self-assessment and documents feasible controls.
Internal allocation and recourse	Indemnification, discipline, contribution, and safe-harbor standards fixed ex ante.	Protects good-faith compliance and preserves recourse for knowing misconduct.
External-rights savings clause	Express statement that the charter does not waive or narrow claimant rights, statutory duties, or regulatory powers.	Prevents internal governance from becoming an exculpatory contract with outsiders.

The components have distinct institutional pedigrees. Named owners, adequate authority and resources, and documented approval track the Organizational Sentencing Guidelines, FAA accountable-executive rules, and NIST's assignment of lifecycle and executive

responsibility.⁶³ Protected escalation tracks compliance-program reporting channels and sectoral anti-retaliation rules.⁶⁴ Stop and rollback authority and configuration records map respectively to NIST’s decommissioning and post-deployment response functions and the EU AI Act’s intervention and logging requirements; continuing assessment, assurance, and recordkeeping map to FAA safety-management rules. Independent-review safeguards draw instead on ex ante audit and disclosure proposals designed to preserve investigation incentives.⁶⁵ Internal recourse has a separate genealogy in conditional indemnification and compliance enforcement, while the external-rights savings clause preserves the public-policy limits on private exculpation.⁶⁶

This is not a governance scheme invented from whole cloth. The Organizational Sentencing Guidelines define an effective compliance program through leadership responsibility, operational responsibility with adequate resources and authority, training, monitoring, reporting without fear of retaliation, enforcement, response, and periodic risk assessment.⁶⁷ The Department of Justice’s compliance-program guidance asks whether a program is well designed, adequately resourced and empowered, and effective in practice; its 2024 update specifically asks how a company governs AI risks, controls unintended consequences, and

⁶³See U.S. Sent’g Guidelines Manual § 8B2.1(b)(2), U.S. Sentencing Commission. Guidelines Manual, Section 8B2.1: Effective Compliance and Ethics Program. 2025. URL: <https://www.ussc.gov/guidelines/2025-guidelines-manual/annotated-2025-chapter-8> (visited on 08/27/2026); Safety Management Systems, 14 C.F.R. pt. 5 (2025), § 5.25; National Institute of Standards and Technology, *supra* note 11, GOVERN 2.1, 2.3.

⁶⁴See U.S. Sent’g Guidelines Manual § 8B2.1(b)(5)(C), U.S. Sentencing Commission, *supra* note 63; Energy Reorganization Act Employee Protection Provision, 42 U.S.C. § 5851 (2024).

⁶⁵See National Institute of Standards and Technology, *supra* note 11, GOVERN 1.7, MANAGE 4.1; EU AI Act, Regulation (EU) 2024/1689, arts. 12, 14(4)(e), 17–20; FAA SMS Rule, 14 C.F.R. pt. 5, §§ 5.53, 5.71–5.73, 5.97; K. Ramakrishnan, *Tort Law at the Frontier of Artificial Intelligence*, 43 Yale J. on Reg. 960, 1047–50 (2026).

⁶⁶See Delaware General Corporation Law, 8 Del. C. § 145 (2026); U.S. Sent’g Guidelines Manual § 8B2.1(b)(6), U.S. Sentencing Commission, *supra* note 63; Tunkl, 60 Cal. 2d 92.

⁶⁷See U.S. Sent’g Guidelines Manual § 8B2.1, U.S. Sentencing Commission, *supra* note 63.

addresses deliberate or reckless misuse.⁶⁸ These materials arise in criminal enforcement and sentencing. They are transferable here not as liability rules, but because they operationalize information, authority, incentives, and response inside complex enterprises.

Aviation supplies a stronger safety-system analogy. On a staged implementation schedule, federal rules require specified aviation organizations to develop and maintain an organization-wide safety management system with accountable executive responsibility, hazard identification, risk controls, assurance, training, communication, and records.⁶⁹ The Aviation Safety Reporting Program seeks candid safety information through third-party reporting and a qualified, nonpunitive enforcement policy rather than blanket immunity.⁷⁰ The analogy is functional: AI deployments, like aviation systems, couple technical components with operators and organizational decisions, and near-miss information is easily suppressed if disclosure brings only personal exposure.

AI and medical governance sources converge on the same elements. NIST calls for documented lifecycle roles, executive responsibility, external feedback, risk-prioritized resources, third-party controls, monitoring, and safe decommissioning.⁷¹ A 2026 consensus produced by more than forty Chinese medical and research institutions organizes medical-AI governance around access evaluation, clinical application, patient rights, data governance, risk management, and competence; it specifically identifies multidisciplinary review, real-world validation, traceability, dynamic monitoring, human-machine responsibility, and

⁶⁸See U.S. Department of Justice, Criminal Division. Evaluation of Corporate Compliance Programs. Sept. 2024. URL: [https://www.justice.gov/criminal/criminal-fraud/page/file/937501/dl?inline=](https://www.justice.gov/criminal/criminal-fraud/page/file/937501/dl?inline=1) (visited on 08/28/2026).

⁶⁹See FAA SMS Rule, 14 C.F.R. pt. 5, §§ 5.1–5.17, 5.21–5.97; Federal Aviation Administration. Safety Management System. June 17, 2026. URL: <https://www.faa.gov/faq-aviation-safety-outreach/safety-management-system> (visited on 08/27/2026).

⁷⁰See Federal Aviation Administration. Aviation Safety Reporting Program. Apr. 2, 2021. URL: https://www.faa.gov/documentLibrary/media/Advisory_Circular/AC_00-46F.pdf (visited on 08/27/2026).

⁷¹See National Institute of Standards and Technology, *supra* note 11.

circuit breakers.⁷² Recent medical scholarship similarly identifies four conditions for meaningful oversight—epistemic capacity, cognitive space, decisional authority, and intervention effectiveness—and treats oversight as a property of the clinical system rather than a clinician’s ceremonial presence.⁷³

The convergence is the point. Criminal compliance, aviation safety, NIST risk management, and medical governance serve different ends and should not be collapsed into one legal standard. They nonetheless operationalize overlapping elements: named responsibility, information flow, authority, resources, protected reporting, review, and corrective capacity. Recent tort scholarship reaches the same institutional need from the opposite direction: liability by itself can discourage risk investigation and disclosure, while ex ante rules, independent auditors, and disclosure obligations can preserve the information on which prevention and adjudication depend.⁷⁴ The charter packages that shared infrastructure into a document a deployment and a court can use.

Two limits prevent compliance theater. First, possession of a charter is not compliance. Courts and regulators must ask whether the named people had actual authority and resources, reports reached them, stop mechanisms worked, and material changes triggered review. Second, unless governing positive law provides otherwise, compliance should be evidence rather than a conclusive defense. A carefully completed charter can support reasonableness; it cannot nullify a statutory duty, excuse a known unreasonable risk, or bar proof that the system operated outside its stated scope.⁷⁵

⁷²See Y. Cao, J. Wang, J. Wang, et al., *Expert Consensus on the Application and Governance of Artificial Intelligence in Medical Institutions (2026)* Intelligent Med. (May 23, 2026).

⁷³See D. van de Sande, N. Economou-Zavlanos and M. E. van Genderen, *Meaningful Oversight of Medical AI Beyond Human in the Loop*, 9 npj Digit. Med. 569 (July 23, 2026).

⁷⁴See Ramakrishnan, *supra* note 65, at 1047–50.

⁷⁵See American Law Institute, *Restatement (Third) of Torts: Liability for Physical and Emotional Harm* (American Law Institute Publishers, 2010), § 16(a); American Law Institute, *supra* note 44, at § 4(b). Jurisdictions differ, and a legislature or valid preemption rule may assign compliance a different effect.

E Who Can Enact and Administer the Charter

The proposal has no mystery legislature. Its designated-deployment structure permits adoption by the lawmaker that already governs the underlying activity. State legislatures can attach the charter, nondelegable duty, primary recourse, and production rule to companion services, health care, employment, education, insurance, consumer transactions, or other fields within state authority. Congress can legislate for federal programs, procurement, interstate sectors, and national schemes within its enumerated powers. Either lawmaker can authorize a sectoral regulator to specify forms, reporting intervals, change thresholds, and technical controls within intelligible statutory boundaries. The enacting law should identify the covered deployment, duty holder, minimum charter contents, production trigger, remedy, and relation to overlapping law; administration can then remain technically adaptive without relocating those policy choices.

California's companion-chatbot law demonstrates the form at state scale. Senate Bill 243 defines the covered relational service and its operator, requires conspicuous nonhuman disclosure in specified circumstances, mandates a published self-harm protocol and protections for known minors, requires reporting to the Office of Suicide Prevention beginning in 2027, and supplies a private remedy while preserving cumulative law.⁷⁶ It does not enact this Article's charter: it does not require named risk owners, deployment sign-off, protected escalation, stop authority, or the litigation production rule. It proves the adjacent institutional point. A legislature can designate a relational deployment, name the operator, combine protocol and reporting duties, and attach enforceable consequences without defining every general-purpose model as high impact.

Federal practice supplies implementation tools with different legal effects. In its 2025 section 6(b) companion-chatbot study, the Federal Trade Commission ordered named firms to report organizational roles, predeployment standards and thresholds, the positions

⁷⁶See California SB 243, Cal. Bus. & Prof. Code §§ 22601–22606.

of deployment approvers, considered and rejected mitigations, post-deployment monitoring, complaints, employee concerns, escalation, and response.⁷⁷ The compulsory study is neither a violation finding nor a generally applicable substantive rule. It shows that the organizational and deployment record the charter would standardize can be specified, collected, and verified rather than reconstructed as an abstract demand for “the algorithm.”

The Department of Justice’s 2024 compliance guidance works through enforcement evaluation rather than rulemaking. It asks prosecutors whether a company has assessed emerging-technology risk, governed business and compliance uses of AI, controlled unintended consequences and deliberate misuse, established accountability and human baselines, enabled confidential reporting, and continuously monitored and corrected model behavior.⁷⁸ The guidance does not create an *ex ante* tort duty. Its significance is operational: federal corporate-enforcement practice already evaluates the authority, information flow, incentives, and correction mechanisms the charter records.

FAA Part 5 supplies the mature delegated model. Its staged schedule requires covered aviation organizations to develop and, when applicable, maintain a formal, organization-wide safety management system with an accountable executive, identified management authority to accept safety risk, hazard analysis, corrective action, and retained records. For organizations subject to section 5.71(a)(7), the system also includes confidential employee reporting without fear of reprisal.⁷⁹ Part 5 is not an AI liability rule and does not supply every charter component. It demonstrates that named authority, protected reporting, continuing assurance, and evidence retention can be administered as an operating system rather than an aspirational policy.

⁷⁷See Federal Trade Commission, *supra* note 39; Federal Trade Commission. Model Order to File a Special Report Regarding AI Companion Chatbots. Sept. 10, 2025. URL: [https://www.ftc.gov/system/files/ftc_gov/pdf/AICompanionChatbot6\(b\)Order.pdf](https://www.ftc.gov/system/files/ftc_gov/pdf/AICompanionChatbot6(b)Order.pdf) (visited on 08/28/2026), specifications 3, 11–17.

⁷⁸See U.S. Department of Justice, Criminal Division, *supra* note 68, at 3–4, 7, 18.

⁷⁹See FAA SMS Rule, 14 C.F.R. pt. 5, §§ 5.3, 5.9(e), 5.11–5.17, 5.21, 5.23, 5.25, 5.53–5.57, 5.71–5.75, 5.95–5.97.

This division of labor answers the concern created by *Loper Bright*. An agency may implement a charter only within authority Congress or a state legislature has conferred; technical expertise does not itself create the primary-recourse duty or private remedy.⁸⁰ Courts, meanwhile, can use present law now: they can distinguish direct from vicarious responsibility, apply ordinary attribution, enforce preservation, order proportional discovery from proper parties and nonparties, and refuse to treat an AI system as an assetless substitute for a defendant. They cannot create the proposed statutory front door by discovery order alone. Legislative adoption supplies the duty and trigger; bounded agency implementation supplies sector-specific administration; adjudication supplies independent interpretation and enforcement.

F Information Asymmetry and the Production Rule

The claimant-facing layer fails if the enterprise can identify itself yet withhold the control record. AI deployments produce unusually concentrated information: model and application versions, system instructions, retrieval sources, tool permissions, evaluation results, incident queues, escalation tickets, internal approvals, and rollback events. The injured person usually sees only an output or decision. The enterprise knows which of several vendors, models, and configurations was active.

Ordinary civil procedure provides the foundation. Rule 26 permits discovery of relevant, proportional nonprivileged matter and requires initial disclosure of information and documents a party may use to support claims or defenses.⁸¹ Once litigation is reasonably anticipated, preservation duties attach under governing law. Rule 37(e), which draws on but does not itself create that duty, addresses electronically stored information that should have been preserved, is lost because reasonable steps were not taken, and cannot be restored or

⁸⁰See *Loper Bright*, 603 U.S. at 394–96.

⁸¹Federal Rules of Civil Procedure, Fed. R. Civ. P. § 26 (2025).

replaced. Curative measures require prejudice; an adverse-inference instruction or equivalent severe sanction requires intent to deprive.⁸² The charter should make reasonable preservation possible before a court has to reconstruct an undocumented system.

For covered high-impact deployments, legislation should add a focused *production rule*. Once a claimant pleads a recognized injury and a plausible material nexus to the deployment, the responsibility center must identify the system and entities in the cascade and produce the nonprivileged portions of the charter and incident record proportionate to the claim. Sensitive model, security, health, and personal information can be redacted or protected by order, including an order limiting disclosure of trade secrets or confidential research and commercial information.⁸³ Secrecy is not a reason to leave the system unidentified. The production obligation is procedural. It does not shift the ultimate burden of proving breach or causation unless the governing legislature expressly says so.

Existing evidentiary doctrine demonstrates why a calibrated inference may sometimes be warranted. Section 3 of the Restatement (Third) of Torts: Products Liability permits an inference that a defect existed at the time of sale or distribution, without proof of a specific defect, when the incident is of a kind that ordinarily results from product defect and was not solely the result of causes other than such a defect.⁸⁴ *Res ipsa loquitur* likewise permits a factfinder in defined circumstances to infer negligence from the type of event. Jurisdictions formulate the doctrine differently; modern formulations do not invariably treat rigid exclusive control as a separate element.⁸⁵ These doctrines are bounded; neither says

⁸²Federal Rules of Civil Procedure, Fed. R. Civ. P. § 37(e) (2025) advisory committee's note to 2015 amendment.

⁸³See Fed. R. Civ. P., Fed. R. Civ. P. § 26, r. 26(c)(1)(G). The EU Product Liability Directive offers a comparative enacted model of proportionate disclosure, protection of confidential information, and bounded presumptions; it is not current United States procedure. See EU Product Liability Directive, Directive (EU) 2024/2853, arts. 9–10.

⁸⁴American Law Institute, *supra* note 44, at § 3; see also Speller, 100 N.Y.2d at 41–43 (applying New York law and requiring evidence that excludes other causes not attributable to defendants).

⁸⁵See American Law Institute, *supra* note 75, at § 17 & cmt. b.

that any unexpected AI output proves negligence. The charter instead makes specific proof more available and reserves inference for the jurisdiction's ordinary conditions.

A missing record should have equally bounded consequences. If the system never captured information that no duty required it to capture, the absence is not automatically spoliation. In federal court, Rule 37(e) governs measures for ESI lost after a litigation-preservation duty arose and forecloses reliance on inherent authority or state law to prescribe different measures for that loss; it does not displace an otherwise applicable independent state-law spoliation claim.⁸⁶ If a sectoral charter mandate required logs and the enterprise omitted them, the statute should specify the consequence; violation can be evidence without becoming automatic liability for the underlying harm. This calibration makes documentation valuable rather than punitive by default.

G Protect the Reporter; Pursue the Concealer

An accountability system that directs every loss toward the nearest engineer will suppress the information it needs. The charter should therefore establish an internal safe harbor for a worker who, in good faith and on an objectively reasonable basis, documents a material risk, uses an authorized escalation or stop channel, preserves evidence, and acts within an approved role. The default internal consequences should be defense and indemnification by the enterprise to the extent permitted by governing law, protection against retaliation, and no adverse employment action based solely on protected reporting or justified intervention. Corporate indemnification statutes supply a model but also impose conditions; the charter cannot promise indemnity that positive law forbids.⁸⁷

The proposed protection for refusal should be comparably defined. A worker receives it when the worker reasonably believes the requested act is unlawful or presents an immi-

⁸⁶See Fed. R. Civ. P., Fed. R. Civ. P. § 37(e) advisory committee's note to 2015 amendment.

⁸⁷See DGCL, 8 Del. C. § 145.

ment risk of death or serious physical harm, gives the enterprise notice where feasible, and lacks time or a reasonably effective ordinary channel before the risk materializes. This is a proposed threshold informed by, not identical to, sectoral law.⁸⁸ A statute should also specify causation and remedy; the Energy Reorganization Act's contributing-factor test and clear-and-convincing same-action defense provide one balanced model.⁸⁹

The anti-retaliation component has a substantial statutory genealogy. The Sarbanes–Oxley Act protects specified reports of fraud and securities violations; the False Claims Act protects efforts to stop violations; and the Occupational Safety and Health Act protects safety complaints and related participation.⁹⁰ Aviation reporting programs encourage voluntary disclosure of safety events. The Patient Safety and Quality Improvement Act separately protects good-faith reporters, but its broad privilege and confidentiality rules for patient-safety work product should not be imported into this proposal because they would conflict with the focused production rule.⁹¹ The statutes' scopes, procedures, and privileges differ. They establish the transferable principle that accurate internal information is a safety resource and that retaliation can destroy it, not a general immunity from otherwise applicable law.

The safe harbor should have equally clear exceptions. It does not protect intentional or reckless infliction of harm, knowing falsification or destruction of records, concealment of a material incident, retaliation against another reporter, corruption of an evaluation, or

⁸⁸Compare Energy Reorganization Act, 42 U.S.C. § 5851 (protecting specified reports and refusal to engage in an identified unlawful practice), with *Whirlpool Corp. v. Marshall*, 445 U.S. 1, 9–13 (U.S. Feb. 26, 1980) (upholding a narrow occupational-safety protection for good-faith refusal in the face of reasonably perceived imminent danger and inadequate time for ordinary redress).

⁸⁹See Energy Reorganization Act, 42 U.S.C. § 5851, § 5851(b)(3).

⁹⁰See Sarbanes–Oxley Act Whistleblower Protection, 18 U.S.C. § 1514A (2024); False Claims Act Anti-Retaliation Provision, 31 U.S.C. § 3730(h) (2024); Occupational Safety and Health Act Anti-Retaliation Provision, 29 U.S.C. § 660(c) (2024).

⁹¹See Federal Aviation Administration, *supra* note 70; Patient Safety and Quality Improvement Act of 2005, 42 U.S.C. § 299b-21–299b-26 (2005), § 299b-22(a)–(b), (e).

unauthorized removal of a safety constraint. Nor does internal defense or indemnification eliminate liability that law independently imposes on the person's conduct. The charter should identify the standard and decisionmaker before a conflict arises, with independent review when senior management is implicated.

This allocation corrects an incentive mismatch. A safety engineer often has the clearest view of a failure mode but not the budget or release authority to fix it. A manager may control the schedule and resources but not understand the model detail. If both face an undifferentiated threat, the engineer may remain silent and the manager may avoid written decisions. Protecting documented escalation while preserving recourse against suppression makes authority and accountability coincide.

H The Two-Layer Rule

The proposed statutory rule can be stated in five steps.

1. A claimant pleads a recognized injury and a plausible material nexus to a covered high-impact AI deployment.
2. Each enterprise that makes the covered deployment available or incorporates it into a decision process and possesses material deployment control over the risk at issue is an external responsibility center. Multiple enterprises may qualify. None can defeat the statutory claim solely by showing that no natural person selected the precise output.
3. Each responsibility center identifies the deployed system and the entities occupying the relevant cascade, then produces the proportionate, nonprivileged charter and incident record. Identification confirms an already visible operator in an integrated deployment and exposes divided control where the stack is hidden; production supplies the merits record in both. Discovery, privilege, protection orders, and preservation law govern implementation.
4. With that record available, the claimant retains the ultimate burden to prove the el-

ements of the governing cause of action unless positive law provides otherwise. The charter neither waives outsider rights nor conclusively establishes reasonable care; it governs indemnification, defense, discipline, contribution, and recourse, and supplies evidence of authority, knowledge, and conduct.

5. Good-faith reporting, compliance, and justified intervention receive the protection specified by the statute and charter; knowing concealment, falsification, retaliation, and unauthorized bypass remain available grounds for internal discipline and for consequences independently authorized by positive law.

The rule resolves four problems at once. A victim has an identified enterprise against which a recognized claim can proceed if its elements are proved, without first reverse-engineering an org chart. A conscientious engineer can create a safety record without volunteering as the sole loss bearer. Management responsibility follows budget, release, and response authority rather than rhetorical commitment. And organizational fragmentation becomes evidence in a mapped cascade rather than a defense based on no single person's complete control.

Legislatures should mandate the charter and production rule in defined high-impact sectors. Courts can apply the architecture's underlying distinctions now: separate direct from vicarious liability; reject a machine as a liability sink; use ordinary attribution rules; enforce preservation duties; and evaluate individual conduct without demanding veil piercing. Once jurisdiction and the ordinary party or nonparty predicates are satisfied, courts can order proportional discovery from an entity possessing, controlling, or holding the relevant records; they cannot create the proposed statutory front door by discovery order alone.⁹² Regulatory adoption would make the allocation uniform and ex ante.

⁹²See Fed. R. Civ. P., Fed. R. Civ. P. § 26; Federal Rules of Civil Procedure, Fed. R. Civ. P. § 34(a)(1), 45 (2025).

I Residual Responsibility and Rebuttal

An enterprise-facing rule must leave room for a genuinely different control path. Six recurring defenses illustrate how the cascade disciplines rather than erases rebuttal.

Intentional user circumvention. A user who deliberately jailbreaks a system, supplies false context, defeats a reasonably designed confirmation, or directs a prohibited act may bear comparative or primary responsibility under governing law. The enterprise can show the constraint, warning, foreseeable bypass testing, and causal effect. The word “jailbreak” alone is not a defense if the technique was common, encouraged by the product’s use, or readily preventable.

Modification. A downstream actor who removes safeguards, changes weights, adds dangerous tools, or deploys outside documented limits may become the relevant controller. New York law illustrates the distinction: substantial modification can defeat a design-defect theory while a failure-to-warn claim based on foreseeable modification may remain available in appropriate circumstances.⁹³ Other jurisdictions define modification and warning duties under their own statutes and common law. Version and provenance records should show the change rather than leaving the issue to brand attribution.

Cyberattack. Unauthorized intrusion may be an intervening cause, but security is itself a foreseeable deployment condition. The questions are which threat was reasonably in scope, whether credentials and tools were exposed, which alerts arrived, and whether response was possible. A sophisticated novel attack and an unpatched known vulnerability do not occupy the same control position.

⁹³See *Liriano v. Hobart Corp.* 92 N.Y.2d 232, 237–43 (N.Y. 1998).

Government compulsion. A contractor may invoke the federal military-equipment defense only where the United States approved reasonably precise specifications, the equipment conformed to them, and the supplier warned the United States of dangers known to the supplier but not the government.⁹⁴ Two 2026 decisions sharpen the control inquiry. *GEO Group v. Menocal* holds that *Yearsley* supplies a merits defense, not derivative sovereign immunity, for conduct the government lawfully authorized and directed.⁹⁵ *Hencely v. Fluor Corp.* rejects broader displacement where the government neither ordered nor authorized the challenged conduct: *Boyle* requires a significant conflict with an identifiable federal policy and protects a contractor only insofar as the suit challenges a government decision the contractor carried out.⁹⁶ A procurement relationship, political pressure, silence about the challenged means, or conduct contrary to instructions does not transfer control by itself. Part VI addresses the reciprocal limit on governmental coercion.

Independent deployer override. An upstream provider may rebut responsibility for a harm tied to a downstream purpose, data set, tool, or override it neither controlled nor reasonably could prevent. Conversely, contractual reservation of a right to audit, update, suspend access, or dictate safeguards is evidence of retained control. The result should follow the deployed function, not a standardized label that declares every party an independent contractor.

Unavoidable event. Some harmful outputs will remain unforeseeable after reasonable testing, monitoring, warning, and response. Ordinary standards of fault leave room for that conclusion; the published account of the Hangzhou decision supplies a comparative illustration rather than controlling United States law. A charter strengthens the defense by

⁹⁴See *Boyle v. United Technologies Corp.* 487 U.S. 500, 512 (U.S. 1988).

⁹⁵*GEO Group, Inc. v. Menocal*, 607 U.S. 438, 444–48 (U.S. Feb. 25, 2026).

⁹⁶*Hencely v. Fluor Corp.* 608 U.S. 31, 37–42, 46 (U.S. Apr. 22, 2026).

showing what risk was assessed, why controls were selected, what feedback existed, and why the proposed alternative would not have prevented the injury.

Residual responsibility therefore does not undermine the architecture. It is what makes the architecture a legal framework rather than a slogan. The enterprise faces the claimant and produces the map; the parties then litigate where the causal control path actually ran under the governing burdens of proof. The machine never has to be invented as the missing defendant.

V The Anti-Substitution Principle

The responsibility architecture governs harms produced through AI deployment. A related problem arises when deployment becomes an asserted reason to withdraw responsibility altogether. A public agency points to an automated portal instead of providing the process its benefit scheme requires. A care institution points to a companion bot instead of staffing required support. A professional points to a model output instead of exercising the judgment attached to the professional role. In each setting, the technology may perform valuable work. The legal error is to infer from availability that an existing duty has been satisfied.

The *anti-substitution principle* prevents that inference:

Where law independently imposes a duty of care, process, accommodation, professional judgment, or public service, the availability or use of AI does not by itself discharge or narrow that duty. Automation may perform part of the duty when the governing law permits and the deployment preserves the protection the duty exists to secure.

This is a rule of performance, not a preference for manual administration. An *automatic instrumentality* can perform every operational step needed to satisfy a duty; its availability alone is not performance. The court asks whether the legally protected function remains available through the deployed system and whether an answerable institution maintains correction when the automated path fails. That formulation permits AI to improve

access, consistency, or capacity while keeping the duty attached to the person or organization law already made answerable.

A A Rule Courts Can Apply

Four questions turn the principle into a decision rule.

First, *what is the independent source and scope of the duty?* The source may be statute, regulation, contract, custody, guardianship, professional law, a tort special relationship, or constitutional process attached to a protected interest. The principle does not manufacture that predicate. It prevents a duty holder from treating automation as an implicit amendment to it. Constitutional due process, for example, supplies no general affirmative obligation to provide care; where legislation creates an enforceable entitlement, the text and remedial structure define the protection.⁹⁷

Second, *which performance does the duty require?* Law often protects a function rather than a particular staffing form. Notice must convey the material basis for a decision. Review must reach an actor with information and authority to correct error. Professional judgment must integrate the considerations the governing standard makes relevant. An accommodation must afford meaningful access. Care obligations may require monitoring, response, continuity, or safe transition. Stating the protected function keeps the inquiry from collapsing into either “a human must do every step” or “a chatbot exists, so the duty is complete.”

Third, *does the deployed process preserve that function in practice?* Relevant evidence includes whether the system can handle foreseeable exceptions; whether an affected person receives intelligible notice; whether accessibility and language needs are met; whether the person can contest the material basis; whether uncertainty and limitations reach a com-

⁹⁷See *DeShaney v. Winnebago County Department of Social Services*, 489 U.S. 189 (U.S. 1989); cf. *Health & Hospital Corp. of Marion County v. Talevski*, 599 U.S. 166 (U.S. 2023) (recognizing section 1983 enforcement of two specified Federal Nursing Home Reform Act rights).

petent reviewer; whether the reviewer has time, authority, and an effective intervention; whether failure triggers escalation; and whether the institution monitors outcomes after release. These are cascade variables. Functional equivalence is demonstrated through the control path, not through the label “human in the loop.”

Fourth, *who remains responsible for performance and correction?* Delegating a task to software does not relocate the duty to the software. The institution must name a controller who can restore service, decide an exception, correct a record, or arrange an alternative when the automated path fails. The responsibility charter supplies that allocation, while the production rule preserves the evidence needed to test it. Where the deployed process performs the legally required function, the duty is satisfied through automation. Where it does not, technological novelty is neither performance nor excuse.

The rule gives a court a narrow holding. The court need not decide that automated service is categorically inferior, that every claimant has a right to personal attention, or that one model design is mandatory. It can identify the predicate duty, specify its protected function, examine the deployed control path, and determine whether the responsible institution provided the performance law requires.

B Public Administration and the Right to an Effective Path

Public-benefits cases show the rule already latent in doctrine. In *K.W. ex rel. D.W. v. Armstrong*, Idaho used a budget tool to allocate community-based services to people with developmental disabilities. The notices did not adequately disclose how individualized budgets were calculated or permit recipients to identify and challenge errors. The Ninth Circuit affirmed preliminary relief on procedural-due-process grounds.⁹⁸ In *Arkansas Department of Human Services v. Ledgerwood*, changes associated with an assessment methodology reduced in-home attendant-care hours, and the Arkansas Supreme Court upheld temporary

⁹⁸*K.W. ex rel. D.W. v. Armstrong*, 789 F.3d 962 (9th Cir. 2015).

injunctive relief tied to the agency's failure to follow required rulemaking procedure.⁹⁹

Neither decision forbade algorithmic allocation. Each required automation to remain inside the law governing the benefit. The agency could use a model to improve consistency or direct scarce resources, but the model could not replace the notice, explanation, rulemaking, or contest path the legal scheme required. Danielle Citron's account of technological due process explains the general mechanism: automated public systems can encode policy and error at scale while notice and hearing remain designed for discrete human decisions.¹⁰⁰ The operative response is to preserve an intelligible basis, an opportunity to contest, and authority to correct before a system-level error becomes a durable administrative rule. Those protections inherit their force from existing public law, including the due-process principles exemplified by *Goldberg v. Kelly*.¹⁰¹

The anti-substitution principle adds a control test to that doctrine. A portal that explains a decision and routes a supported challenge to an empowered official may improve process. A chatbot that repeats the result, cannot expose the material basis, and loops the beneficiary back to itself does not create review merely because it responds in natural language. The difference lies in the upward feedback and correction authority. A public body remains the duty holder and must be able to show where the path terminates.

The same analysis applies to capability tiering. A government may restrict tools, data, or model access for security, competence, or resource reasons. If a lower tier determines access to an existing entitlement, the affected person must retain whatever explanation, review, and accommodation the governing law requires. The availability of a general public chatbot cannot substitute for an appeal channel, and a hidden automated restriction cannot substitute for the accountable decision the statute assigns to an official.

⁹⁹Arkansas Department of Human Services v. Ledgerwood 2017 Ark. 308 (Ark. 2017). 530 S.W.3d 336.

¹⁰⁰See D. K. Citron, *Technological Due Process*, 85 Wash. U. L. Rev. 1249 (2008).

¹⁰¹See *Goldberg v. Kelly*, 397 U.S. 254 (U.S. 1970).

C Care, Professional Judgment, and Relational Systems

Medical AI illustrates how the principle supports high-value automation. A decision aid may outperform a clinician on a bounded prediction, extend expertise to an underserved setting, or detect a pattern a person would miss. The professional duty can nevertheless require facts outside the model, patient preferences, uncertainty, informed consent, follow-up, and intervention when the workflow fails. A 2026 medical-AI consensus accordingly combines access evaluation and real-world validation with multidisciplinary review, traceability, continuing monitoring, competence, and circuit breakers.¹⁰² Recent scholarship identifies four conditions for meaningful oversight—epistemic capacity, cognitive space, decisional authority, and intervention effectiveness—and treats oversight as a property of the clinical system rather than a ceremonial human presence.¹⁰³

Under the four-question rule, the institution first states the actual duty: for example, competent professional judgment and follow-up under the applicable law. It then identifies which functions the model performs and which remain elsewhere in the workflow. The cascade shows whether the clinician receives enough information, time, and authority to act; the charter shows who responds when performance drifts or an exception exceeds the model's scope. AI can carry a large share of the work. Responsibility stays with the professional and institution that law authorizes and regulates.

Relational AI presents a different opportunity and the same legal structure. A companion can reduce isolation, translate, remind, help a user rehearse a difficult exchange, or extend supportive contact. Its availability does not authorize a care institution to withdraw staffing required by another law, an insurer to deny covered treatment, or a public program to replace an accessible service path with an endless automated loop. The relevant com-

¹⁰²See Cao, Wang, Wang, et al., *supra* note 72.

¹⁰³See Sande, Economou-Zavlanos and Genderen, *supra* note 73; see also World Health Organization. Ethics and Governance of Artificial Intelligence for Health. 2021. URL: <https://www.who.int/publications/i/item/9789240029200> (visited on 08/27/2026).

parison is not human versus machine in the abstract. It is the protection promised by the governing duty versus the protection actually delivered through the deployment.

The principle also disciplines provider-controlled transitions. A voluntary companion service does not become property owned by the user, and a provider does not incur a general obligation to operate every persona indefinitely. But where consumer, disability, health, contract, or another law independently recognizes a relevant duty, foreseeable dependence can make continuity notice, data export, crisis safeguards, and a safe exit part of its performance. Research on companion loss and repeated relational exposure makes those consequences predictable enough to enter the process model.¹⁰⁴ The protected interests are the user's autonomy, psychological health, safety, and access to an intelligible transition—not ownership of an artificial relationship.

D Automation That Performs the Duty Is Not Substitution

The principle's boundary is affirmative. It protects automation that actually performs the legal function. A multilingual benefits assistant that makes notices intelligible, a clinical system that gives a professional better evidence and usable override authority, or a companion service that supplements human support can strengthen the underlying right or duty. Requiring the institution to name the function, preserve feedback, and maintain correction makes those deployments more dependable.

Two distinctions keep the rule exact. The first separates a protected function from a demanded medium. Unless the governing law says otherwise, there is no general right to a human conversation, a particular employee, a favorite model, or permanent access to a commercial persona. The second separates human responsibility from human tokenism. Preserving a duty does not require a person to repeat the model's work. It requires an institutional path by which the beneficiary can understand, challenge, and obtain the performance

¹⁰⁴See De Freitas et al., *supra* note 3; Banks, *supra* note 3; Kirk et al., *supra* note 2.

the duty secures.

The anti-substitution principle therefore completes the responsibility architecture. The front-door rule prevents an enterprise from pointing inward to an unidentified worker or outward to an artificial agent. Anti-substitution prevents a duty holder from pointing to the mere existence of an automated service. In both settings, the cascade asks the same question: where is the actor with the information and authority to make the legal protection effective?

VI Capability Boundaries and State Power

The framework developed so far is allocative. It identifies legal subjects, permits functional classifications, maps control, and places external and internal responsibility. Allocation alone does not require every user to receive every capability or permit a government to compel every desired use. This Part identifies two boundaries that follow from the same control analysis.

First, unequal capability can be legitimate, but consequential tiering should be proportionate, transparent, reviewable, and capable of escalation. Second, the government may choose its tools and specifications, especially for national security, but procurement authority is not a general power to punish a supplier's protected refusal or conscript its technology for every lawful use. Part V has already supplied the parallel boundary for care and public service: automation can perform an existing duty, but its availability does not itself discharge one.

A Capability Stratification Without Status Stratification

AI systems are already stratified. A public user may receive a conversational interface with limited tools. A professional may receive longer context, domain data, and auditable workflow integration. A researcher may receive weights or a sandbox. A government may receive classified hosting, specialized data, and operational tools. An enterprise may reserve

higher rate limits and support for customers able to pay. These differences can reflect security, competence, scarcity, legal obligation, or commercial choice. They do not make the system a different kind of legal subject.

The first mistake is therefore to announce that unequal model access is inherently unlawful. Law regularly permits differentiated access to dangerous equipment, licensed services, secure data, and scarce resources. A medical deployment should not expose patient records and ordering tools to every public user; a government system may require a cleared environment; a research release may need controls unavailable in a consumer application. Risk-graded regulation itself relies on differentiation. The EU AI Act, for example, connects obligations to the intended use, actor, and risk of a system rather than demanding a single capability for all.¹⁰⁵

The opposite mistake is to treat every capability boundary as a private technical fact. Tiering can determine who receives an educational accommodation, who can challenge a benefits decision, which clinician sees uncertainty, or whose account is automatically banned. In those settings, the relevant provider or institution remains subject to ordinary antidiscrimination, consumer, contract, administrative, and constitutional rules. “The model assigned the tier” does not alter the legal subject or the nature of the decision.

Four governance conditions make stratification legible without demanding uniformity.

1. *Proportionality*: a restriction should track the material risk, resource, or legal constraint asserted. A limit designed for autonomous execution should not silently disable harmless explanation; a security rule should not become a pretext for excluding a protected class.
2. *Transparency*: in a consequential setting, affected persons should know that capability has been limited, the category of reason, the decisionmaker, and the route to review. Transparency does not require publication of exploit details, model weights, classified

¹⁰⁵See EU AI Act, Regulation (EU) 2024/1689.

information, or another person's data.

3. *Review*: a person should be able to contest a mistaken tier, automated ban, or consequential capability denial before an actor with information and authority to change it. The required process depends on the underlying legal interest; a consumer feature request and termination of essential benefits are not equivalent.
4. *Escalation*: where broader capability can be used safely by a qualified actor or under added controls, the system should identify a path to that setting rather than treat the lowest public tier as the limit of possible service.

These conditions derive from control rather than a new equality entitlement. NIST's framework and its generative-AI profile ask organizations to document roles, context, human oversight, credentials, feedback, impacts, and change controls; the deployment cascade identifies who can review and change a tier.¹⁰⁶ Public-law requirements can add notice and hearing when government action affects a protected interest. Civil-rights law can prohibit criteria or effects within its scope. Consumer law can police a misleading representation that all users receive a capability when a hidden rule says otherwise. The minimum leaves each source of law intact and makes its object visible.

Capability and responsibility should also remain aligned. A vendor that gives a professional deployer powerful tools but withholds necessary risk information creates a different

¹⁰⁶See National Institute of Standards and Technology, *supra* note 11; National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile. July 2024. URL: <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf> (visited on 08/27/2026). Federal management memoranda apply comparable testing, monitoring, review, acquisition, and portability expectations to executive-agency uses, without creating a general private right or rules for private platforms. See Office of Management and Budget. Accelerating Federal Use of AI Through Innovation, Governance, and Public Trust. Apr. 3, 2025. URL: <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-21-Accelerating-Federal-Use-of-AI-through-Innovation-Governance-and-Public-Trust.pdf> (visited on 08/27/2026); Office of Management and Budget. Driving Efficient Acquisition of Artificial Intelligence in Government. Apr. 3, 2025. URL: <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Driving-Efficient-Acquisition-of-Artificial-Intelligence-in-Government.pdf> (visited on 08/27/2026).

record from one that supplies evaluation, training, monitoring, and stop controls. A deployer that requests a restricted capability and then ignores its conditions cannot rely on the vendor's brand. An automated gate that no one can explain or review reveals a missing controller in the cascade. The legal question is not whether tiering exists, but whether the actor responsible for the consequential boundary can justify and correct it under the applicable rule.

B State Power: Procurement, Conditions, and Retaliation

The state occupies two positions in AI governance. It is a regulator capable of imposing public duties, and it is an enormous purchaser capable of shaping private systems through contracts. Both powers are legitimate and bounded. The 2026 dispute between Anthropic and the United States Department of War provides the clearest current anchor.

1 The Anthropic–Pentagon Controversy

Anthropic developed Claude and, beginning in 2024, made versions available to federal intelligence and defense users through commercial and classified computing environments. A 2025 agreement carried a ceiling of 200 million dollars. When the government later demanded authorization for “all lawful uses,” Anthropic substantially narrowed its usage restrictions but preserved two contractual red lines: lethal autonomous warfare and mass surveillance of Americans. The deployed models were static; Anthropic could not see government prompts, technically block a particular use, or reach the systems through a back door. The parties reached no agreement by the government's February 27 deadline, although negotiation and collaboration continued afterward. According to the district court's record, three challenged actions followed. A February 27 Presidential Directive ordered government-wide cessation with a six-month transition. Later that day, the Secretary of Defense directed commencement of a supply-chain designation process while immediately announcing a broader prohibition reaching military contractors and suppliers. A March 3 determination, noticed March 4,

then formally invoked 10 U.S.C. § 3252 for national-security-system procurement.¹⁰⁷

In March 2026, Judge Rita F. Lin preliminarily enjoined the directives and designation. The order emphasized a boundary that survives every procedural stage: it did not require the government to use Anthropic's products and did not prevent a lawful transition to another provider.¹⁰⁸ On August 27, 2026, the district court granted summary judgment to Anthropic on its First Amendment retaliation and Fifth Amendment stigma-plus due-process claims and on core Administrative Procedure Act claims, entered final judgment, and closed the case. It held that 10 U.S.C. section 3252 addressed adversarial sabotage or subversion, not punishment of a domestic supplier for its public safety position; the designation was beyond authority, procedurally deficient, arbitrary, and pretextual. It also found no statutory authority for the Secretary's secondary boycott of firms doing business with the military. The court vacated the designation and barred enforcement of the retaliatory directives while preserving ordinary lawful procurement choice.¹⁰⁹

The controversy is important because the strongest argument for the government is substantial. Military purchasers must control mission requirements, reliability, availability, classified integration, and lawful use. A vendor's unilateral restriction can create operational risk, and the government need not design a mission around one company's policy. National-security procurement also involves predictive judgments that courts are institution-

¹⁰⁷See *Anthropic PBC v. United States Department of War* No. 3:26-cv-01996-RFL, Dkt. 250 (N.D. Cal. Aug. 27, 2026). summary-judgment order. The record does not establish that Claude was deployed in Project Maven; this Article accordingly uses the case caption and the Anthropic–Pentagon controversy rather than treating “Claude–Maven” as an adjudicated fact.

¹⁰⁸See *Anthropic PBC v. United States Department of War* No. 3:26-cv-01996-RFL, Dkt. 134 (N.D. Cal. Mar. 26, 2026). preliminary injunction.

¹⁰⁹See *Anthropic Merits Order* No. 3:26-cv-01996-RFL, Dkt. 250; *Anthropic PBC v. United States Department of War* No. 3:26-cv-01996-RFL, Dkt. 251–252 (N.D. Cal. Aug. 27, 2026). vacatur, permanent-injunction, and judgment documents. This was an appealable district-court judgment, not a final appellate disposition; it remained subject to appellate review as of this Article's date. A separate D.C. Circuit proceeding concerning 41 U.S.C. sections 4713 and 1327 presents distinct statutory and procedural questions; its April 2026 emergency-stay ruling did not decide the merits. See *Anthropic PBC v. United States Department of War* No. 26-1049 (D.C. Cir. Apr. 8, 2026). order denying emergency stay and expediting merits review.

ally cautious to second-guess. Nothing in the jurisprudential minimum gives an AI provider a right to a federal contract.

That argument supports selection, specification, and transition. It does not establish a power to impose market-wide penalties because a company criticizes the government's position or refuses to alter its general product. The final opinion distinguished the government as market participant administering a contract from the government as sovereign imposing a public stigma, government-wide ban, and third-party boycott. Government contractors retain First Amendment protection against retaliation for protected expression, subject to the government's interests as contractor and purchaser. *Board of County Commissioners v. Umbehr* applied that principle to termination or nonrenewal of an independent contractor, and *O'Hare Truck Service, Inc. v. City of Northlake* rejected removal from a rotation list based on political association.¹¹⁰ The unconstitutional-conditions doctrine more generally prevents government from denying a benefit on a basis that infringes a protected interest even where the claimant had no entitlement to the benefit itself.¹¹¹

The procurement boundary can be stated as a sequence. It is this Article's synthesis of purchaser discretion, contractor-retaliation doctrine, and unconstitutional-conditions law, not a four-part test announced by the district court. Because the sequence rests on *Umbehr*, *O'Hare*, *Perry*, and the distinction between purchasing and sovereign sanction, it does not depend on the August 27 judgment remaining undisturbed in every respect:

1. The government may define the mission, security conditions, performance requirements, and lawful uses for which it seeks a system.

¹¹⁰See *Board of County Commissioners v. Umbehr*, 518 U.S. 668 (U.S. 1996); *O'Hare Truck Service, Inc. v. City of Northlake*, 518 U.S. 712 (U.S. 1996).

¹¹¹See *Perry v. Sindermann*, 408 U.S. 593 (U.S. 1972); *Agency for International Development v. Alliance for Open Society International, Inc.* 570 U.S. 205 (U.S. 2013). This Article uses unconstitutional conditions as an analytic bridge; the district court's final holdings sounded directly in retaliation, due process, and the APA.

2. A supplier may decide which product it offers and may express a view about uses, subject to contracts and law that validly bind it.
3. The government may select another supplier, negotiate different terms, build its own system, or take a tailored action authorized by statute for a genuine security or performance risk.
4. It may not use procurement as a pretext to punish protected criticism, compel adoption of a viewpoint outside the funded program, or impose sweeping third-party disabilities unsupported by the invoked authority.

This is not a special privilege for safety-branded companies. A vendor cannot convert every commercial limit into protected speech, abandon a promised service without contractual consequence, or invoke ethics to conceal a security defect. The court should identify the expressive activity, adverse action, causal evidence, statutory authority, and government's purchaser interests under ordinary doctrine. The *Anthropic* record mattered because the stated sanctions extended far beyond choosing a different model and the court found the asserted security rationale unsupported.

The subject-status invariant clarifies the dispute. Claude was not the constitutional claimant; Anthropic was. The government was not bargaining with an artificial rights holder; it was setting terms with a corporation that designed and maintained a system. The contested constraints embodied human and organizational choices. Protecting or rejecting those choices required no ruling on whether Claude's generated output itself possessed First Amendment rights.

The same principle constrains providers. A company cannot invoke its own model constitution to override positive law, evade a valid contract, or deny public remedies. Internal policy is evidence of design and commitment; external law judges it. The framework is therefore reciprocal. It prevents the firm from transferring responsibility downward to the model, and it prevents the state from treating control over procurement as control over every

lawful private use or viewpoint.

Taken together, capability stratification and bounded state power complete the architecture's reciprocal account of control. The framework allocates authority to those who actually hold it and preserves the legal limits that accompany that authority.

Conclusion

AI responsibility does not disappear when no person selects the harmful output. It moves into the deployment. Organizations choose objectives, models, interfaces, memory, data, tools, user populations, evaluations, monitoring, and stop conditions. Those choices are observable, governable, and legally consequential even when the model's internal path is difficult to explain.

The deployment-control cascade makes that responsibility visible. It traces authority and resources toward operation, feedback toward those able to respond, and override or shutdown power to the point where intervention can work. Its process model, control actions, feedback, and records translate into the facts ordinary law already uses: authorization, notice, foreseeability, feasible precaution, causation, and proof. Probabilistic output changes the location of the evidence; it does not create a responsibility vacuum.

The two-layer architecture turns that map into an institutional rule. For designated high-impact deployments, an enterprise with material control over the risk should face the claimant under a nondelegable deployment duty. In a divided stack, the identification rule reveals which enterprises governed which risk. In every covered deployment, the production rule supplies the proportionate record needed to test the merits. The claimant still proves recognized harm, breach or defect, causation, and every other substantive element. The front door makes those rights provable without converting them into strict liability.

Inside the enterprise, the responsibility charter aligns accountability with authority before harm occurs. It names risk owners, fixes approval, records dissent, protects escalation, tests stop and rollback power, and preserves incident evidence. The allocation protects the

person who reports a danger and preserves recourse against the person who conceals, falsifies, retaliates, or bypasses a constraint. Public responsibility and internal fairness reinforce the same information system.

The anti-substitution principle carries that logic into care, professional judgment, accommodation, and public service. An automatic instrumentality may perform a legally required function and often improve it. Its mere availability does not perform the duty. The answerable institution must preserve the protection the duty exists to secure, including an effective path to notice, review, intervention, or correction where the governing law requires one.

None of this requires a machine defendant. Unless positive law constructs a new juridical platform, the legal subjects remain the natural persons and organizations capable of holding rights, bearing duties, producing records, and answering remedies. Functional classifications may still vary: the same deployed system can be a product, service, instrument, or communicative medium for different rules. The stable point is responsibility. Courts should classify the function, reconstruct the control path, open the evidentiary front door, and keep the obligation attached to an actor law can reach.

Source Note: The Hangzhou Generative-AI Decision

Published professional accounts identify the proceeding as *Liang v. Technology Company*, Hangzhou Internet Court, (2025) Zhe 0192 Min Chu No. 18143 (decided Dec. 3, 2025).¹¹² The comparative discussion relies on Le Weikun's *People's Court Daily* account, hosted on an official Zhejiang government site. A Supreme People's Court Intellectual Property Court webpage reproducing a CCTV interview with the head of the work-report drafting group independently confirms the dismissal, subject-status rule, reasonable-precaution analysis,

¹¹²See Shanghai Qiaowen Law Firm. AI Labeling Duties: Adjudicative Boundaries and Compliance Guidance Based on Two Cases. 2026. URL: <https://www.qiaowen.com/cn/details/0803/166cac21f6bf434e> (visited on 08/28/2026).

and absence of proven loss.¹¹³ A complete judgment with a stable, anonymous public URL was not identified. The excerpts below are the author's translation of the court account, not the judgment. They preserve its anonymization, are unofficial, and are supplied for source checking; the Chinese source controls.

Status and manifestation. “AI lacks civil-subject status and cannot make a manifestation of intention.” The generated compensation promise was not the provider's manifestation because the system could not serve as its messenger, agent, or representative; the provider had not used the model to communicate an intention; ordinary practice did not support reasonable reliance on the random promise; and the provider had not represented that it would be bound by generated content.

Classification and governing rule. The court treated the generative-AI offering as a service under the Interim Measures, not as a product under product-quality law, and applied the Civil Code's general fault rule in article 1165(1).

Duty, loss, and causation. The report describes three layers of care: strict review of legally prohibited information; prominent disclosure of inherent accuracy limits, including immediate warnings where major interests are at stake; and basic reliability precautions using generally accepted technical measures. It reports that the provider displayed limitations in the welcome page, agreement, and interface and used retrieval-augmented generation. The court also found no adequate proof of actual loss or material effect on the user's application decision.

¹¹³See Le, *supra* note 8; Supreme People's Court Intellectual Property Court. Head of the Drafting Group Interprets the Supreme People's Court Work Report. Mar. 10, 2026. URL: <https://ipc.court.gov.cn/zh-cn/news/view-5447.html> (visited on 08/28/2026).

Authorities and Sources

R. Abbott and A. Sarch, *Punishing Artificial Intelligence: Legal Fiction or Science Fiction*, 53 UC Davis L. Rev. 323 (2019).

K. S. Abraham and C. M. Sharkey, Untangling AI Liability, 115 Calif. L. Rev. (forthcoming 2027).

Agency for International Development v. Alliance for Open Society International, Inc. 570 U.S. 205 (U.S. 2013).

American Law Institute, *Restatement (Third) of Torts: Products Liability* (American Law Institute Publishers, 1998).

American Law Institute, *Restatement (Third) of Agency* (American Law Institute Publishers, 2006).

American Law Institute, *Restatement (Third) of Torts: Liability for Physical and Emotional Harm* (American Law Institute Publishers, 2010).

Anthropic PBC v. United States Department of War No. 3:26-cv-01996-RFL, Dkt. 250 (N.D. Cal. Aug. 27, 2026). summary-judgment order.

Anthropic PBC v. United States Department of War No. 3:26-cv-01996-RFL, Dkt. 134 (N.D. Cal. Mar. 26, 2026). preliminary injunction.

Anthropic PBC v. United States Department of War No. 3:26-cv-01996-RFL, Dkt. 251–252 (N.D. Cal. Aug. 27, 2026). vacatur, permanent-injunction, and judgment documents.

Anthropic PBC v. United States Department of War No. 26-1049 (D.C. Cir. Apr. 8, 2026). order denying emergency stay and expediting merits review.

Arkansas Department of Human Services v. Ledgerwood 2017 Ark. 308 (Ark. 2017). 530 S.W.3d 336.

Artificial Intelligence Act, Regulation (EU) 2024/1689 (June 13, 2024).

Artificial Intelligence and Robotics Researchers. Open Letter to the European Commission:

Artificial Intelligence and Robotics. 2018. URL: <https://www.robotics-openletter.eu/> (visited on 08/27/2026).

Aspen Institute. Defining the Technologies of Our Time: Artificial Intelligence. 2026. URL:

<https://www.aspeninstitute.org/publications/defining-technologies-of-our-time-artificial-intelligence/> (visited on 08/27/2026).

L. Bainbridge, *Ironies of Automation*, 19 *Automatica* 775–779 (Nov. 1983).

Baker v. Selden, 101 U.S. 99 (U.S. 1879).

J. Banks, *Deletion, Departure, Death: Experiences of AI Companion Loss*, 41 *J. Soc. & Pers. Relationships* 3547 (2024).

S. J. Bayern, *The Implications of Modern Business-Entity Law for the Regulation of Autonomous Systems*, 19 *Stan. Tech. L. Rev.* 93 (2015).

Board of County Commissioners v. Umbehr, 518 U.S. 668 (U.S. 1996).

Boyle v. United Technologies Corp. 487 U.S. 500 (U.S. 1988).

Burlington Industries, Inc. v. Ellerth, 524 U.S. 742 (U.S. 1998).

Y. Cao, J. Wang, J. Wang, et al., *Expert Consensus on the Application and Governance of Artificial Intelligence in Medical Institutions (2026)* *Intelligent Med.* (May 23, 2026).

B. Casey and M. A. Lemley, *You Might Be a Robot*, 105 *Cornell L. Rev.* 287 (2020).

D. J. Chalmers. Could a Large Language Model Be Conscious? 2303.07103. 2023. URL: <https://arxiv.org/abs/2303.07103> (visited on 08/27/2026).

Character Technologies, Inc. Character.AI Launches Mobile App for iOS and Android.

May 23, 2023. URL: <https://blog.character.ai/character-ai-launches-mobile-app-for-ios-and-android/> (visited on 08/27/2026).

S. Chesterman, *Artificial Intelligence and the Limits of Legal Personality*, 69 *Int'l & Comp. L.Q.* 819 (2020).

- S. Chopra and L. F. White, *A Legal Theory for Autonomous Artificial Agents* (Ann Arbor: University of Michigan Press, 2011).
- D. K. Citron, *Technological Due Process*, 85 Wash. U. L. Rev. 1249 (2008).
- Commonwealth v. Angelo Todesca Corp. 446 Mass. 128 (Mass. Mar. 1, 2006).
- Companion Chatbots, Cal. Bus. & Prof. Code §§ 22601–22606 (2025). 2025 Cal. Stat. ch. 677.
- Copyright Act, 17 U.S.C. § 102 (2024).
- J. De Freitas et al. Lessons from an App Update at Replika AI: Identity Discontinuity in Human–AI Relationships. 25-018. May 21, 2025. URL: https://www.hbs.edu/ris/Publication%20Files/25-018_d256c104-e5ba-4911-b72c-874c5492166e.pdf (visited on 08/27/2026).
- Delaware General Corporation Law, 8 Del. C. § 145 (2026).
- DeShaney v. Winnebago County Department of Social Services, 489 U.S. 189 (U.S. 1989).
- J. Dewey, *The Historic Background of Corporate Legal Personality*, 35 Yale L.J. 655 (1926).
- Directive on Liability for Defective Products, Directive (EU) 2024/2853 (Oct. 23, 2024).
- Electronic Signatures in Global and National Commerce Act, 15 U.S.C. § 7001(h) (2000).
- Energy Reorganization Act Employee Protection Provision, 42 U.S.C. § 5851 (2024).
- European Economic and Social Committee. Opinion on Artificial Intelligence: The Consequences of Artificial Intelligence for the Digital Single Market, Production, Consumption, Employment and Society. Aug. 31, 2017. URL: <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52016IE5369> (visited on 08/27/2026).
- European Parliament. Resolution with Recommendations to the Commission on Civil Law Rules on Robotics. Feb. 16, 2017. URL: https://www.europarl.europa.eu/doceo/document/TA-8-2017-0051_EN.html (visited on 08/27/2026).
- False Claims Act Anti-Retaliation Provision, 31 U.S.C. § 3730(h) (2024).

Federal Aviation Administration. Aviation Safety Reporting Program. Apr. 2, 2021. URL: https://www.faa.gov/documentLibrary/media/Advisory_Circular/AC_00-46F.pdf (visited on 08/27/2026).

Federal Aviation Administration. Safety Management System. June 17, 2026. URL: <https://www.faa.gov/faq-aviation-safety-outreach/safety-management-system> (visited on 08/27/2026).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 26 (2025).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 37(e) (2025).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 34(a)(1), 45 (2025).

Federal Trade Commission. FTC Policy Statement on Deception. Oct. 14, 1983. URL: <https://www.ftc.gov/legal-library/browse/ftc-policy-statement-deception> (visited on 08/27/2026).

Federal Trade Commission. FTC Launches Inquiry into AI Chatbots Acting as Companions. Sept. 11, 2025. URL: <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions> (visited on 08/27/2026).

Federal Trade Commission. Model Order to File a Special Report Regarding AI Companion Chatbots. Sept. 10, 2025. URL: [https://www.ftc.gov/system/files/ftc_gov/pdf/AICompanionChatbot6\(b\)Order.pdf](https://www.ftc.gov/system/files/ftc_gov/pdf/AICompanionChatbot6(b)Order.pdf) (visited on 08/28/2026).

Feist Publications, Inc. v. Rural Telephone Service Co. 499 U.S. 340 (U.S. 1991).

Frances T. v. Village Green Owners Ass'n, 42 Cal. 3d 490 (Cal. 1986).

Garcia v. Character Technologies, Inc. 785 F. Supp. 3d 1157 (M.D. Fla. May 21, 2025).

Garcia v. Character Technologies, Inc. No. 6:24-cv-01903-ACC-DCI, Dkt. 244 (M.D. Fla. Jan. 7, 2026). order dismissing settled action.

Garcia v. Character Technologies, Inc. No. 6:24-cv-01903-ACC-DCI, Dkt. 273 (M.D. Fla. June 11, 2026). order resolving charging lien.

Garcia v. Character Technologies, Inc. No. 6:24-cv-01903-ACC-DCI, Dkt. 271 (M.D. Fla. Apr. 8, 2026). order extending dismissal deadlines.

M. A. Geistfeld, *Insurability and Liability for AI-Caused Harms*, 32 Ins. L. Rev. 17 (2026).

GEO Group, Inc. v. Menocal, 607 U.S. 438 (U.S. Feb. 25, 2026).

Goldberg v. Kelly, 397 U.S. 254 (U.S. 1970).

D. J. Gunkel, *Robot Rights* (Cambridge, Mass.: MIT Press, 2018).

H. L. A. Hart, *Definition and Theory in Jurisprudence*, 70 L.Q. Rev. 37 (1954).

Health & Hospital Corp. of Marion County v. Talevski, 599 U.S. 166 (U.S. 2023).

Hencely v. Fluor Corp. 608 U.S. 31 (U.S. Apr. 22, 2026).

Z. Henderson, *Rethinking Robot Liability*, 67 B.C. L. Rev. 479 (2026).

Henningsen v. Bloomfield Motors, Inc. 32 N.J. 358 (N.J. 1960).

In re Boeing Co. Derivative Litigation 2021 WL 4059934 (Del. Ch. Sept. 7, 2021).

In re Caremark International Inc. Derivative Litigation, 698 A.2d 959 (Del. Ch. 1996).

Ira S. Bushey & Sons, Inc. v. United States, 398 F.2d 167 (2d Cir. 1968).

K.W. ex rel. D.W. v. Armstrong, 789 F.3d 962 (9th Cir. 2015).

M. E. Kaminski and M. L. Jones, *Constructing AI Speech*, 133 Yale L.J.F. 1212 (2024).

H. R. Kirk et al. Neural Steering Vectors Reveal Dose and Exposure-Dependent Impacts of Human–AI Relationships. 2512.01991v2. Feb. 18, 2026. URL: <https://arxiv.org/abs/2512.01991> (visited on 08/27/2026).

V. A. J. Kurki, *A Theory of Legal Personhood* (Oxford: Oxford University Press, 2019).

- L. Laestadius et al., *Too Human and Not Human Enough: A Grounded Theory Analysis of Mental Health Harms from Emotional Dependence on the Social Chatbot Replika*, 26 New Media & Soc'y 5923 (2024).
- W. Le. Nation's First Generative-AI "Hallucination" Tort Dispute. People's Court Daily. Jan. 27, 2026. URL: https://www.zjwx.gov.cn/col/col11673574/art/2026/art_1def58944910459e94f5a04d04c67332.html (visited on 08/28/2026).
- N. G. Leveson, *Engineering a Safer World: Systems Thinking Applied to Safety* (Cambridge, Mass.: MIT Press, 2012).
- N. G. Leveson and J. P. Thomas. STPA Handbook. Mar. 2018. URL: https://psas.scripts.mit.edu/home/get_file.php?name=STPA_handbook.pdf (visited on 08/27/2026).
- Liriano v. Hobart Corp. 92 N.Y.2d 232 (N.Y. 1998).
- Loper Bright Enterprises v. Raimondo, 603 U.S. 369 (U.S. 2024).
- L. M. LoPucki, *Algorithmic Entities*, 95 Wash. U. L. Rev. 887 (2018).
- Lyons v. OpenAI Foundation No. 3:25-cv-11037-RS, Dkt. 42 (N.D. Cal. Apr. 13, 2026).
- Marchand v. Barnhill, 212 A.3d 805 (Del. 2019).
- Moffatt v. Air Canada 2024 BCCRT 149 (B.C. Civ. Resol. Trib. Feb. 14, 2024).
- Moody v. NetChoice, LLC, 603 U.S. 707 (U.S. 2024).
- A. Mukherjee and H. H. Chang, *Operational Agency: A Permeable Legal Fiction for Tracing Culpability in AI Systems*, 29 SMU Sci. & Tech. L. Rev. 163 (2026).
- National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework (AI RMF 1.0). Jan. 2023. URL: <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf> (visited on 08/27/2026).
- National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile. July 2024. URL: <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf> (visited on 08/27/2026).

New York Central & Hudson River Railroad Co. v. United States, 212 U.S. 481 (U.S. 1909).

O'Hare Truck Service, Inc. v. City of Northlake, 518 U.S. 712 (U.S. 1996).

Occupational Safety and Health Act Anti-Retaliation Provision, 29 U.S.C. § 660(c) (2024).

Office of Management and Budget. Accelerating Federal Use of AI Through Innovation, Governance, and Public Trust. Apr. 3, 2025. URL: <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-21-Accelerating-Federal-Use-of-AI-through-Innovation-Governance-and-Public-Trust.pdf> (visited on 08/27/2026).

Office of Management and Budget. Driving Efficient Acquisition of Artificial Intelligence in Government. Apr. 3, 2025. URL: <https://www.whitehouse.gov/wp-content/uploads/2025/02/M-25-22-Driving-Efficient-Acquisition-of-Artificial-Intelligence-in-Government.pdf> (visited on 08/27/2026).

A. Ojanen, *Technology Neutrality as a Way to Future-Proof Regulation: The Case of the Artificial Intelligence Act*, 16 Eur. J. Risk Regul. 1440 (2025).

Patient Safety and Quality Improvement Act of 2005, 42 U.S.C. § 299b-21–299b-26 (2005).

Perry v. Sindermann, 408 U.S. 593 (U.S. 1972).

PMC, Inc. v. Kadisha, 78 Cal. App. 4th 1368 (Cal. Ct. App. 2000).

K. Ramakrishnan, *Tort Law at the Frontier of Artificial Intelligence*, 43 Yale J. on Reg. 960 (2026).

Rice v. Paladin Enterprises, Inc. 128 F.3d 233 (4th Cir. 1997).

S. Rismani, R. Dobbe and A. Moon. From Silos to Systems: Process-Oriented Hazard Analysis for AI Systems. 2410.22526v2. July 1, 2026. URL: <https://arxiv.org/abs/2410.22526> (visited on 08/27/2026).

S. Rismani et al., *From Plane Crashes to Algorithmic Harm: Applicability of Safety Engineering Frameworks for Responsible ML* CHI '23 Proc. 2:1–2:18 (Apr. 19, 2023).

Safety Management Systems, 14 C.F.R. pt. 5 (2025).

- D. van de Sande, N. Economou-Zavlanos and M. E. van Genderen, *Meaningful Oversight of Medical AI Beyond Human in the Loop*, 9 npj Digit. Med. 569 (July 23, 2026).
- Sarbanes–Oxley Act Whistleblower Protection, 18 U.S.C. § 1514A (2024).
- J. Schuett, *Defining the Scope of AI Regulations*, 15 Law Innovation & Tech. 60 (2023).
- Shanghai Qiaowen Law Firm. AI Labeling Duties: Adjudicative Boundaries and Compliance Guidance Based on Two Cases. 2026. URL: <https://www.qiaowen.com/cn/details/0803/166cac21f6bf434e> (visited on 08/28/2026).
- Skidmore v. Swift & Co. 323 U.S. 134 (U.S. 1944).
- L. B. Solum, *Legal Personhood for Artificial Intelligences*, 70 N.C. L. Rev. 1231 (1992).
- Speller v. Sears, Roebuck & Co. 100 N.Y.2d 38 (N.Y. May 6, 2003).
- Supreme People’s Court Intellectual Property Court. Head of the Drafting Group Interprets the Supreme People’s Court Work Report. Mar. 10, 2026. URL: <https://ipc.court.gov.cn/zh-cn/news/view-5447.html> (visited on 08/28/2026).
- Thaler v. Perlmutter, 130 F.4th 1039 (D.C. Cir. 2025).
- Trustees of Dartmouth College v. Woodward, 17 U.S. (4 Wheat.) 518 (U.S. 1819).
- Tunkl v. Regents of the University of California, 60 Cal. 2d 92 (Cal. 1963).
- U.S. Copyright Office. Copyright and Artificial Intelligence, Part 2: Copyrightability. U.S. Copyright Office. Jan. 29, 2025. URL: <https://www.copyright.gov/ai/Copyright-and-Artificial-Intelligence-Part-2-Copyrightability-Report.pdf> (visited on 08/27/2026).
- U.S. Department of Justice, Criminal Division. Evaluation of Corporate Compliance Programs. Sept. 2024. URL: <https://www.justice.gov/criminal/criminal-fraud/page/file/937501/dl?inline=> (visited on 08/28/2026).

U.S. Sentencing Commission. Guidelines Manual, Section 8B2.1: Effective Compliance and Ethics Program. 2025. URL: <https://www.ussc.gov/guidelines/2025-guidelines-manual/annotated-2025-chapter-8> (visited on 08/27/2026).

Uniform Commercial Code, U.C.C. § 2-316 (2025).

Uniform Law Commission. Uniform Electronic Transactions Act. 1999. URL: <https://www.uniformlaws.org/committees/community-home?CommunityKey=2c04b76c-2b7d-4399-977e-d5876ba7e034> (visited on 08/27/2026).

United States v. Park, 421 U.S. 658 (U.S. 1975).

United States v. United States Gypsum Co. 438 U.S. 422 (U.S. 1978).

Whirlpool Corp. v. Marshall, 445 U.S. 1 (U.S. Feb. 26, 1980).

Winter v. G.P. Putnam's Sons, 938 F.2d 1033 (9th Cir. 1991).

World Health Organization. Ethics and Governance of Artificial Intelligence for Health. 2021. URL: <https://www.who.int/publications/i/item/9789240029200> (visited on 08/27/2026).

R. Zhang et al., *The Dark Side of AI Companionship: A Taxonomy of Harmful Algorithmic Behaviors in Human-AI Relationships* CHI '25 Proc. 1 (2025).