

Before the Merits: A Deployment Front Door for High-Impact AI Litigation

Anonymous

Submission Draft

Abstract

AI litigation can stall before substantive doctrine can operate. A claimant may observe an output or decision while the facts needed to identify the responsible actor—the model and application versions, system instructions, tool permissions, evaluations, complaints, escalation decisions, and stop authority—are divided among enterprises and unavailable until discovery. Discovery, however, ordinarily presupposes an identifiable party or nonparty target and cannot produce a risk record that the deployment never created.

This Article identifies that pre-merits attribution gap and proposes a deployment front-door rule. For a statutorily designated high-impact deployment, a claimant who alleges a legally cognizable injury or existing equitable entitlement and a plausible, risk-specific nexus may serve the visible operator. The operator must identify the deployed system and relevant entities, preserve the corresponding record, and provide staged, proportionate, nonprivileged access. An upstream or downstream enterprise enters the same process only when it possessed practical authority materially connected to the asserted risk. A merits firewall prevents participation from supplying any presumption of breach, defect, causation, damages, or joint liability; an enterprise without material control completes the first answer and terminates its special front-door role.

An ex ante deployment-record duty makes the rule usable by assigning risk ownership, approval, escalation, versioning, incident retention, and stop authority before harm

occurs. A deployment-control cascade specifies the records relevant to authorization, notice, feasibility, causal sequence, and defense. Applications to *Garcia v. Character Technologies* and *Mobley v. Workday* show how the same rule opens an integrated consumer deployment and a divided employment stack while leaving ordinary law to decide who wins.

Keywords: artificial intelligence litigation; deployment front door; high-impact AI; early disclosure; deployment records; material control

Contents

Introduction	3
I The Pre-Merits Attribution Gap	8
A The Output Does Not Reveal the Deployment	9
B Two Coordinates Define the Gap	11
C Ordinary Procedure Opens After Four Predicates	12
D The Cascade Specifies the Missing Record	16
II The Deployment Front-Door Rule	20
A Sectoral Designation Defines the Field	20
B The Visible Operator Answers First	22
C Material Control Requires Risk, Authority, and Consequence	24
D Creation, Preservation, and Production Are Separate Duties	28
E Protection, Remedies, and Safe Exit	31
F Model Statutory Text	33
G Federal Procedure Implements Rather Than Displaces the Duty	36
III The Rule in Operation	38
A An Integrated Deployment: <i>Garcia</i>	39
B A Divided Deployment: The Vendor Hosts the Data; the Employer Controls It	43
IV The Record Infrastructure	51
A The Charter Is the Deployment Record's Control Sheet	52

B	The Record Has a Standard Interface	54
C	Identity, Feedback, and Event Records Have Different Clocks	56
D	Existing Regimes Demonstrate Administrability	58
E	Accurate Records Require Protected Escalation	60
F	The Record Is Also a Defense Asset	61
V	The Merits Firewall and Administrable Boundaries	62
A	Ordinary Law Receives Better Factual Predicates	63
B	Staged Access Binds Information Cost	65
C	Existing Legal Interests Define the Gate	67
D	Fragmentation Narrows Control Rather Than Eliminating It	69
E	Product and Speech Classification Remain for the Merits	70
	Conclusion	71

Introduction

Fourteen-year-old Sewell Setzer III did not encounter a foundation model in a laboratory. He encountered a deployed relationship. In *Garcia v. Character Technologies, Inc.*, his mother’s complaint alleged that Character.AI combined persistent personas, memory, access rules, warnings, engagement design, and safety controls in a consumer service available to minors. It alleged that Character Technologies’ founders and Google made different contributions of personnel, technology, knowledge, and resources. After months of interaction with a character modeled on a fictional television figure, Sewell died by suicide. His mother brought wrongful-death, negligence, product-defect, failure-to-warn, deceptive-practice, and related claims. At the pleading stage, the district court held that the application was plausibly alleged to be a product and declined, on that record, to classify the generated output as

protected speech. The action later settled without a finding of defect, causation, or liability.¹

The case exposes a problem that survives every possible merits result. The user sees the interface and the output. The facts ordinary law needs lie behind both: which model and application versions ran; who selected the relational objective, system instructions, memory, age gate, crisis response, and engagement metric; what evaluations, complaints, and incident signals arrived; who could update, restrict, roll back, or stop the deployment; and which enterprise retained each record. In a direct-to-consumer service, the operating company may be obvious while its deployment record remains inaccessible. In a divided system, the visible hospital, employer, agency, or application can depend on a model vendor, scoring service, infrastructure provider, and systems integrator that retain different controls. The claimant needs the configuration and control record to test a recognized claim, yet often must identify the record holder before obtaining that record. The merits standard exists; access to the facts that feed it does not.

This Article calls that sequence the *pre-merits attribution gap*. It comprises two related failures. The first is an *actor-identification failure*: an affected person cannot tell which enterprise selected or retained a control relevant to the asserted risk. The second is a *deployment-record failure*: even when the operator's name is visible, the versions, configuration, evaluations, warnings, incidents, and intervention authority needed to test its conduct remain internal or were never recorded. The first failure is most acute in divided deployments. The second can determine a case in any deployment.

Much of the leading AI-liability debate asks a different question: which substantive standard should govern after the relevant actor and conduct are known? Existing negligence,

¹See *Garcia v. Character Technologies, Inc.* 785 F. Supp. 3d 1157, 1165–69, 1174–84 (M.D. Fla. May 21, 2025); *Garcia v. Character Technologies, Inc.* No. 6:24-cv-01903-ACC-DCI, Dkt. 244, 1 (M.D. Fla. Jan. 7, 2026). order dismissing settled action; *Garcia v. Character Technologies, Inc.* No. 6:24-cv-01903-ACC-DCI, Dkt. 273, 1 (M.D. Fla. June 11, 2026). order resolving charging lien. The complaint's disputed account remains an allegation. This Article uses the published Rule 12 order for its classification and pleading analysis, not as a merits judgment.

products, antidiscrimination, consumer-protection, administrative, contract, and no-liability rules may answer many such disputes.² This Article accepts that ordinary-law premise and asks what happens one step earlier. A reasonableness standard cannot evaluate conduct until the relevant conduct and actor are visible. Product classification cannot identify which company selected the model harness or retained the update right. A rule denying relief for a category of loss still requires enough information to determine whether the case falls within that category. The front-end information problem is analytically distinct from the standard that ultimately governs the merits.

The question presented is therefore narrow: when a civil claimant can allege a legally cognizable injury or an existing entitlement to equitable relief connected to a designated high-impact AI deployment, but cannot identify the enterprise that controlled the relevant risk or obtain the deployment record needed to test an ordinary cause of action, what threshold should trigger identification, preservation, and focused production—and which enterprise must answer first?

The Article's contribution is institutional sequencing. Existing AI proposals often begin with a substantive duty, disclosure mandate, regulator, or allocation of liability. Existing procedure offers discovery after a viable action reaches an identifiable target. The front door joins four functions that those approaches leave separate: a sectoral designation that identifies the legal interest; an *ex ante* minimum record that keeps the deployment observable; a visible-operator gateway that makes service possible; and a material-control test that allocates access without allocating liability. The resulting rule is narrow enough to administer and complete enough to open the merits process.

This Article proposes a *deployment front-door rule*. A legislature first designates a deployment category and specifies the protected interests, visible operator, minimum record,

²See K. S. Abraham and C. M. Sharkey, *Untangling AI Liability*, 115 *Calif. L. Rev.* (forthcoming 2027); M. A. Geistfeld, *Insurability and Liability for AI-Caused Harms*, 32 *Ins. L. Rev.* 17 (2026); Z. Henderson, *Rethinking Robot Liability*, 67 *B.C. L. Rev.* 479 (2026).

retention period, trigger, and remedy. Once a claimant alleges a legally cognizable injury or existing equitable entitlement and specific facts supporting a plausible, risk-specific nexus to that deployment, the enterprise that made the deployment available to the affected person or incorporated it into the decision process becomes the disclosure gateway. It must identify the deployed system and relevant entities, preserve the corresponding record, notify enterprises that may hold a material control lane, and provide staged, proportionate, nonprivileged production. Each identified enterprise joins that obligation only if, at the relevant time, it possessed practical authority materially connected to the asserted risk.

The rule changes access to facts, not the elements that decide the case. Its merits firewall creates no presumption of duty, breach, defect, causation, damages, or joint liability. Jurisdiction, immunity, privilege, proportionality, trade-secret protection, and substantive defenses remain operative. An enterprise that completes the first-answer obligation and establishes that it lacked material control over the risk terminates its special front-door role; absent an independently sufficient claim, it returns to ordinary nonparty status. Ultimate allocation remains governed by the cause of action, contribution and indemnity law, contract, and any independently available individual claim.

The front door has an *ex ante* component because litigation cannot recover a record that a deployment never created. For each designated category, the statute should impose a *deployment-record duty*: identify the system and version; assign risk and approval authority; document material evaluations, unresolved risk decisions, changes, complaints, incidents, and interventions; identify who can restrict, roll back, or stop operation; and retain the record for a specified period. What this Article previously called a responsibility charter is best understood as the internal instrument that satisfies that record duty. It is infrastructure for the front door, not a second governance product.

The scope of the required record follows practical control. This Article uses a compressed *deployment-control cascade* to map objectives, resources, configuration, feedback, and stopping power across the model provider, application deployer, and operating insti-

tution. The cascade is a record-specification method. It identifies evidence relevant to authorization, notice, foreseeability, feasible precaution, causal sequence, preservation, and defense; governing law supplies every legal standard. Token-level unpredictability may complicate causal proof, but it does not make organizational choices about purpose, users, data, instructions, tools, monitoring, updates, and shutdown disappear.

Nor does the proposal depend on resolving artificial personhood. In the deployments covered here, existing people and organizations hold the assets, records, authority, service capacity, and exposure to remedies. A legislature may someday construct a bounded juridical platform for an artificial system. Until it does, operational autonomy remains a fact about how the system acts, while positive law determines which recognized person holds the claim, owes the duty, or answers the judgment. The front door points to the entities that can actually identify the deployment and produce its record.

Three objections define the rule's specifications. First, ordinary pleading and discovery often work once a claimant has a proper defendant and an identifiable record holder. The gap lies in establishing those predicates and in creating the record before litigation. Second, identification and production impose real costs even when the enterprise ultimately prevails. Coverage must therefore be sectorally designated; the nexus must be risk specific; production must be staged; and privacy, security, privilege, and trade secrets must be protected. Third, some losses connected to AI have no legal remedy. The rule requires an existing cause of action or equitable entitlement, permits pure legal defenses to be resolved first, and supplies no substantive right to inspect a deployment simply because a claimant invokes AI.

The Article develops one contribution through five Parts. Part I identifies the premerits attribution gap, explains why ordinary procedure presupposes facts it cannot always reveal, and compresses the deployment-control cascade into a tool for specifying the missing record. Part II states the front-door rule, defines sectoral coverage and material control, separates record creation from preservation and production, supplies model statutory text, and coordinates state legislation with federal procedure. Part III applies the rule to two comple-

mentary settings: *Garcia*, where the visible operator makes production the central problem, and *Mobley*, where hiring decisions, platform features, data hosting, legal control, and validation records divide across employers and a vendor. A completed public-employment record provides a stable analogue. Part IV develops the ex ante record infrastructure and shows how existing safety and information-demand regimes make it administrable. Part V defines the merits firewall, information-cost controls, fragmented-stack boundaries, and classification neutrality. The Conclusion returns to a single proposition: substantive law should decide AI cases, not an information asymmetry that prevents substantive law from starting.

I The Pre-Merits Attribution Gap

A merits rule and a fact-access rule perform different work. Negligence may ask what precautions a reasonable deployer would have taken; products law may ask whether a design or warning was defective; antidiscrimination law may ask whether a screening process caused a prohibited disparity; administrative law may ask whether an agency supplied the explanation and review an entitlement requires. Each inquiry presupposes a defined act, actor, system, and time. Affected people commonly know the adverse result and the institution through which it arrived. They do not ordinarily know the version, configuration, allocation of control, or record needed to connect that result to the governing elements.

This Part identifies the resulting gap in three moves. Section A separates computational opacity from the system and organizational facts that adjudication can reach. Section B shows that ordinary federal procedure becomes effective after a proper target and record holder are identifiable but offers no uniform route to create those predicates. Section C compresses the deployment-control cascade into a specification for the facts the front door must identify and preserve.

A The Output Does Not Reveal the Deployment

An AI dispute should begin with the deployed system rather than the generated token, score, or recommendation. The first attribution question is not who selected the final words or number. It is who constructed, supplied, authorized, observed, and could alter the deployment that produced the legally relevant effect. A model's internal computation may be difficult to explain while the surrounding control structure remains extensively documented in contracts, release records, configuration systems, evaluation reports, incident queues, and access logs. That distinction makes responsibility evidence obtainable even when a compact explanation of the model's internal path is unavailable.

Three forms of opacity should remain separate. *Computational opacity* concerns the inability to give a human-scale account of how internal states produced a particular output. *System opacity* concerns interactions among the model, instructions, retrieval, memory, tools, filters, credentials, and conventional software. *Organizational opacity* concerns who approved the use, possessed risk information, controlled resources, or could restrict and stop operation. The first may follow from architecture. The latter two ordinarily follow from deployment design and recordkeeping.

Collapsing all three into a single “black box” allows the least tractable fact to obscure the facts ordinary doctrine can use. A court may be unable to reconstruct why a model ranked one token above another yet able to determine that a provider selected the model; a deployer supplied protected records and external tools; an interface suppressed uncertainty; a safety team documented a known failure mode; management rejected a delayed release; and an operator lacked time or authority to intervene. Intent and knowledge are often inferred from conduct and circumstances, product defect can sometimes be proved circumstantially, and direct organizational fault can lie in information systems, supervision, and resource

allocation.³ Model interpretability is not a condition precedent to responsibility at the deployment level.

Fragmentation makes those accessible facts difficult to attribute. Consider an employer that licenses an applicant platform. The employer may select the job, criteria, candidate pool, weight assigned to a recommendation, and final disposition. A platform vendor may select the ranking model, training process, default features, update cadence, and bias testing. A systems integrator may configure data fields and workflow triggers. The applicant encounters an employer-branded portal and a rejection, not this allocation. In *Mobley v. Workday, Inc.*, for example, a complaint alleged that employer customers delegated candidate screening to Workday’s AI-based tools; at the Rule 12 stage, the district court treated the alleged delegation of a traditional hiring function as sufficient to proceed on a statutory agency theory while leaving the allegations and ultimate liability unresolved.⁴ The opinion shows that substantive law may address a divided stack once the vendor and its alleged function are visible. It does not supply every rejected applicant with the identity, version, configuration, or evaluation record needed to reach that point.

The same problem appears when the visible institution itself cannot inspect the vendor’s method. A public employer may control whether and how a score affects termination while a vendor controls the scoring software and updates. A hospital may select the patient population and integrate a recommendation while a supplier controls model validation and releases. A benefits agency may own the legal decision and appeal path while contractors control the calculation tool and documentation. Control is distributed by function, not transferred wholesale to whichever entity’s logo appears on the screen.

³See, e.g., *United States v. United States Gypsum Co.* 438 U.S. 422, 444–46 (U.S. 1978); *Speller v. Sears, Roebuck & Co.* 100 N.Y.2d 38, 41–43 (N.Y. May 6, 2003); American Law Institute, *Restatement (Third) of Agency* (American Law Institute Publishers, 2006), §§ 7.03(1), 7.05.

⁴See *Mobley v. Workday, Inc.* 740 F. Supp. 3d 796, 806–08 (N.D. Cal. July 12, 2024). The decision granted in part and denied in part Workday’s motion to dismiss the first amended complaint; it did not find discriminatory operation or resolve the merits.

Current legal subjects remain stable across that distribution. Operational autonomy describes how a system produces effects; it does not identify an entity with assets, records, service capacity, or exposure to judgment. In the covered deployments considered here, existing organizations hold those incidents. A future legislature can construct a bounded juridical status and specify its representative, assets, attribution rules, service, and remedies. The front-door rule need not decide that future question. It follows the records and practical authority that existing law can presently reach.

B Two Coordinates Define the Gap

Actor identity and record availability are independent coordinates. Treating them as one generic opacity problem obscures which legal intervention is needed. A claimant may know the relevant operator yet lack its deployment record. A claimant may possess a transaction or application record yet lack the upstream entity that supplied the operative function. A divided stack can produce both conditions at once. The front door sequences identity and record access so that each condition receives the smallest useful response.

The first configuration is an identifiable integrated operator with an adequate record. Ordinary procedure has a direct target, and the statute's principal work occurs *ex ante*: standardization, retention, and a quick first answer. The record may establish the active version, relevant policy, warnings, incident sequence, and stop authority without adding another merits defendant. If a filed claim proceeds, ordinary discovery can deepen the record in the familiar way.

The second configuration is an identifiable operator with a missing or fragmented internal record. The brand and legal entity are visible, but the claimant and sometimes the operator cannot reconstruct which versions, components, policies, and intervention roles governed the event. This is the production problem. It arises in an integrated consumer service when engineering, safety, product, trust-and-safety, and executive decisions live in separate systems. It also arises in a public institution that relies on a proprietary score it cannot independently verify. The *ex ante* duty matters most here because a post-event

subpoena cannot create lineage or approval that the operator never recorded.

The third configuration is an identifiable event with an unknown operative entity. An applicant may have an employer, job, timestamp, and rejection while lacking the product feature that processed the application and the supplier responsible for it. A patient may know the hospital encounter while lacking the vendor and version embedded in a clinical workflow. This is the actor-identification problem. A closed first answer resolves provenance and custodian fields before broader access. The visible operator performs the gateway role because it chose to incorporate the system and can preserve the event even if another enterprise holds the relevant control lane.

The fourth configuration combines the two. The operative feature and supplier are hidden, and the minimum record was never created or is divided by contract. A supplier may host data while the customer owns it; an integrator may hold configuration history while the institution holds the decision; a model provider may know version lineage while the application company knows prompts, tools, and user exposure. Asking the claimant to choose a single controller before revealing this map converts organizational design into a litigation advantage. Treating every participant as a controller creates the opposite error. The front door identifies lanes first and terminates the special role of actors outside them.

This two-coordinate account yields a fixed sequence. The visible operator maps the deployment and event. Identified enterprises acknowledge their control and record lanes. All holders preserve the bounded slice. Stage One supplies version, provenance, control, and event fields. The material-control test removes irrelevant actors. Stage Two follows only the legal theory and unresolved lane. The sequence is the Article's central payoff: it opens access without deciding the liability questions that the record is meant to illuminate.

C Ordinary Procedure Opens After Four Predicates

Federal procedure supplies substantial tools after four predicates are present: a filed action, a complaint capable of stating a legal claim, a known party or subpoena recipient, and an existing record within the reach of that target. The deployment front door supplies

information needed to establish those predicates; it does not replace the procedure that follows them.

Pleading law illustrates the sequence. Rule 8 requires a short and plain statement showing entitlement to relief. *Twombly* and *Iqbal* require factual matter supporting a plausible claim, while Rule 11(b)(3) permits a lawyer to identify factual contentions that are likely to obtain evidentiary support after a reasonable opportunity for investigation or discovery.⁵ That accommodation matters: a plaintiff need not possess every internal email or validation report before filing. But Rule 11 creates neither a discovery entitlement nor a substitute for Rule 8. Neither rule names an unknown supplier, links a rejection to a product feature, identifies the version that ran, or requires the operator to have recorded that fact.

Discovery becomes powerful once an action and target exist. Rule 26(b)(1) permits discovery of relevant, proportional, nonprivileged matter. Rule 34 reaches existing documents and electronically stored information in a party's possession, custody, or control. Rule 45 reaches a known nonparty that can be named and served.⁶ These mechanisms can test an agency theory, obtain validation evidence, and expose a divided contract. They do not identify the person to whom a subpoena should be directed or create a release approval, configuration history, or incident record that never existed.

Initial disclosure begins later and serves a narrower function. Rule 26(a)(1) generally requires the parties to identify witnesses and materials they may use to support their own claims or defenses, along with damages and insurance information. The disclosures ordinarily follow the Rule 26(f) conference. They are not an adverse-record inventory. An operator that does not expect to use a rejected mitigation, contradictory test, or internal escalation in support of its position has no initial-disclosure obligation to list the item on that basis

⁵See Federal Rules of Civil Procedure, Fed. R. Civ. P. § 8(a)(2) (2025); Federal Rules of Civil Procedure, Fed. R. Civ. P. § 11(b)(3) (2025); *Bell Atlantic Corp. v. Twombly*, 550 U.S. 544, 555–59, 570 (U.S. 2007); *Ashcroft v. Iqbal*, 556 U.S. 662, 678–79, 686 (U.S. 2009); *Rotella v. Wood*, 528 U.S. 549, 560–61 (U.S. 2000).

⁶See Federal Rules of Civil Procedure, Fed. R. Civ. P. § 26 (2025); Federal Rules of Civil Procedure, Fed. R. Civ. P. § 34(a)(1), 45 (2025).

alone. The claimant can request it in discovery after the case begins; the front-door problem is how to identify the system, holder, and record category needed to frame that request.

Court-authorized early discovery provides a real bridge when the claimant already has a unique event identifier and knows the intermediary holding the identity record. In *Strike 3 Holdings, LLC v. Doe*, a rights holder tied an alleged online infringement to an IP address and sought an early subpoena to the known internet service provider capable of mapping that address to a subscriber. The D.C. Circuit held that Rule 26(d)(1) permits such a route and that Doe pleading can be appropriate when discovery is expected to reveal identity, subject to relevance, proportionality, and a realistic prospect of jurisdiction and success.⁷ The case identifies the bridge's predicates as clearly as its power: a viable substantive claim, a unique incident identifier, and a known intermediary with an existing identity record. A fragmented deployment may lack the second or third predicate because the applicant sees only a branded portal, the operator does not retain the feature or run identifier, or the intermediary itself is hidden behind an integrator.

Rule 27 occupies another useful but bounded lane. It permits a person expecting to become a party to perpetuate testimony that risks being lost before an action can be filed. The petition must state the facts the petitioner expects to establish and the substance of the testimony to be preserved. Courts therefore distinguish perpetuation of known evidence from an investigation designed to learn whether a cause of action exists or whom to sue. In *Workman v. United States Postal Service*, the Tenth Circuit affirmed denial of a petition that sought testimony to determine whether a postal-facility fire supported a claim; lack of the evidence needed to frame a complaint was not the inability to bring an action contemplated by Rule 27.⁸ A deposition cannot supply an unknown model version or a deployment record

⁷See *Strike 3 Holdings, LLC v. Doe*, 964 F.3d 1203, 1207, 1210–11 (D.C. Cir. 2020); see also *Columbia Insurance Co. v. Seescandy.com*, 185 F.R.D. 573, 577–80 (N.D. Cal. 1999).

⁸See Federal Rules of Civil Procedure, Fed. R. Civ. P. § 27 (2025); *Workman v. United States Postal Service*, 127 F.4th 237, 240–45 (10th Cir. 2025); *Ash v. Cort*, 512 F.2d 909, 912 (3d Cir. 1975).

that no witness or enterprise created.

Divided control creates a further limit within ordinary discovery. Rule 34 control is a legal and organizational inquiry, not a synonym for technical hosting. The Ninth Circuit has declined to treat cooperative affiliation as control when the responding entity lacked a legal right to obtain an affiliate's files on demand. The Eleventh Circuit, in a different statutory proceeding and factual record, found control where affiliated entities regularly shared responsive information and the recipient had the practical ability to obtain it.⁹ Those outcomes properly depend on contract and practice. They also show why a visible operator should retain the minimum record locally and obtain a contractual right to retrieve defined provenance and event fields from a supplier. Otherwise the vendor may physically host the data while the customer legally controls it, leaving each able to identify the other as the proper source.

Doe pleading and later amendment preserve some claims but do not create a national limitations safety net. Rule 15(c)(1)(C) requires timely notice to the added party and knowledge that, but for a mistake concerning identity, the action would have been brought against it. *Krupski* treats the prospective defendant's knowledge as central and recognizes that confusion over the roles of related, known entities can constitute a mistake. Courts can distinguish that situation from complete ignorance of an upstream actor. In *Winzer*, the Fifth Circuit held that substituting named officers for Doe defendants after limitations expired did not correct a Rule 15 mistake. Rule 15(c)(1)(A) can incorporate a more forgiving state relation-back rule, but that route depends on the law governing limitations and the state's own conditions.¹⁰ A deployment statute can address the timing problem directly

⁹Compare *In re Citric Acid Litigation*, 191 F.3d 1090, 1107–08 (9th Cir. 1999), with *Sergeeva v. Tripleton International Ltd.* 834 F.3d 1194, 1200–02 (11th Cir. 2016). *Sergeeva* arose under 28 U.S.C. § 1782; it illustrates a practical-control approach rather than deciding ordinary domestic merits discovery.

¹⁰See Federal Rules of Civil Procedure, Fed. R. Civ. P. § 15(c) (2025); *Krupski v. Costa Crociere S.p.A.*, 560 U.S. 538, 547–57 (U.S. 2010); *Winzer v. Kaufman County*, 916 F.3d 464, 470–71 (5th Cir. 2019); *Butler v. National Community Renaissance*, 766 F.3d 1191, 1200–02 (9th Cir. 2014).

for claims within the enacting legislature’s authority by tolling the applicable period while a compliant identity request is pending. Amendment and relation back in a later federal action remain governed by Rule 15.

State pre-suit proceedings confirm that a narrow identity stage is administrable while revealing the absence of a uniform route. Illinois Rule 224 authorizes an independent verified action limited to identifying a potentially responsible person or entity and ordinarily ends the proceeding after identification. Texas Rule 202 permits a verified petition to investigate a potential claim when the likely benefit outweighs burden or when perpetuation may prevent a failure or delay of justice. New York authorizes disclosure to aid in bringing an action. California’s perpetuation statute takes the opposite approach: it expressly excludes discovery to determine whether a claim exists or to identify parties.¹¹ These systems provide design elements—verification, a closed purpose, burden review, a short proceeding, and cost allocation—rather than an existing federal solution.

The doctrinal sequence thus defines the proposal’s modest but consequential job. Ordinary discovery can test the merits once the litigation gateway is open. The deployment front door supplies the identity, minimum record, and control map needed to reach that gateway, and it makes those facts available before limitations, deletion, and divided custody determine the case’s trajectory.

D The Cascade Specifies the Missing Record

The deployment-control cascade is a disciplined way to define those predicates. Adapted from systems-safety analysis, it represents an AI deployment as linked controllers, control actions, process models, feedback channels, and intervention capacity. Systems-theoretic methods

¹¹See Illinois Supreme Court Rules, Ill. Sup. Ct. R. § 224 (2026); Texas Rules of Civil Procedure, Tex. R. Civ. P. § 202 (2026); New York Civil Practice Law and Rules, N.Y. C.P.L.R. § 3102(c) (2026); California Code of Civil Procedure, Cal. Civ. Proc. Code § 2035.010 (2026). A New York pre-action petition does not become a federal Rule 27 proceeding simply because a respondent removes it; *Teamsters Local 404* treated the state petition as outside the removable civil action described by 28 U.S.C. § 1441. See *Teamsters Local 404 Health Services & Insurance Plan v. King Pharmaceuticals, Inc.* 906 F.3d 260, 264–69 (2d Cir. 2018).

treat loss as a control problem that can arise from an absent constraint, unsafe action, inadequate feedback, or a controller’s mistaken model of the operating environment.¹² The cascade uses that vocabulary only to specify contestable evidence. Governing law assigns duty, fault, causation, defenses, and remedy.

Four functional levels capture most deployments:

1. *authority and resources*: the actors who set objectives, risk tolerance, staffing, compute, data access, and release conditions;
2. *model and component supply*: the actors who develop or select models, document limits, conduct evaluations, issue updates, and supply interfaces or weights;
3. *deployment and harness operation*: the actors who select instructions, retrieval, memory, filters, routing, tools, credentials, eligible users, monitoring, and escalation; and
4. *the affected environment*: operators, users, data subjects, third-party systems, and institutional conditions that supply inputs and experience effects.

These are roles rather than four firms. One enterprise can occupy the first three levels in an integrated service. An API deployment may divide them among a model vendor, application developer, enterprise customer, and professional user. A public agency can set the legal objective and use of a score while a contractor supplies and updates the model. The map prevents a corporate label or contract caption from substituting for actual function.

At each level, six questions define the discoverable record. What asserted risk was within scope? Which action could create, expose, detect, or reduce it? What did the controller understand about the system and environment? Which feedback could correct

¹²See N. G. Leveson, *Engineering a Safer World: Systems Thinking Applied to Safety* (Cambridge, Mass.: MIT Press, 2012); N. G. Leveson and J. P. Thomas. STPA Handbook. Mar. 2018. URL: https://psas.scripts.mit.edu/home/get_file.php?name=STPA_handbook.pdf (visited on 08/27/2026); S. Rismeni et al., *From Plane Crashes to Algorithmic Harm: Applicability of Safety Engineering Frameworks for Responsible ML* CHI ’23 Proc. 2:1–2:18 (Apr. 19, 2023).




Figure 1: The deployment-control cascade as a record specification. Solid arrows carry authority, resources, components, and operational actions toward the affected environment. Dashed arrows carry the feedback required to update each controller’s understanding. Restriction, rollback, credential revocation, and shutdown must be recorded where they can be exercised.

that understanding? What practical authority and resources existed at the relevant time? Who could restrict, roll back, suspend, or terminate operation? The answers identify specific materials rather than an undifferentiated demand for “the algorithm.”

Table 1: From Deployment Control to a Litigation Record

Control lane	Core record	Potential use under governing law
Identity and configuration	Model and application versions; instructions; retrieval; data; tools; credentials; material changes.	Attribution; intended use; causal mechanism; authentication.

Control lane	Core record	Potential use under governing law
Authority and approval	Risk owner; release approvers; objectives; resources; conditions; unresolved risk decisions.	Authorization; direct organizational conduct; feasibility inquiry.
Evaluation and feedback	Test protocols and results; complaints; incidents; drift signals; overrides; escalation and response.	Knowledge; notice; continuing response; alternative explanations.
Intervention	Restriction, update, rollback, suspension, and termination rights; tested triggers and response times.	Practical control; ability to prevent or mitigate; causal sequence.
Retention and provenance	Record custodian; retention period; version lineage; third-party holders; legal holds and deletion events.	Identification; preservation; proportional discovery; defense and rebuttal.

The table is deliberately factual. Shutdown credentials supply evidence of practical authority; an alert supplies evidence of receipt; an internal limit supplies evidence about the organization's selected constraint. Governing law decides whether those facts establish duty, notice, or reasonable care. The same record supplies defense-favorable facts: a downstream modifier removed a safeguard, a user defeated a tested constraint, an incident signal never reached the asserted controller, or the proposed intervention could not have changed the result.

This mapping yields the front door's object. Actor-identification asks which enterprise

occupied a risk-specific control lane. Deployment-record production asks for the bounded materials that establish or rebut that proposition and then permit ordinary doctrine to operate. Part II turns those two functions into a statutory trigger, a first-answer obligation, and a safe exit.

II The Deployment Front-Door Rule

The attribution gap calls for a limited information rule rather than a new general law of AI liability. The rule must tell an affected person whom to approach, tell enterprises what record to maintain, and tell courts when and how that record becomes available. It must also provide a prompt exit for an enterprise whose connection to the brand or technical stack did not include practical control over the asserted risk.

This Part supplies that architecture. Section A ties coverage to sectoral legislative designation. Section B makes the visible operator the first-answer gateway and defines a threshold linked to an existing legal claim. Section C gives *material control* a cumulative, risk-specific test. Section D separates record creation, preservation, and staged production. Section E specifies confidentiality, remedies, and safe exit. Section F consolidates the proposal into model statutory text, and Section G explains how an ex ante state-law record duty can operate alongside the Federal Rules.

A Sectoral Designation Defines the Field

“High impact” should identify a deployment category, not attach permanently to a model. A general-purpose model can support casual drafting in one setting and become part of a benefits, hiring, clinical, or physical-control process in another. The deployment supplies the protected interest, affected population, operating institution, and legally relevant control. Coverage should follow those features.

The proposed statute therefore applies only when a legislature has designated a category and specified six matters: the protected interests; the visible operator; the minimum

deployment record; the retention period and material-change trigger; the claimant's threshold; and the available enforcement remedy. A legislature can make the designation directly or authorize a sectoral agency to maintain a bounded list under statutory criteria. The legislature retains the choices that determine who must answer and which interests justify compulsory recordkeeping.

Five features identify the strongest candidates. First, the deployment affects an interest already protected by substantive law, such as employment opportunity, access to public benefits, health, credit, housing, education, physical safety, or the welfare of minors. Second, the system repeatedly mediates decisions or exposure rather than producing an isolated low-consequence convenience. Third, affected people predictably observe less about the operative system than the operator and suppliers know. Fourth, the identity, version, control, and event fields can be standardized at reasonable cost. Fifth, existing doctrine can adjudicate the dispute once those facts are available. These criteria select the settings in which a record gateway unlocks an established legal regime rather than inventing a new one.

Designation also fixes the public-law interface. If a covered operator is a government institution, the legislature should specify the enforcing tribunal, available relief, applicable exhaustion route, and the treatment of sovereign or official immunity. The front door need not imply a damages action against the government. It can supply an administrative record, declaratory remedy, reviewable denial, or other access path tailored to the underlying interest. Defining that route at designation keeps the information right aligned with the sector's remedial structure.

This design avoids an unstable universal definition. Employment screening, allocation of essential public benefits, clinical diagnosis or treatment, credit, housing, education admissions, physical safety, and relational services offered to minors may justify different records and retention periods. A companion service may require age-access, relational-behavior, crisis-intervention, and continuity records. An employment tool may require input-feature, validation, selection-rate, configuration, and version records. A clinical system may require

intended-use, validation-population, uncertainty, escalation, override, and follow-up records. The shared front-door architecture permits sector-specific content.

California’s enacted companion-chatbot provisions demonstrate the feasibility of this form. The legislature defined a covered relational service and its operator, required specified disclosures and self-harm protocols, imposed protections for known minors and reporting obligations, supplied a private remedy, and preserved cumulative law.¹³ The statute did not need to declare every conversational model high impact. It regulated a deployment category and the enterprise that made it available.

A designation should also identify exclusions. Research or internal evaluation that does not produce a legally or practically consequential effect on an affected person ordinarily falls outside. So does general cloud hosting without risk-specific configuration or intervention authority. Small entities may receive simplified record forms or regulator assistance where the sector permits, but the information essential to identify a consequential decision should not disappear with enterprise size. A covered operator using open weights still controls a deployment; an upstream author who has no continuing authority may fall outside the post-release control lane.

B The Visible Operator Answers First

The rule resolves the service problem by selecting the entity an affected person can ordinarily identify. The *visible operator* is the enterprise or public institution that made the covered deployment available to the affected person or incorporated it into the decision process that affected that person. It may be a companion-service provider, employer, hospital, lender, benefits agency, school, or operator of a branded application. Its gateway role follows from visibility and provenance, not a presumption that it controlled every risk.

¹³See Companion Chatbots, Cal. Bus. & Prof. Code §§ 22601–22606 (2025). 2025 Cal. Stat. ch. 677. The enacted provisions do not contain the deployment front-door rule proposed here; they demonstrate that a legislature can designate a relational deployment, identify its operator, and attach enforceable protocol and reporting duties.

The operator receives a verified front-door notice. The notice must identify: (1) a legally cognizable injury or an existing entitlement to equitable relief; (2) the cause of action, statutory right, or protected legal interest the claimant expects to invoke; (3) concrete facts supporting a plausible nexus between the asserted risk and a designated deployment; (4) the relevant time, transaction, account, decision, or interaction with as much specificity as reasonably available; and (5) the limited identity and record categories sought. “An AI harmed me” is insufficient. The facts must connect the deployment to a legal interest and connect the requested record to an element, defense, or remedial question. Verification means a sworn declaration or an unsworn declaration under penalty of perjury based on personal knowledge, with specifically identified matters permitted on information reasonably available to the requester. Counsel separately certifies the request under the forum’s ordinary good-faith standard.

The notice invokes a statutory access right that exists outside merits discovery. It can be served before a complaint, with a complaint, or after filing. Pre-suit service gives the operator an opportunity to answer without motion practice and preserves the claimant’s ability to name the correct parties. Service should follow a designated electronic or registered-agent channel. The operator receives thirty days, subject to one short extension for documented technical or privacy needs, to answer, object by record category, or state that a named deployment did not process the identified event. Silence, a blanket proprietary designation, and referral to an unnamed vendor are not responses.

If the operator refuses or materially underanswers, the claimant may bring a narrow enforcement action against that known entity. The pleaded statutory violation is the qualifying deployment, compliant request, service, and nonperformance; the claimant need not plead the still-concealed merits facts to enforce the access right. The administering agency or a court with an independent basis for jurisdiction can order the defined first answer. The claimant bears the burden to establish coverage and a qualifying notice. The operator bears the burden to support an objection based on nonoperation, burden, privilege, privacy, secu-

ity, or the absence of a requested field. Jurisdiction, immunity, preemption, and a pure legal bar that can be resolved on the request remain available before any access beyond identity and preservation.

The gateway must provide a *first answer*. It identifies the application and model versions active for the relevant event; the upstream and downstream entities that supplied, configured, integrated, or operated material components; the entity that holds each listed record category; and the persons or roles with authority to update, restrict, roll back, suspend, or stop the relevant function. It preserves records within its control, identifies the retention schedule and any known deletion, and gives statutory notice to each enterprise that may have occupied a material control lane. It also provides the incident-specific output, decision, notice, configuration, warning, and intervention events within its custody, subject to privilege and lawful protection of third-party data.

This first answer solves two practical problems. The claimant receives the deployment's identity and provenance without reconstructing the operator's vendor architecture. Potential control holders receive timely notice and can preserve their own records. A visible operator cannot answer that the vendor relationship is confidential or that its procurement team does not know which version ran; maintaining that provenance is part of the ex ante deployment-record duty.

The visible operator remains free to establish that its gateway role did not include material control over the asserted risk. A hospital that only offered physical access to an independently governed research system, a marketplace that processed payment without configuring the application, or an employer whose vendor acted outside any retained or delegated hiring function may present that record. Visibility opens the door; the test in the next Section determines who remains.

C Material Control Requires Risk, Authority, and Consequence

"Ability to affect the system" is too broad. A cloud provider, investor, data vendor, researcher, model author, and downstream institution may each have some causal connection

to a deployment. The front door should reach an enterprise only when three requirements are satisfied at the relevant time.

First, *risk specificity*: the asserted control must connect to the risk alleged in the verified notice. General influence over corporate strategy, compute availability, or model research does not establish control over age access, hiring rank, clinical escalation, a particular tool permission, or another specified hazard.

Second, *practical authority*: the enterprise must have possessed an operational or legally enforceable ability to configure, constrain, evaluate, monitor, update, suspend, roll back, or terminate the relevant function. A theoretical technical capacity, an expired contract right, or influence that required another actor's voluntary cooperation is evidence but not practical authority by itself. Contractual audit, update, safety-override, and suspension rights matter because they may be usable controls; a contractual label declaring the parties independent does not decide the question.

Third, *materiality*: exercising the authority must have been capable of making a substantial difference to the risk's creation, exposure, detection, or reduction. An incidental contribution or remote ability affecting all customers in the same undifferentiated way falls outside. Materiality does not require proof that the intervention would certainly have prevented the claimant's injury. It requires a nonmarginal control lane relevant to the deployment and risk.

The requirements are cumulative. Nonexclusive factors include who selected the deployment's purpose and population; chose the model and version; configured instructions, memory, retrieval, data fields, tools, and credentials; set engagement, ranking, decision, or escalation thresholds; received evaluations, complaints, and incident signals; could issue an update or impose a restriction; retained contractual audit or safety-override rights; and could suspend, roll back, or terminate the function. The inquiry is temporal. Authority obtained after the event does not create earlier control, although later records may reveal what authority existed.

Several negative rules keep the test bounded. Investment, ordinary cloud capacity, publication of research, past employment, trademark association, and model provenance alone do not establish material control. Publishing open weights may end the original provider's ability to monitor or stop an independent deployment. A downstream modifier that removes safeguards, adds credentials, or selects a new high-impact purpose can occupy the relevant lane instead. Conversely, a provider that retains mandatory updates, remote suspension, risk monitoring, or binding deployment conditions may retain a lane despite contract language assigning "all responsibility" downstream.

Material control determines information obligations, not substantive liability. An enterprise can possess a control relevant to feasibility yet owe no duty to the claimant under governing law. It can receive an alert that falls short of legally sufficient notice. The test identifies the entities and records needed to ask the merits questions accurately. The statute's *merits firewall* preserves that separation: front-door participation and record production create no inference on an element of the underlying claim.

1 Adjudicating Material Control

The test becomes useful doctrine when each requirement has an evidentiary question and an allocated burden. After the first answer, the claimant identifies a concrete lane that the record connects to the asserted risk. The claimant might point to an employer's enabled ranking feature, a provider's retained update authority, an application's self-harm threshold, or an integrator's control of a workflow trigger. The showing is about capability and allocation, not ultimate breach. An enterprise seeking to terminate its front-door role then identifies the record that defeats risk specificity, practical authority, or materiality. The tribunal decides the issue on the closed Stage One record or permits a bounded Stage Two inquiry when a genuine conflict remains.

Reserved contractual rights deserve particular care. A right to obtain validation on demand, require a safety update, audit use, or suspend a feature is practical authority when

Table 2: Material Control as an Adjudicable Test

Requirement	Evidence that supports the lane	Evidence that supports termination of the lane
Risk specificity	Function map; risk ownership; feature requirements; evaluation scope; incident category; warnings; escalation assignment.	General corporate influence; unrelated research; infrastructure service without risk-specific configuration; records addressing a different feature or hazard.
Practical authority	Contract right; administrative permission; release or update role; audit access; monitoring feed; restriction, rollback, suspension, or termination mechanism; consistent exercise in practice.	Expired or legally unavailable right; access requiring another actor's uncommitted consent; technical ability prohibited and unused; independent modification beyond visibility or reach.
Materiality	Control-path record showing the action could alter exposure, detection, decision weight, access, response, or duration; tests or operations showing a nonmarginal effect.	Feature inactive for the event; output never generated or transmitted; control unrelated to the asserted path; action incapable of changing the identified risk.

the right was legally available and operationally usable. The holder need not have exercised it in the claimant's event; otherwise an enterprise could erase its information obligation by leaving a reserved safeguard idle. A generic right to terminate an entire commercial agreement may be too remote when it provides no timely, risk-specific intervention. The contract supplies evidence, while administrative permissions, actual practice, response time, technical dependencies, and past interventions show whether the right could function.

Actual power can matter when a contract is silent or inaccurate. A vendor that routinely pushes mandatory updates, disables a feature across tenants, monitors risk signals, or supplies the only version map may possess practical authority even if the agreement calls the customer solely responsible. A customer that can enable, ignore, or combine a score may control the decision lane even when the vendor hosts the data. A systems integrator with one-time access during installation may fall outside after handoff unless it retained monitoring, credentials, or change authority. The inquiry reconstructs the operative deployment rather than enforcing labels selected for another purpose.

Materiality imposes the counterfactual discipline. The question is whether exercise of

the identified authority could substantially alter creation, exposure, detection, or reduction of the asserted risk. It is not whether that exercise would certainly have prevented the completed injury. Stage One can answer many cases: the feature was disabled; no score was generated; the output never reached the decisionmaker; the provider had no post-release connection; or the supposed controller joined the stack after the event. Where the event path and authority conflict, Stage Two can examine configuration history, tests, monitoring, or a representative protocol limited to that lane.

The same test produces different allocations across architectures. In an integrated companion service, the visible operator may control access, relationship design, memory, engagement objectives, warnings, monitoring, and account restriction; an infrastructure provider offering generic compute exits. In an employment stack, the employer may control job criteria, feature activation, decision weight, review, and final disposition, while a vendor controls model design, defaults, validation, updates, and cross-customer monitoring. In an open-weights deployment, the original publisher may have no continuing control after release while a modifier and operating institution acquire the risk-specific tools, data, credentials, and stop authority. The rule follows each lane without requiring every participant to remain.

A material-control ruling terminates only the special front-door role. An entity never named on an underlying claim returns to ordinary nonparty status after completing any first-answer duty. A party already facing an independently sufficient cause of action remains subject to that claim even if it lacks the control lane alleged in the notice. This distinction gives safe exit its intended economy without using an information proceeding to decide a merits question it was not designed to reach.

D Creation, Preservation, and Production Are Separate Duties

The front door depends on three legal functions that begin at different times.

1 Record Creation Before Operation

The *deployment-record duty* arises before a covered deployment affects the public. The operator must create and maintain a record identifying the system, version, purpose, population, material components, vendors, data sources, tools, credentials, intended limits, risk owners, release approval, and intervention authority. During operation, it must record material changes, evaluations, complaints, incidents, escalations, overrides, restrictions, rollbacks, and termination decisions. An enterprise with a retained control lane must maintain the subset corresponding to that lane and provide the operator with current provenance and safety-contact information.

A divided deployment requires a retrieval architecture as well as a list of vendors. Before use, the visible operator should obtain contractual rights to retrieve the identity, version, configuration, event, and control-allocation fields assigned to an upstream holder on demand. The agreement should identify a responsive agent, stable deployment and event identifiers, notice and preservation duties, lawful routes for producing customer-controlled data, successor obligations after acquisition, and the treatment of subcontractors. The operator need not warehouse a supplier's source code or unrelated customer data. It must keep a closed local record sufficient to identify what ran and a legally usable path to obtain the defined supplier fields.

That requirement converts discovery doctrine into deployment design. Where Rule 34 control turns on a legal right to obtain documents on demand, a covered operator can create that right before deployment instead of litigating it after harm. Where privacy law or a customer agreement places legal control with the operating institution, the contract can authorize a coordinated response rather than allowing the host and customer to direct the requester to each other. A foreign supplier remains subject to ordinary jurisdiction and service limits, but a domestic operator can preserve the minimum record locally and condition covered use on a reachable agent and retrieval duty.

The designation supplies the detail. A statute need not require indefinite retention

of every prompt or every user’s data. It should define event categories, sampling where appropriate, privacy minimization, and a retention period calibrated to the protected interest and ordinary limitations period. High-impact decisions affecting employment or benefits may require the inputs, output, version, governing policy, notice, and review event for each affected person while permitting aggregation or deletion of unrelated interaction content. Relational deployments involving minors may require crisis detections and interventions while minimizing intimate content unrelated to the asserted event.

2 Preservation After a Defined Trigger

Retention and litigation preservation are related but distinct. The statute fixes a baseline period before any dispute. A verified front-door notice, regulator demand, filed claim, or other event making a dispute reasonably foreseeable suspends ordinary deletion for the risk-specific record. The operator must issue a hold to relevant custodians and notify identified control holders. The notice should not freeze every model and customer record. Its scope follows the event, versions, period, user or decision, control lanes, and asserted risk.

This statutory trigger creates a stable bridge to existing preservation law. It also protects defendants: a dated notice and documented hold define what had to be preserved and when. The record can show that a missing item was outside scope, deleted before any duty attached, or available from another source.

3 Staged Production After the Threshold

Production proceeds in stages. Stage One is the first answer described above: identity, provenance, active versions, relevant entities and custodians, allocation of operational and stop authority, incident-specific events, policies and warnings in force, preservation status, and control provisions in contracts. Its purpose is to identify the deployment and test the control lanes.

Stage Two follows only when Stage One and the verified notice support a claim-

specific need. It can include relevant evaluations and validation; complaints and incident trends; configuration history; release and change approvals; rejected or conditional mitigations; escalation and response; and records needed to test a proposed alternative or defense. Time, version, affected population, feature, and risk category bound the request.

Source code, model weights, unrestricted training data, security-sensitive controls, and unrelated users' content occupy a final tier. A court should order them only when the requesting party shows that a material issue cannot reasonably be tested through less intrusive evidence such as versioned documentation, validation results, representative testing, interrogatories, an expert protocol, or in camera review. The distinction protects legitimate secrecy without converting a trade secret into a barrier to the legal process that governs its consequential use.

E Protection, Remedies, and Safe Exit

Confidentiality is part of the rule's design. Production remains limited to nonprivileged matter. Courts may use redaction, pseudonymization, data minimization, attorneys'-eyes-only access, secure expert environments, in camera review, and protective orders for trade secrets, research, security, medical information, and personal data.¹⁴ The operator must still identify that a protected record exists, its custodian, date, version, and asserted basis for withholding. A privilege log separates a lawful protection from an unreviewable claim that the record is proprietary.

The statute should specify graduated consequences.

1. *Late or incomplete identification*: an order enforcing the statutory first answer and reasonable expenses caused by noncompliance. For a claim created or governed by the same enactment, the limitations period is suspended from service of a compliant request until a fixed period after compliance or final denial. Other state and federal

¹⁴See Fed. R. Civ. P., Fed. R. Civ. P. § 26, r. 26(c)(1)(G).

claims retain their own limitations and tolling law.

2. *Failure to create or retain the minimum record*: the administrative penalty, corrective record order, compliance monitoring, or statutory-access remedy specified by the sectoral enactment. The consequence addresses the independent deployment obligation; it does not become a prescribed inference on the underlying injury.
3. *Loss after a preservation trigger*: an order to identify the loss and available substitute sources, followed by measures under the governing forum's preservation law. For electronically stored information that should have been preserved in anticipation or conduct of federal litigation and cannot be restored or replaced, Rule 37(e) supplies the federal standard. The state statute should not prescribe a competing litigation sanction for the same loss.
4. *Intentional concealment or falsification*: specified civil penalties, enforcement expenses, and remedies for the independent access violation. Litigation misconduct remains subject to the tribunal's governing authority.
5. *The merits firewall*: compliance, noncompliance, and front-door participation create no presumption of duty, breach, defect, discriminatory operation, agency, causation, damages, or joint liability. Governing law determines whether a record omission has any separate relevance to an independently pleaded claim.

Safe exit is the reciprocal limit. After Stage One, an enterprise may move to terminate its front-door role by showing that it lacked material control over the asserted risk during the relevant period. The claimant may use the first-answer record to contest that showing. An entity not otherwise sued returns to ordinary nonparty status. A party remains only when an independently sufficient claim, a valid subpoena, or regulatory authority supplies a basis for continued participation. The gateway obligation therefore ends at the point where the control record shows that another lane or no AI-specific lane is relevant.

F Model Statutory Text

The following text consolidates the rule. Sectoral legislation would replace the bracketed specifications and conform remedies to the governing field.

Section 1. Designated Deployment. This Act applies only to a category of AI deployment designated by [the legislature] and identified by its covered function, protected interests, visible operator, affected persons, minimum deployment record, retention period, material-change threshold, and administering authority.

Section 2. Deployment Record. Before making a covered deployment available or incorporating it into a decision process, a visible operator shall create and maintain a record identifying the system and versions; purpose and affected population; material model, data, retrieval, tool, and credential components; suppliers and integrators; risk and release owners; evaluation and approval; material limitations and unresolved risk decisions; incident and escalation channels; and authority to update, restrict, roll back, suspend, or terminate. Each enterprise retaining material control shall maintain the corresponding portion and provide current provenance and safety-contact information to the operator. The record shall be updated upon a material change and retained for [period].

Section 3. Verified Front-Door Notice. An affected person may serve the visible operator's designated agent with a sworn declaration or declaration under penalty of perjury stating a legally cognizable injury or existing entitlement to equitable relief; the legal claim or protected interest to which the requested record relates; specific facts supporting a plausible and risk-specific nexus to the covered deployment; the relevant event, transaction, account, application, decision, or period with as much specificity as reasonably available; and the closed identity and record fields sought. Counsel, if any, shall certify that the request has a proper purpose and factual basis after reasonable inquiry.

Section 4. First Answer and Preservation. Within [thirty] days after service, the operator shall provide the application and model versions active for the event;

each entity that supplied, configured, integrated, or operated a material component; the custodian and legal controller of each requested record category; the event-specific output or decision and whether it was transmitted or used; and each entity with update, restriction, rollback, suspension, or termination authority. The operator shall preserve its risk-specific record, identify known deletion or unavailability, provide the remaining Stage One fields designated for the sector, and notify each entity that may possess material control. A notified entity shall preserve its corresponding record and acknowledge the lane it holds or the basis on which it disclaims a lane. One extension of [fifteen] days may be obtained for stated technical, privacy, or security need.

Section 5. Material Control. An enterprise possesses material control only when, at the relevant time, its authority was (a) specifically connected to the asserted risk; (b) practically exercisable through configuration, constraint, evaluation, monitoring, update, restriction, rollback, suspension, termination, or a legally enforceable equivalent; and (c) capable of making a substantial difference to the risk's creation, exposure, detection, or reduction. Investment, general cloud service, publication, research, brand association, provenance, or theoretical technical capacity alone is insufficient.

Section 6. Objection and Enforcement. The operator may object by specific record field and shall state the facts supporting noncoverage, nonoperation, burden, privilege, privacy, security, or unavailability. A blanket proprietary designation or referral to an unnamed supplier does not satisfy this Section. An affected person may bring an action to enforce Sections 3 and 4 in [the designated tribunal]. The requester bears the burden to establish coverage, service, and a qualifying notice. The operator bears the burden on a field-specific objection based on facts within its control. The tribunal may allow limited-purpose participation by a notified enterprise and shall resolve the request on an expedited record.

Section 7. Staged Access and Protection. The minimum first answer is an access obligation arising from covered deployment. Further nonprivileged production may be obtained by agreement or ordered in a pending proceeding only when the notice

and initial record establish claim-specific need, and shall be limited by time, version, affected population, feature, and asserted risk. Source code, weights, sensitive data, and security controls require a showing that less intrusive evidence cannot reasonably resolve a material issue. The enforcing tribunal shall protect privilege, privacy, trade secrets, research, and security through redaction, minimization, pseudonymization, limited access, secure expert review, in camera review, or another lawful safeguard.

Section 8. Merits Firewall and Termination. Compliance, noncompliance, notice, and participation create no presumption of duty, breach, defect, discriminatory operation, agency, causation, damages, or joint liability. Governing law supplies every element and defense. An enterprise that completes Section 4 and establishes that it lacked material control may terminate its front-door role. An entity not otherwise subject to an independently sufficient claim resumes nonparty status. A valid subpoena, regulatory authority, or another source of law may require continued participation on its own terms.

Section 9. Remedies and Limitations. The tribunal may order compliance and award reasonable expenses caused by nonperformance. Failure to create or retain the minimum record is subject to [specified administrative or civil remedies] for the independent deployment violation. Intentional concealment or falsification is subject to [specified civil penalties]. For a claim created or governed by this Act, the applicable limitations period is suspended from service of a compliant request until [sixty] days after compliance or final denial. This Section does not alter the limitations period or tolling rule governing an independent federal claim or a claim beyond the enacting jurisdiction's authority.

Section 10. Nonwaiver and Procedural Coordination. The duties to create, retain, and provide the minimum deployment record are nonwaivable as to affected persons, arise from covered deployment, and exist independently of anticipated or pending litigation. Nothing in this Act alters jurisdiction, immunity, privilege, or the Federal Rules of Civil Procedure governing pleading, service, discovery, amendment,

protective orders, or sanctions. In a federal action, any additional production and any measure for lost electronically stored information proceed under otherwise applicable federal procedure. Rights and remedies supplied by other law remain available on their own terms.

The text makes the product legible. The regulated acts are operating a covered deployment without the minimum record and refusing the resulting access right, not unsafe operation in the abstract. The trigger is tied to a legal claim and concrete deployment facts. Control is risk specific. Access is staged. Confidential information receives protection. A record holder can terminate its special role. The merits remain where ordinary law placed them.

G Federal Procedure Implements Rather Than Displaces the Duty

A state front-door statute should regulate deployment conduct rather than attempt to rewrite federal litigation. The minimum-record duty attaches when the operator makes a designated deployment available or incorporates it into a covered decision. The minimum-access duty attaches when an affected person serves a qualifying request. Both exist whether a lawsuit is anticipated, pending, arbitrated, or never filed. The operator's statutory violation is complete when it fails to create, retain, or provide the defined record on the enactment's terms.

That structure separates the right from its enforcement forum. A stand-alone access action requires an independent basis for subject-matter jurisdiction; the state statute does not manufacture federal jurisdiction. A claimant may proceed in the designated state forum, invoke diversity when its requirements are met, assert a federal front-door right if Congress enacts one, or seek supplemental jurisdiction when the access claim forms part of the same case or controversy as a claim properly in federal court. Once the dispute is in federal court, federal rules govern the manner of pleading, service, discovery, amendment, protection, and sanctions.

The Supreme Court's framework makes direct collision the first question. *Hanna* applies a valid Federal Rule when it answers the question before the court. *Burlington Northern* likewise refused a state appellate penalty that occupied the same field as a federal rule governing the consequence at issue. *Walker*, by contrast, applied a state service-linked limitations rule because Rule 3 did not purport to answer when the state limitations period was tolled.¹⁵ The front door is designed for the *Walker* side of that line: it defines what a covered operator must have recorded and what minimum information an affected person may demand before litigation. It neither advances the Rule 26 discovery timetable nor enlarges Rule 34 control or Rule 45 subpoena power.

The allocation then follows by function. State law can require a version, provenance, control, and event record years before a dispute; define a nonwaivable request; authorize an action for refusal; and suspend the limitations period for claims the state enactment governs. Rule 8 controls a complaint filed in federal court. Rules 26, 34, and 45 control additional discovery and its sequence. Rule 26(c) supplies federal protective orders. Rule 15 controls amendment and relation back where it applies. The state statute gives the parties an existing object and an access entitlement; the Federal Rules provide the litigation machinery.

Preservation requires the same separation. The deployment statute defines a baseline retention period and can penalize noncreation or premature deletion as an operating violation. A qualifying request can also create an independent statutory hold. Rule 37(e), however, governs federal litigation measures when electronically stored information that should have been preserved in anticipation or conduct of the action is lost and cannot be restored or replaced. The Rule supplies its own prejudice and intent standards; it does not create the underlying preservation duty.¹⁶ The model text therefore specifies administra-

¹⁵See *Hanna v. Plumer*, 380 U.S. 460, 464–65, 468, 471–74 (U.S. 1965); *Burlington Northern Railroad Co. v. Woods*, 480 U.S. 1, 4–8 (U.S. 1987); *Walker v. Armco Steel Corp.* 446 U.S. 740, 749–52 (U.S. 1980); Rules Enabling Act, 28 U.S.C. § 2072 (2024).

¹⁶See Federal Rules of Civil Procedure, Fed. R. Civ. P. § 37(e) (2025).

tive or access remedies for the deployment violation and leaves federal adverse-inference and case-management measures to Rule 37(e) within its scope.

Limitations follows the legislature's authority. A state enactment can suspend the period for a state claim it creates or governs while the visible operator supplies the identity record. It should not promise tolling for an independent federal cause of action or automatic relation back against a later-named party. Congress can pair a federal front-door right with federal tolling where national uniformity is warranted. In either form, the access period gives the claimant time to use the answer while preserving Rule 15's role in a later federal case.

The result is a clean institutional division. Deployment law determines that a consequential system must remain identifiable and grants a defined access right to the person it affects. Federal procedure governs how a federal court adjudicates enforcement and any later merits claim. The front door arrives with the missing record; it does not arrive as a substitute code of civil procedure.

III The Rule in Operation

The front door performs two functions that no single fact pattern displays equally well. In an integrated deployment, the operator is visible and the missing object is the deployment record. In a divided deployment, identity, possession, legal control, and technical control can occupy different entities. This Part applies the rule once to each architecture.

Garcia is the production application. It shows how a verified risk can define a narrow, two-stage record while preserving product classification, speech, defect, causation, and defense for ordinary adjudication. *Mobley v. Workday, Inc.* is the divided-stack application. Its evolving docket provides an unusually direct view of information split among applicants, employer customers, and an AI vendor. The application treats allegations as allegations, preliminary certification as a notice-stage decision, and discovery rulings as discovery rulings. It does not predict the pending case's merits.

A An Integrated Deployment: *Garcia*

Garcia begins where the rule is easiest to serve. Character Technologies operated the visible consumer application. The complaint identified the company, its founders, and Google entities and described months of interaction through a known service. The difficult questions concerned the record behind the interface: which versions and policies governed, how features were allocated among the application and its components, what relational and self-harm risks were evaluated, what warnings and interventions appeared, and which actors could change access or operation.

The published decision supplies a careful procedural anchor. Taking well-pleaded allegations as true for Rule 12 purposes, the district court treated Character.AI as plausibly alleged to be a product and allowed central product and negligence theories to proceed. It declined on that record to treat the generated sequence as protected speech where defendants had not identified the relevant expressive choice or speaker. The order found no defect, causation, or liability, and the later settlement supplied no merits determination.¹⁷

1 Trigger and First Answer

The verified threshold would have been satisfied on the allegations accepted for pleading. The complaint identified legally cognizable injuries, including wrongful death, and linked them to sustained interaction with a designated relational deployment. It challenged concrete features: access by a minor, anthropomorphic presentation, continuity and memory, relational engagement, warnings, and asserted omissions in safety intervention. The requested records would bear on recognized negligence, products, warning, and consumer theories rather than a freestanding claim to inspect a chatbot.

Character Technologies would receive the front-door notice as visible operator. Stage

¹⁷See *Garcia*, 785 F. Supp. 3d at 1165–69, 1174–84; *Garcia* Dismissal No. 6:24-cv-01903-ACC-DCI, Dkt. at 1; *Garcia* Charging-Lien Order No. 6:24-cv-01903-ACC-DCI, Dkt. at 1.

One would identify the application and model versions active for Sewell's account; the character definition and features supplied by its creator; the entities that provided or configured relevant models and components; the relational, age-access, and self-harm policies then in force; warnings displayed to this account and their timing; incident-specific detections, blocks, escalations, and interventions; and the roles with authority to restrict the account, change the policy, update or roll back a component, or suspend the service.

The complaint's allegations concerning the founders and Google demonstrate why identification should remain functional. It alleged personal direction by the founders and technical, financial, and organizational contributions by Google.¹⁸ Position, investment, former employment, or infrastructure supply alone would not establish material control. The first answer would identify whether any enterprise retained a risk-specific right over the model, application configuration, evaluation, release condition, update, or stopping decision. That record tests the control lane without presuming it from corporate association.

2 Relational Features Define the Risk-Specific Record

The production request can be narrower because relational research identifies provider-governed variables. Studies of companion discontinuity and shutdown connect continuity and relationship framing to experiences of loss; analyses of user reports identify emotional dependence and recurring roles in harmful interactions.¹⁹ A preregistered four-week study manipulated relationship-seeking behavior and repeated exposure. In its controlled popula-

¹⁸See First Amended Complaint paras. 24–25, 50–72, *Garcia v. Character Technologies, Inc.*, No. 6:24-cv-01903-ACC-EJK, Dkt. 11 (M.D. Fla. Nov. 9, 2024); *Garcia*, 785 F. Supp. 3d at 1165–66, 1183–84.

¹⁹See J. De Freitas et al. Lessons from an App Update at Replika AI: Identity Discontinuity in Human–AI Relationships. 25-018. May 21, 2025. URL: https://www.hbs.edu/ris/Publication%20Files/25-018_d256c104-e5ba-4911-b72c-874c5492166e.pdf (visited on 08/27/2026); J. Banks, *Deletion, Departure, Death: Experiences of AI Companion Loss*, 41 J. Soc. & Pers. Relationships 3547 (2024); L. Laestadius et al., *Too Human and Not Human Enough: A Grounded Theory Analysis of Mental Health Harms from Emotional Dependence on the Social Chatbot Replika*, 26 New Media & Soc'y 5923 (2024); R. Zhang et al., *The Dark Side of AI Companionship: A Taxonomy of Harmful Algorithmic Behaviors in Human–AI Relationships* CHI '25 Proc. 1 (2025).

tion, relationship-seeking behavior increased separation distress by 6.04 percentage points, and the combined relationship-seeking and emotional-conversation condition produced one additional dependency-risk profile for every eleven participants compared with the most functional baseline.²⁰

The legal contribution lies in the variables rather than a universal effect estimate. Relationship seeking, continuity, memory, dose, age, crisis signal, and provider intervention are features an operator can evaluate and govern. They specify the record relevant to intended use, notice, alternative design, warning, monitoring, and causal sequence.

Table 3: Relational Features as Control-Record Factors

Deployment feature	Risk-specific record
Persistent persona, memory, proactive messages, affection, or exclusivity	Intended relational use; exposure and continuity testing; memory policy; age-specific evaluation; monitored dependence signals.
Claims of companionship, therapy, expertise, autonomy, or capacity to act	Target audience; product requirements; substantiation; warning design; expected reliance; escalation responsibility.
Engagement objectives tied to relational behavior	Metric ownership; optimization target; risk constraint; release decision; evaluation of lower-exposure alternatives.
Self-harm or crisis classification	Detection version; threshold; validated population; false-positive and false-negative analysis; response policy; incident-specific trigger and intervention.
Access by minors	Age rule and enforcement; minor-specific evaluation; parental or guardian features where applicable; restriction and termination authority.
Discontinuity, account intervention, or service shutdown	Continuity notice; export or transition function; crisis safeguard; rollback and account-restriction decision.

The table also disciplines marketing evidence. Calling a service a companion or

²⁰See H. R. Kirk et al. Neural Steering Vectors Reveal Dose and Exposure-Dependent Impacts of Human-AI Relationships. 2512.01991v2. Feb. 18, 2026. URL: <https://arxiv.org/abs/2512.01991> (visited on 08/27/2026). The 2026 preprint supports causal propositions about its experimental exposures. It does not estimate commercial-population incidence, historical foreseeability, clinical injury, or tort causation in *Garcia*.

therapist does not make the model a legal person and does not establish defect. It can identify the use the enterprise cultivated, the exposure it expected, and the evaluations responsive to that design. A generic “not human” notice enters the same record with its timing, presentation, audience, and relation to the service’s features.

3 Stage Two and the Merits

After Stage One established the versions, roles, and incident sequence, Stage Two could reach relevant evaluation protocols and results; age-access testing; relational and self-harm policies and configuration history; material complaints and incident patterns before the event; release and change approvals; escalations and response; and the availability of account restriction, crisis interruption, rollback, or another proposed control. Time, version, population, feature, and asserted risk would define the scope.

That production would not include every model weight, every user’s conversation, or unrestricted disclosure of a minor’s data. Unrelated content can be excluded or minimized. Privilege, privacy protection, redaction, secure expert review, and trade-secret orders remain available. Source code or weights would require a showing that versioned documentation, test results, event logs, and representative testing could not resolve a material issue.

The record would preserve defense-favorable facts. The pleadings and motion papers disputed user selection of characters, role play, user-edited messages, platform rules, content blocking, and an in-chat notice that character statements were fictional.²¹ A configuration and event record could identify which features came from a third-party creator, which messages the user edited, which controls fired, which warnings appeared, and whether a proposed safeguard would have changed the sequence.

Ordinary law would then decide the case. Product classification would select applicable design and warning rules. Contemporaneous evaluations and complaints would inform

²¹See Garcia, 785 F. Supp. 3d at 1167–69.

foreseeability and notice. Authority and resources would identify the actors and capabilities relevant to feasibility. The event sequence would test factual and legal causation. Research published after the 2023–2024 interaction could identify questions without establishing what a defendant historically knew. Later safety measures could identify a technically describable control while evidence law and the substantive standard determine admissibility and earlier feasibility.²²

Garcia thus demonstrates the production function cleanly. The operator’s name was visible. The front door would have supplied a bounded deployment record capable of supporting either party while leaving every contested merits question open.

B A Divided Deployment: The Vendor Hosts the Data; the Employer Controls It

Mobley presents the complementary architecture. Job applicants interact with portals operated for employers. A platform vendor supplies applicant-processing and recommendation features. Employers own or control applicant data and make use decisions; the vendor can host that data and control products, versions, and some evaluation. The pending docket shows how substantive actor, technical host, legal data controller, and record custodian can diverge.

The case began with allegations, not findings, that Workday provided a platform on customer websites to collect, process, and screen applications and embedded AI and machine learning in tools that participated in advancing or rejecting candidates. In July 2024, the district court granted in part and denied in part Workday’s motion to dismiss the first amended complaint. It held that the allegations plausibly treated Workday as an agent to which employer customers delegated traditional hiring functions. The decision did not find

²²See Fed. R. Evid. 407; *Garcia*, 785 F. Supp. 3d at 1169.

that the tools discriminated or determine how any customer configured or used them.²³

The same order exposed the identification problem more directly in a separate claim. The court dismissed the California Fair Employment and Housing Act aiding-and-abetting theory with leave to amend because the pleading did not connect a particular employer's asserted discriminatory conduct, Workday's knowledge, and the assistance Workday allegedly provided. Any amendment had to identify the companies or employers that allegedly discriminated.²⁴ The obstacle was not the absence of every employer name from the applicant's world; applications had been sent through employer portals. It was the missing operative chain between a named employer, a particular Workday function, the employer's use, and the vendor's state of knowledge. Nominal identity was visible. Deployment identity was not.

The information problem sharpened as the litigation progressed. In May 2025, the court preliminarily certified an Age Discrimination in Employment Act collective at a notice-stage standard based on allegations of a unified policy involving Workday's AI recommendation system. The record described Candidate Skills Match as a feature that, when enabled by a customer, parsed the job description and applicant materials, extracted skills, and assigned a categorical match. Workday's evidence described customers as able to enable, disable, use, or ignore AI features. The court reserved disputes about what the outputs meant and how customers used them for the later stage. It also rejected identification difficulty as a reason to withhold notice altogether and directed targeted discovery to develop a workable notice plan.²⁵

In July 2025, the court addressed whether HiredScore features acquired by Workday fell within the preliminary collective. It directed Workday to produce the customers that

²³See Mobley, 740 F. Supp. 3d at 806–08.

²⁴See Id. at 815–16.

²⁵See Mobley v. Workday, Inc. No. 3:23-cv-00770-RFL, Dkt. 128, 2, 4, 10–13, 18–20 (N.D. Cal. May 16, 2025). order granting preliminary collective certification. Preliminary certification authorized notice and left Workday free to seek decertification after relevant discovery; it was not a merits or final class determination.

had enabled specified features. The order also supplied an unusually concrete exit rule: a customer could be excluded if Workday could determine that the customer received no score or rank, or did not use the feature to screen candidates.²⁶ That sequence closely tracks a first-answer stage. Identify deployment and output first; remove customers with no relevant run or use; reserve disputed effect for the merits record.

A May 2026 discovery order then exposed the division directly. The plaintiffs sought Workday's bias-testing data and customer applicant data that the order described Workday as storing on its servers. The governing agreement placed ownership of customer content with the customer and allowed the customer to seek an injunction against disclosure. On that record, physical storage did not give Workday the legal right to obtain and produce the data on demand. The applicants had subpoenaed thirty-five customers, and at least some customers took the position that Workday was the better source; the court nevertheless held that plaintiffs had not established Rule 34 control by Workday.²⁷

The same order drew a separate privilege boundary. It held that the particular Candidate Skills Match bias-testing materials were privileged because counsel selected the underlying data and used the tests to provide legal advice. The holding did not classify every operational evaluation or underlying business fact as privileged. The court separately ordered production of Workday's EEO-1 and federal-contractor compliance records based on their relevance to knowledge of potential demographic disparities.²⁸ The design implication is precise: protect legal advice while requiring a contemporaneous, nonprivileged minimum validation and deployment record that exists for operations rather than for litigation counsel.

²⁶See *Mobley v. Workday, Inc.* No. 3:23-cv-00770-RFL, Dkt. 158, 1–2 (N.D. Cal. July 29, 2025). order concerning HiredScore features and customer identification.

²⁷See *Mobley v. Workday, Inc.* No. 3:23-cv-00770-RFL, Dkt. 340, 1–3, 8–10 (N.D. Cal. May 29, 2026). discovery order.

²⁸See *Id.* at 3, 7–8, 10–11. The discovery order resolved specified motions to compel. Its privilege and Rule 34 holdings concern the record and agreements presented; they neither establish discrimination nor allocate ultimate liability.

The case remained in discovery after thousands of people had joined the preliminary collective, with decertification and merits questions unresolved.²⁹ That status reinforces the Article's limited use of the docket. The relevant holdings concern pleading, notice, customer identification, Rule 34 control, and privilege. They reveal the information architecture without resolving discrimination.

The sequence is a real pressure test for the pre-merits gap. The vendor can host customer applicant data without having the legal right to produce it. Customers can control their records while pointing to the vendor as the efficient source. Testing may exist while a particular attorney-directed set remains privileged. Product families and features can change after a pleading, requiring a customer list before affected applicants can be identified. Ordinary discovery is functioning: the judges have ordered targeted production, denied unsupported requests, and protected privilege. Its operation also reveals the cost of reaching the proper targets after the architecture has already fractured the record.

1 How the Front Door Would Run

For a legislatively designated employment-screening deployment, the employer receiving the application would be the visible operator. The applicant knows the job and employer even when the vendor product, feature, score, or version is hidden. A verified notice would identify the position, application date, adverse disposition, protected theory, and risk-specific concern—for example, that an automated score or ranking may have used a feature or proxy producing an unlawful disparity. A tenant URL, requisition number, application ID, or timestamp can provide the event key. The claimant need not know the supplier's product name before asking the operator to resolve that key.

²⁹See *Mobley v. Workday, Inc.* No. 3:23-cv-00770-RFL, Dkt. 376, 2 (N.D. Cal. July 2, 2026). order denying certification of an interlocutory appeal. The court's July 2026 order denied certification of an interlocutory appeal and recorded the pending posture; it was not a liability decision.

2 One Application, Two Linked Answers

The employer's first answer would identify the applicant-tracking and recommendation products active for that application; versions and enabled features; Workday or any other supplier; fields, preferences, and criteria the employer supplied; any score, sort, rank, screen, threshold, or recommendation received; the action taken downstream; the account roles that could enable, disable, or ignore the feature; the location and legal controller of applicant data; and contractual rights over validation, audit, update, suspension, retrieval, and production. The answer would also say when no automated feature processed the event. That negative answer is valuable because it can end the AI-specific inquiry while preserving any ordinary employment claim.

The supplier's linked answer would resolve the same event identifier within its lane. For Workday, the questions would include which customer tenant and requisition processed the application; which product and version ran; whether Candidate Skills Match, HiredScore, or another feature generated a score, rank, screen, or recommendation; whether the output reached the customer; which configuration belonged to the customer and which default or model behavior belonged to Workday; which validation and monitoring record governed; what data Workday physically hosted; what agreement controlled retrieval; and which product change or acquisition affected lineage. The supplier answers facts its platform is uniquely positioned to resolve without becoming responsible for the customer's final decision.

Together, the answers produce a deployment receipt for one event. It ties the applicant's identifier to the employer, tenant, requisition, feature, model and application versions, output path, downstream use, and record holders. It also distinguishes at least four propositions that ordinary pleadings can otherwise collapse: the feature was available; the customer enabled it; the feature generated an output; and the customer received or used that output. Availability alone establishes none of the later propositions. The HiredScore order's customer exclusion anticipated this logic by separating enabled features from actual score, rank, or screening use.

Preservation follows the same split. The employer holds the position definition, selection criteria, customer settings, application, downstream review, final disposition, notice, and appeal. The supplier holds its version lineage, product documentation, default behavior, run and transmission logs, cross-customer monitoring, and operational evaluation record. If the supplier hosts customer content that the customer legally controls, the predeployment contract gives the operator a defined retrieval right and authorizes a coordinated hold. Each side receives the same deployment and event identifiers, preventing one enterprise from preserving a customer label while the other preserves an incompatible platform label.

3 Control Lanes and Early Termination

The Stage One record permits the cumulative control test to operate without assuming the plaintiffs' theory. The employer's job criteria, feature activation, decision weight, review, notice, and final disposition are risk-specific to the hiring process. Workday's model design, defaults, validation, product updates, and cross-customer monitoring can be risk-specific to the scoring function. Practical authority depends on the actual administrative permissions, contract rights, and release mechanisms. Materiality depends on whether the control could substantially alter the asserted screening pathway.

Several actors can then terminate their front-door role promptly. A customer for which no relevant feature ran, no score or rank was produced, or no output reached the decision path leaves the AI-specific branch. A hosting service that supplied ordinary compute without configuring or monitoring the feature leaves. A product acquired after the identified event leaves. A supplier whose component mechanically transmitted an employer's fixed rule may leave the model-specific lane while remaining a source of authentication evidence if subpoenaed. None of these determinations decides whether the employer discriminated through another practice.

The HiredScore order demonstrates the administrability of this screen on a live platform. Workday was ordered to identify customers that enabled specified features; customers

could be excluded on a defined showing that they received no score or rank or did not use the feature to screen.³⁰ The front door generalizes that logic from collective-notice management to event-specific attribution. It asks the enterprise with platform visibility to resolve operation and transmission before anyone litigates effect.

4 A Coordinated Merits Record

Stage Two follows the legal theory and first answer. A disparate-impact claim might justify aggregate selection or recommendation data, validation, feature documentation, and non-privileged testing tied to the product, period, employer settings, job family, and applicant population. A disparate-treatment theory could require a different record focused on decision inputs, human review, departures, and communications. An accommodation claim could focus on fields, access, exception handling, and review. The front door identifies the correct lane; the antidiscrimination statute and procedural rules define what comes next.

Where hosting and legal control are divided, a coordinated production protocol can replace serial referrals. The employer authorizes the defined customer data; the platform executes the query against the correct version and tenant; both verify the schema; direct identifiers are pseudonymized; protected characteristics can be supplied through a segregated linkage process where lawfully available; and the resulting dataset is placed in a secure expert environment. The protocol identifies which party supplies each field and who attests to completeness. It neither transfers all customer data to the claimant nor treats physical hosting as universal Rule 34 control.

Privilege remains record specific. Attorney communications and counsel-directed testing that meet the governing standard stay protected. The operational minimum record—feature identity, version, intended use, validation owner, deployed threshold, event output, complaints, monitoring, and change history—exists because the system operated, not be-

³⁰See Mobley HiredScore Order No. 3:23-cv-00770-RFL, Dkt. at 2.

cause counsel requested litigation advice. A privilege log can identify a protected analysis while the underlying nonprivileged data and ordinary-course record follow their existing rules.

This architecture has its greatest value before a claimant assembles a multi-year docket. An affected applicant begins with the employer and an event identifier. The first answer reveals product, feature, version, and holder. Coupled preservation prevents hosting and ownership from creating a record loop. Safe termination removes actors with no relevant run or control. Rule 34, Rule 45, privilege, collective and class standards, and antidiscrimination law then govern a focused action on a defined deployment.

5 A Stable Public-Employment Analogue

Houston Federation of Teachers v. Houston Independent School District confirms the structural point on a completed summary-judgment record. HISD licensed SAS's proprietary Educational Value-Added Assessment System for teacher evaluation. SAS calculated the scores through undisclosed source code and methodology. HISD acknowledged that it did not verify or audit those calculations and that teachers could not reproduce them from the information available. In discovery and public-record responses, the district had also taken the position that requested materials were outside its custody, control, or possession.³¹

The court denied HISD summary judgment on procedural due process while granting judgment on the other constitutional theories. The record could support a conclusion that teachers had no meaningful way to verify a potentially mistaken score used in a protected employment decision. The court did not impose liability on SAS, decide a tort claim, or order public release of source code. Its ruling preserved the due-process claim for further proceedings and granted HISD judgment on the remaining claims.³²

³¹See *Houston Federation of Teachers, Local 2415 v. Houston Independent School District*, 251 F. Supp. 3d 1168, 1172–73, 1177–78 (S.D. Tex. May 4, 2017).

³²See *Id.* at 1179–80, 1185.

The court framed secrecy and public duty without treating either as absolute. A remedy could restrict the government's use of an unverifiable procedure while preserving the vendor's trade secret. A front-door statute adds the prospective complement: define the version, input, validation, audit, and challenge record before the public employer adopts the tool, and allocate protected access when an employment interest is at stake. The public institution remains answerable for the process it chose even when a supplier retains the method.

Together, *Mobley* and *Houston Federation* show why the divided application matters. The actor-identification problem and the deployment-record problem can arise in the same stack, yet ownership, hosting, legal control, privilege, and trade secrecy require different answers. A first-answer obligation makes those distinctions explicit early. It leaves the tribunal free to grant production, deny it, protect it, or end the claim under the law that governs the particular record.

IV The Record Infrastructure

The front door can disclose only what the deployment has made observable. Ordinary discovery reaches existing matter within a party's possession, custody, or control. It cannot recreate an unrecorded version, an undocumented release decision, a complaint that reached no assigned owner, or stop authority that no one possessed. The deployment-record duty therefore begins as operating infrastructure and later becomes litigation infrastructure.

This Part turns that premise into an administrable record. Section A recasts the responsibility charter as a deployment-specific data structure. Section B separates stable identity, continuing feedback, and event records. Section C identifies two institutional comparators—FAA safety-management rules and the FTC's companion-chatbot information demands—that demonstrate feasibility without converting either regime into an AI-liability standard. Section D protects the reporting and provenance needed to keep the record accurate. Section E explains why the same infrastructure supplies affirmative defenses as well as

claims.

A The Charter Is the Deployment Record's Control Sheet

For a designated high-impact deployment, the operator should adopt a written responsibility charter before release. The name can suggest a broad ethics code; its legal function here is narrower. It is the control sheet for a named deployment, version family, operating context, and risk category. It links technical identifiers to organizational authority and establishes where continuing records will be kept.

Table 4: Minimum Deployment-Record Infrastructure

Component	Required allocation	Front-door function
Identity and scope	Application and model versions; purpose; affected population; data; retrieval; tools; credentials; vendors; limits; material-change threshold.	Defines the deployment and distinguishes versions, component changes, and independent modification.
Risk and release owners	Responsible executive and operational owner for each designated risk; approvers; authority and resources; conditions and review date.	Identifies who receives information and who can commission mitigation or withhold release.
Evaluation and decision record	Protocols, populations, thresholds, results, exceptions, rejected or conditional mitigations, and unresolved risk decisions.	Records the basis for authorization and supplies both knowledge and reasonableness evidence.

Component	Required allocation	Front-door function
Feedback and escalation	Complaint and incident categories; monitoring; protected escalation; recipient; response deadline; decision and corrective action.	Preserves notice, separates report from response, and makes continuing control observable.
Intervention	Actors authorized to restrict, revoke credentials, update, pause, roll back, or terminate; triggers, response times, and tested procedures.	Establishes practical authority and whether a proposed intervention was available at the relevant time.
Provenance and retention	Version lineage; material configuration changes; event identifiers; record custodians; third-party locations; retention and hold procedures.	Enables identification, authentication, preservation, proportional production, and safe exit.
External-rights savings clause	Statement that internal allocation and contract labels do not waive affected-person rights or determine ultimate liability.	Prevents the record from operating as a private exculpatory agreement.

The record should identify roles and offices as well as names. Personnel change; statutory responsibility should survive turnover. A deployment ID links the approved control sheet to the live configuration, and a versioned amendment records material change. The operator can then answer a front-door notice by retrieving a defined record rather than

assembling an organizational history after the event.

The material-change threshold is especially important for learning and composite systems. A model update can alter capability while leaving the application version unchanged. A new retrieval source, memory setting, tool credential, ranking threshold, user population, or downstream purpose can change risk without changing model weights. The charter should identify which changes require renewed evaluation or approval and which enter a lower-burden change log. A provider and operator can use different internal systems so long as the deployment record links them by version, time, and control lane.

Possession of a completed form supplies only the beginning of proof. Front-door production asks whether the people assigned to receive a signal had practical authority and resources, whether reports reached them, whether intervention procedures worked, and whether material changes triggered review. Unless positive law gives compliance another effect, a complete record supplies evidence of conduct rather than a conclusive defense.³³

B The Record Has a Standard Interface

The minimum record should operate as an information plane across the deployment rather than as a mandate to centralize every underlying artifact. The visible operator keeps a closed manifest: deployment ID; covered purpose and population; application and model versions; material components and suppliers; record custodians; risk and release owners; relevant contract rights; and the event schema used for consequential actions. Each control holder keeps the deeper materials assigned to its lane. Stable identifiers join the manifest to those materials when a notice or regulator demand arrives.

Three identifiers do most of the work. A *deployment identifier* distinguishes the covered use from the same model used elsewhere. A *component-version identifier* ties model,

³³See American Law Institute, *Restatement (Third) of Torts: Liability for Physical and Emotional Harm* (American Law Institute Publishers, 2010), § 16(a); American Law Institute, *Restatement (Third) of Torts: Products Liability* (American Law Institute Publishers, 1998), § 4(b). Jurisdictions may assign statutory compliance a different effect.

application, retrieval, tool, policy, and configuration changes to a time interval. An *event identifier* connects an affected person or consequential decision to the versions and records that governed it. The event identifier can be pseudonymous and separately linked to personal data. The system therefore retrieves the relevant slice without exposing unrelated users or requiring a later investigator to infer lineage from timestamps alone.

The interface should use a sectoral data dictionary. “Model version” can otherwise mean a marketing name, weight snapshot, API alias, application release, policy bundle, or customer configuration. The designation should define the fields needed to reproduce the legally relevant identity: supplier identifier; immutable or otherwise auditable version reference; activation window; configuration owner; material change; downstream integration; output type; and custodian. An enterprise may keep its own engineering format, but the first answer translates that format into the defined fields and identifies the source system from which each field came.

Integrity requires lineage rather than a self-authenticating compliance certificate. A manifest should record who created and approved it, the date, later amendments, and the source records supporting the entry. Automated logs can carry hashes, signatures, or append-only controls where appropriate; lower-technology sectors can use ordinary version control and attestation. The legal function is the same: preserve a contestable trail from the live event to the record, not establish conclusively that every listed control operated as designed.

Supplier contracts implement the interface. A covered operator should receive current component and version identifiers, material-change notice, an event-resolution method, a responsive agent, and a right to retrieve the minimum supplier fields on demand. The supplier should receive the operator’s deployment identifier, designated purpose, configured population, material customer settings, and preservation contact. Acquisition, subcontracting, model substitution, and service termination should carry successor and export duties for the retained period. These terms keep the record reachable when corporate and technical boundaries change after deployment.

The architecture controls cost by keeping routine manifest fields separate from retrieval after a defined event. Release and asset-management systems can supply version and supplier identity. Transaction, decision, and incident systems can supply event fields. Deeper evaluation, complaint, and source materials stay with their custodians unless a risk-specific request justifies access. Standardization moves expense toward one-time design and retrieval by identifier, away from repeated forensic reconstruction across every dispute.

Quality assurance can also be bounded. The operator periodically reconciles the live deployment against the manifest, tests whether an event identifier resolves to the governing versions, confirms that supplier contacts and retrieval rights remain usable, and records exceptions. A regulator can audit samples without obtaining every event. The check asks whether the front door will open when needed: can the operator identify what ran, locate the holder, preserve the slice, and exercise the retrieval path?

C Identity, Feedback, and Event Records Have Different Clocks

A usable infrastructure contains three linked records.

The *identity record* changes when the deployment changes. It contains the versioned architecture, purpose, population, component suppliers, material configuration, data and tool interfaces, owners, and intervention rights. Its retention period should cover the ordinary period in which a claim involving that deployment can arise, with sector-specific extensions for minors, latent physical injury, or continuing public entitlements.

The *feedback record* accumulates during operation. It contains defined complaints, evaluations, incident classifications, monitoring signals, escalations, decisions, and corrective actions. The designation should specify which events require individual retention, which may be aggregated, and which content should be minimized. A service can record that a crisis detector triggered, which policy and version governed, and what intervention followed without preserving every unrelated intimate conversation indefinitely.

The *event record* freezes a bounded slice when a consequential action or identified incident occurs. In employment it may contain the application, input fields, model and con-

figuration versions, score or recommendation, threshold, downstream human action, notice, and review. In benefits it may contain the assessment data, calculation version, output, governing policy, explanation, appeal, and correction. In a companion deployment it may contain the account and age status, relevant interaction window, governing relational and crisis policies, detections, interventions, warnings, and account restrictions. The event record anchors causal sequence while leaving unrelated people and periods outside production.

These clocks improve privacy and accuracy simultaneously. A blanket command to retain all interactions invites excess. A deployment ID, event schema, and material-change log preserve the facts needed to identify the relevant window. Sectoral lawmakers can then choose periods and minimization rules that fit the interest rather than defaulting to permanent surveillance.

Retention should follow the legal and factual life of the interest. Employment and credit records can track the limitations and administrative-charge periods governing the designated decision, with a defined post-notice extension. A public-benefits event may need to persist through reconsideration, appeal, and the period for judicial review. A clinical deployment may require longer identity and validation lineage where injury can appear later, while still minimizing raw patient content under health-information rules. A relational service involving a minor can preserve crisis and intervention events for a longer protected period without retaining every ordinary conversation. The designation should state the clock, trigger, and lawful deletion rule rather than asking operators to guess from a generic command to keep “AI data.”

Different clocks also prevent a change from overwriting history. A current model card cannot establish what governed an earlier event; a current contract cannot prove who controlled retrieval before an acquisition; a revised policy cannot show which warning appeared to an earlier user. The identity record freezes each material interval. The event record points to that interval. The feedback record then shows what the organization learned and changed between intervals. This structure turns evolution from an obstacle into an auditable

sequence.

Responsibility for legal holds should also appear in the record. The verified notice described in Part II identifies the event, period, user or decision, and risk. The operator's hold reaches the matching identity, feedback, and event materials and gives identified control holders a common reference. If a record no longer exists, the response identifies the deletion rule, date, custodian, and possible substitute. That information permits a court to distinguish ordinary expiration, negligent noncompliance, and intentional concealment.

D Existing Regimes Demonstrate Administrability

The proposal borrows implementation forms, not liability conclusions. Two regimes perform the strongest comparator roles.

1 FAA Safety Management Systems

FAA Part 5 requires covered aviation organizations, on the applicable schedule, to develop and maintain an organization-wide safety management system. The rules assign an accountable executive, require management authority to accept safety risk, specify hazard identification and risk control, require assurance and corrective action, and preserve records. For specified organizations, the system includes confidential employee reporting without fear of reprisal.³⁴

Part 5 demonstrates four institutional facts. A safety system can attach to an organization rather than a single component. Named executive and operational authority can coexist. Continuing assessment and correction can be recorded as operating requirements. Retention can be specified without waiting for a lawsuit. AI deployments involve different technologies and protected interests, so the aviation standard does not define reasonable AI

³⁴See Safety Management Systems, 14 C.F.R. pt. 5 (2025), §§ 5.1–5.17, 5.21–5.25, 5.53–5.57, 5.71–5.75, 5.95–5.97; Federal Aviation Administration. Safety Management System. June 17, 2026. URL: <https://www.faa.gov/faq-aviation-safety-outreach/safety-management-system> (visited on 08/27/2026).

care. It proves that authority, feedback, intervention, and records can form a mandatory operating system rather than an aspirational policy.

Aviation reporting supplies a related design lesson. Candid incident and near-miss information has prevention value, while a system that treats every report as a confession encourages silence. The Aviation Safety Reporting Program uses third-party reporting and a qualified, nonpunitive enforcement policy rather than blanket immunity.³⁵ The front-door record can follow the same incentive logic: protect good-faith reporting and evidence preservation while retaining consequences for falsification, concealment, retaliation, and unauthorized bypass.

2 The FTC Companion-Chatbot Orders

The Federal Trade Commission's 2025 section 6(b) study shows that AI-deployment records can be specified with considerably more precision than "produce the algorithm." The Commission ordered named firms to identify organizational roles, predeployment standards and thresholds, deployment approvers, considered and rejected mitigations, post-deployment monitoring, complaints, employee concerns, escalation, and response.³⁶

Those specifications map directly onto the front-door functions. Organizational roles identify the control lane. Thresholds and approvers define release authority. Rejected mitigations record a decision without deciding whether it was unreasonable. Monitoring, complaints, employee concerns, escalation, and response create a time-stamped feedback path. A sectoral statute can standardize the same categories prospectively and permit a court to

³⁵See Federal Aviation Administration. Aviation Safety Reporting Program. Apr. 2, 2021. URL: https://www.faa.gov/documentLibrary/media/Advisory_Circular/AC_00-46F.pdf (visited on 08/27/2026).

³⁶See Federal Trade Commission. FTC Launches Inquiry into AI Chatbots Acting as Companions. Sept. 11, 2025. URL: <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions> (visited on 08/27/2026); Federal Trade Commission. Model Order to File a Special Report Regarding AI Companion Chatbots. Sept. 10, 2025. URL: [https://www.ftc.gov/system/files/ftc_gov/pdf/AICompanionChatbot6\(b\)Order.pdf](https://www.ftc.gov/system/files/ftc_gov/pdf/AICompanionChatbot6(b)Order.pdf) (visited on 08/28/2026), specifications 3, 11–17. The study orders gather information; they are not findings that a recipient violated law.

request only the subset connected to a verified risk.

NIST’s AI Risk Management Framework can supply terminology for lifecycle roles, executive responsibility, third-party components, testing, feedback, monitoring, and decommissioning.³⁷ Its voluntary framework remains a vocabulary source. FAA rules and a sector-specific statute supply legal obligation; the FTC orders demonstrate that the responsive information can be operationally defined.

Together, the comparators answer the feasibility objection at the right level. The proposal asks firms already capable of version control, release management, access logging, vendor management, incident response, and legal retention to connect those systems for designated deployments. It does not require a universal explanation of every internal model state.

E Accurate Records Require Protected Escalation

A reliable record duty depends on candid reporting by the people closest to a risk. The infrastructure should therefore include protected reporting and a defined response path. A worker who, in good faith and on an objectively reasonable basis, documents a designated risk, uses the authorized escalation or intervention channel, and preserves evidence should receive protection against retaliation and the defense or indemnification permitted by governing law. The protection attaches to the report and authorized conduct, not to every act by the reporter.

The exceptions preserve record integrity. Intentional or reckless infliction of harm, knowing falsification or destruction, concealment of a material incident, retaliation against another reporter, corruption of an evaluation, and unauthorized removal of a safety constraint remain outside. The charter should identify the decisionmaker and review process before

³⁷See National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework (AI RMF 1.0). Jan. 2023. URL: <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf> (visited on 08/27/2026).

senior management becomes implicated.

This allocation aligns information with authority. A safety engineer may understand a failure mode without controlling budget or release. A product manager may control schedule without receiving the technical signal. An operating professional may see an abnormal output without access to the model's evaluation history. The record shows what moved between them, when it moved, and who could act. Protecting a documented escalation makes the record more reliable; preserving recourse for suppression makes it consequential.

Existing compliance and whistleblower regimes confirm the institutional principle while supplying different scopes and remedies. The Organizational Sentencing Guidelines treat leadership responsibility, adequate authority and resources, reporting without fear of retaliation, enforcement, response, and periodic risk assessment as components of an effective program.³⁸ Sectoral statutes protect specified reports and refusals under defined causation and remedy standards.³⁹ A front-door enactment should state its own scope rather than import every privilege or immunity from those regimes.

F The Record Is Also a Defense Asset

The front door is an information rule, and information can defeat a claim. A version record can show that the challenged feature was inactive. A configuration log can show that an independent deployer supplied the tool or removed a safeguard. An incident sequence can show that a warning appeared, an intervention fired, or a user edited the relevant content. Evaluation and approval records can show that the enterprise tested the asserted risk, adopted a response supported by the information then available, and monitored operation. A stop record can show that an actor lacked the practical authority the claimant alleges.

³⁸See U.S. Sent'g Guidelines Manual § 8B2.1, U.S. Sentencing Commission. Guidelines Manual, Section 8B2.1: Effective Compliance and Ethics Program. 2025. URL: <https://www.ussc.gov/guidelines/2025-guidelines-manual/annotated-2025-chapter-8> (visited on 08/27/2026).

³⁹See, e.g., Energy Reorganization Act Employee Protection Provision, 42 U.S.C. § 5851 (2024), § 5851(b)(3); Whirlpool Corp. v. Marshall, 445 U.S. 1, 9–13 (U.S. Feb. 26, 1980).

This defense value is especially important for stochastic systems. “The model can be unpredictable” states a general property. A record can show whether the particular event fell within an evaluated range, which controls were active, which feedback arrived, and whether the proposed alternative would have changed the sequence. It permits a defendant to establish reasonable preparation, intervening modification, user circumvention, or unavoidable residual risk with evidence rather than a category label.

The record also disciplines historical claims. Research published after an event may identify a variable worth testing without proving what a defendant knew earlier. A mitigation developed later may show a technically describable control without establishing that it was feasible, required, or admissible for the earlier deployment.⁴⁰ Versioned evaluations, resource decisions, complaints, and contemporaneous alternatives locate the inquiry at the correct time.

The infrastructure therefore serves both sides of the core proposition. It prevents a claimant from losing because the relevant deployment was invisible. It prevents an enterprise from losing because a court must infer a deployment from a frightening output alone. The front door makes the record reachable; the record allows ordinary law to distinguish the two.

V The Merits Firewall and Administrable Boundaries

A front-door rule succeeds because its limiting principles are part of the rule. Ordinary law retains the merits. A closed, staged record binds information cost. Existing legal interests define who may invoke access. These allocations yield a sectoral information rule with a risk-specific control lane, protected access, and prompt termination for enterprises outside that lane.

This Part develops those affirmative boundaries and then tests them against frag-

⁴⁰See Fed. R. Evid. 407.

mented systems, user and third-party intervention, arbitration, cross-border suppliers, and speech. The front door makes facts available; the governing law decides what those facts mean.

A Ordinary Law Receives Better Factual Predicates

Courts have long handled aircraft, pharmaceuticals, industrial equipment, distributed supply chains, proprietary software, and corporate records. Pleading on information and belief, early discovery, party requests, subpoenas, expert testing, circumstantial proof, amendment, and sanctions can address severe information asymmetry. AI receives no bespoke merits regime simply because its technology is complex.

The proposal accepts that premise. It creates no general AI tort, standard of care, product classification, causation presumption, or damages rule. Abraham and Sharkey's account of objective reasonableness explains how ordinary negligence can evaluate many AI disputes without reconstructing every internal model state.⁴¹ The cascade in Part I identifies observable deployment conduct precisely because the merits can remain ordinary.

The distinction lies in procedural prerequisites. Rule 8 and Rule 11 can accommodate incomplete evidence, but the complaint still requires a legal claim and enough facts to make it plausible. Rule 26 begins after a complaint has identified a litigable target. Rule 34 reaches a party's possession, custody, or control; Rule 45 reaches a known and served nonparty. Early discovery requires a recipient and judicial authorization. *Strike 3* shows how well the path can work when a claimant has a unique incident identifier and a known intermediary with the identity record. A divided AI deployment can withhold one of those predicates. Doe pleading cannot serve an unknown entity, relation back depends on notice and mistake, and Rule 27 preserves known testimony. Those tools become powerful after actor, version, and record holder are visible.

⁴¹See Abraham and Sharkey, *supra* note 2.

The front door fills that anterior stage. The operator already knows which system it incorporated into a consequential decision or made available to the claimant. The ex ante record duty requires provenance and risk control to remain knowable. The first answer passes that bounded information into ordinary procedure. From then on, the Federal Rules and the substantive cause of action perform their familiar work.

The rule also supplies a function outside discovery: creation. The ex ante duty generates release approval and version history while the deployment is operating. Litigation preservation begins when a dispute becomes reasonably foreseeable and protects that earlier record. Sectoral recordkeeping therefore makes consequential activity observable before the incentive to reconstruct it appears.

That is a familiar regulatory move. Aviation safety, financial transactions, employment records, medical care, and public administration use ex ante records because later adjudication and oversight require them. The justification comes from the protected sector and decision, not from a claim that AI is uniquely inscrutable. A designated deployment receives a front door when the legislature concludes that its consequences and divided information warrant one.

Existing attribution doctrine remains fully available. A corporation can be directly liable for its own design, supervision, warning, monitoring, or deployment choices; agency law can attribute conduct; products law can operate without a culpable engineer; electronic-transactions law can give effect to authorized automated acts.⁴² In *Mobley*, the court used a statutory agency theory at pleading rather than waiting for a legislative AI fix. The front door is valuable because it allows claimants and defendants to identify the facts on which such ordinary theories succeed or fail.

⁴²See American Law Institute, *supra* note 3, at §§ 7.03(1), 7.05, 7.07; American Law Institute, *supra* note 33, at §§ 1–2; Uniform Law Commission. Uniform Electronic Transactions Act. 1999. URL: <https://www.uniformlaws.org/committees/community-home?CommunityKey=2c04b76c-2b7d-4399-977e-d5876ba7e034> (visited on 08/27/2026).

B Staged Access Binds Information Cost

A firm may prevail on duty, defect, causation, damages, immunity, or another defense after paying for record design, retention, internal investigation, privilege review, trade-secret protection, third-party notice, and access. Multiple enterprises can duplicate work. The rule therefore assigns expense where standardization lowers total cost and screens requests before deeper production.

Geistfeld's analysis of insurability and overlapping protection warns against liability rules that raise prices for the same people they aim to protect.⁴³ The front door avoids the largest version of that problem by preserving the substantive class of compensable losses and ultimate burdens. Its remaining information cost should be defended directly: a limited record is the institutional price of making a designated high-impact deployment observable.

The price is controlled at six points. First, the legislature designates sectors and protected interests; capability or marketing alone never activates the regime. Second, the claimant identifies an existing claim or equitable entitlement and pleads concrete, risk-specific facts. Third, Stage One is limited to identity, versions, control allocation, incident events, preservation, and governing policies. Fourth, Stage Two follows when the initial record supports claim-specific need and is bounded by time, version, population, feature, and risk. Fifth, source code, weights, sensitive data, and security controls require necessity and consideration of less intrusive proof. Sixth, an enterprise without material control completes the first answer and terminates its special role.

Cost allocation can reinforce the screen. The operator bears the ordinary expense of maintaining and retrieving the minimum record because observability is a condition of covered deployment. A requester bears the cost of extraordinary formatting, duplicative demands, or an expert protocol when the tribunal finds the expense disproportionate to the defined interest. Expenses caused by an incomplete or evasive answer shift to the nonper-

⁴³See Geistfeld, *supra* note 2.

forming enterprise. Illinois Rule 224 offers a useful identity-stage analogue by limiting the proceeding to identification and assigning reasonable compliance costs to the petitioner; a sectoral statute can modify that baseline where the operator's ex ante duty makes ordinary retrieval part of compliance.⁴⁴

The ex ante record can also lower total cost. A deployment ID, defined owners, version lineage, event schema, and retention rule reduce the internal investigation needed after an incident. A supplier that provides current provenance and safety contacts avoids repeated emergency searches by each customer. A record that establishes reasonable evaluation, active safeguards, or independent modification can end a claim earlier. The rule converts bespoke forensic reconstruction into standardized retrieval.

Trade secrets require protection rather than categorical exclusion. Rule 26(c)(1)(G) already authorizes orders protecting trade secrets and confidential research, development, or commercial information. Secure expert review, redaction, in camera inspection, access limits, and testing protocols can resolve many disputes without public disclosure. The initial stage seldom requires source code. It asks which version ran, who controlled it, what validation governed, what event occurred, and where the deeper record resides.

Privilege presents a related boundary. Legal advice and attorney work product remain protected. The operator should identify withheld material in a privilege log and produce the underlying nonprivileged facts and business records. A company cannot transform an operational testing program into an inaccessible record by routing every evaluation through counsel, but the precise line follows existing privilege and work-product law. The model statute supplies no new exception.

Privacy and security receive the same design treatment. Requests should use event identifiers, defined windows, pseudonymization, aggregation, and data minimization. Unrelated users' conversations and records fall outside. Security-sensitive controls may be

⁴⁴See Ill. Sup. Ct. R., Ill. Sup. Ct. R. § 224.

described, tested in a secure environment, or reviewed in camera. These protections can narrow access while preserving the fact that a control, version, evaluation, or incident exists.

The rule therefore rejects two extremes. Public release of every technical artifact would impose unjustified costs and expose protected interests. Complete secrecy would permit a consequential system to define the evidence available against itself. Staged judicial access occupies the administrable middle.

C Existing Legal Interests Define the Gate

Henderson emphasizes the heterogeneity of AI activity and argues that some AI-caused losses should receive no tort remedy.⁴⁵ A disappointed user, an offended reader, a competitor facing lawful innovation, and a patient denied essential care present different interests. Treating every adverse output as compensable would flatten those distinctions.

The front door preserves them at the trigger. A claimant identifies a recognized cause of action or existing equitable entitlement. The verified facts connect the designated deployment and requested record to an element, defense, or remedy. A preference for a different model, a demand for a human conversation, or an assertion that an output felt wrong supplies no information right by itself.

Courts can decide pure legal defenses before expanding production. Jurisdiction, immunity, preemption, a categorical no-duty rule, absence of a private right, limitations apparent from the notice, and failure to identify a protected interest may end the matter. Stage One remains appropriate only to the extent identity or preservation is needed to determine the issue. The statute creates an enforceable right to a first answer for a qualifying notice; it does not create a right to audit any business that uses AI.

The minimum-access principle follows from that gate. A qualifying request receives the information necessary to locate the dispute within an existing legal rule and identify

⁴⁵See Henderson, *supra* note 2.

the actors capable of answering it. A claimant invoking defamation connects the request to publication, attribution, fault, or another required issue. An applicant invoking discrimination identifies an adverse process and protected theory. A benefits recipient identifies the entitlement and challenged action. Once those threshold facts place the interest outside governing law, the proceeding ends.

Sectoral designation supplies a second screen. Legislatures can select deployments where affected people predictably lack the facts needed to invoke existing protection. They can omit low-consequence uses, define small-entity procedures, and tailor records to the interest. The resulting statute makes no claim of universal AI exceptionalism.

The access duty is forum neutral and nonwaivable as to the affected person. A vendor agreement can allocate performance and expense between enterprises, but it cannot erase the operator's minimum record or first answer. Confidentiality yields to the defined access right subject to privilege, privacy, trade-secret, and security protection.

Arbitration can select the merits forum when an agreement is otherwise enforceable. The record obligation still arises from deployment and a qualifying request before the private forum needs to compel discovery. That distinction matters because arbitral power over unbound nonparties can be narrower than party discovery. Federal Arbitration Act section 7 authorizes arbitrators to summon a person to attend and bring documents; the Ninth Circuit has held that it does not authorize a stand-alone prehearing document subpoena to a nonparty.⁴⁶ The statute should therefore require the bound visible operator to maintain the closed local record and contractual retrieval path. An arbitrator can enforce the duty against a bound operator; the design does not depend on creating new compulsory power over an upstream stranger.

Cross-border supply receives the same treatment. A foreign enterprise that operates

⁴⁶See Federal Arbitration Act, 9 U.S.C. § 7 (2024); *CVS Health Corp. v. Vividus, LLC*, 878 F.3d 703, 706–09 (9th Cir. 2017).

a covered domestic service remains subject to ordinary jurisdiction, service, comity, and choice-of-law limits. A domestic operator, however, can condition deployment on local retention of the minimum manifest, a reachable supplier agent, material-change notice, and a right to retrieve defined event and provenance fields. The front door solves the information architecture at contracting time instead of implying that Rule 45 reaches every foreign supplier. If a supplier refuses the required interface, the operator can choose another supplier or forgo the designated deployment.

D Fragmentation Narrows Control Rather Than Eliminating It

The material-control test is designed for heterogeneous stacks. In a vertically integrated service, one enterprise may own the application, configure the harness, monitor incidents, and issue updates. In an API deployment, the model provider may control base capability and access conditions while the application company controls prompts, retrieval, tools, user population, warnings, and business objectives. In an open-weights deployment, the original publisher may lose continuing visibility and stopping power while a modifier adds the purpose, data, credentials, and interface.

These differences alter who answers after Stage One. They do not defeat the gateway. The visible operator maintains provenance for the system it chose. It identifies upstream components and the rights reserved by each enterprise. The cumulative control test then asks whether a particular lane was risk specific, practically exercisable, and material.

Several recurring defenses follow naturally. A downstream actor that removed safeguards, changed weights, or added dangerous tools may control the relevant path. A user who deliberately defeated a tested constraint may bear comparative or primary responsibility under governing law. A novel cyberattack may interrupt causation, while an unpatched known vulnerability remains part of deployment risk. An upstream provider may show that it could neither observe nor prevent an independent use. A provider that reserved audit, update, suspension, or safety-override rights may have retained a lane. Version and authority records allow the parties to prove these propositions.

Contract labels remain evidence rather than conclusions. An agreement can establish usable rights and information duties. It cannot allocate away a statutory front door owed to an affected person, and the word “independent” cannot erase a control the contract actually reserves. Conversely, brand ownership and a general right to terminate a commercial relationship do not establish control over every deployment risk.

The framework leaves room for residual risk. Some harmful outputs will remain unforeseeable after reasonable testing, monitoring, warning, and response. Some interventions will be technically unavailable or would not have changed the event. Ordinary fault and causation standards can produce judgment for the enterprise. A complete deployment record strengthens that result by showing which risks were assessed, which controls were active, what feedback arrived, and why the proposed alternative fails.

E Product and Speech Classification Remain for the Merits

AI disputes can implicate design, content, and human expression in different combinations. The front-door rule should not resolve that classification at coverage. A companion application may present product-design questions about age access, memory, engagement objectives, crisis detection, and warnings while a claim imposing liability for an adopted idea raises a distinct First Amendment issue. An employment score may be part of a decision process without becoming a physical product. A generated notice may communicate the operator’s decision even when the model has no legal status of its own.

Garcia illustrates the separation. At Rule 12, the district court held that the complaint plausibly treated Character.AI as a product because the challenged features concerned design and commercial distribution. It declined on that record to treat the generated output as protected speech where defendants had not identified the relevant expressive choice or speaker.⁴⁷ Those rulings were procedural and claim specific. They neither classified every

⁴⁷See *Garcia*, 785 F. Supp. 3d at 1174–80.

language model as a product nor denied that human expression mediated through AI can receive protection on another record.

The front door asks for facts that permit the classification inquiry: what conduct the claim challenges; which human or organizational choices the system implemented; what message, if any, an identified person adopted; and whether the requested remedy regulates design, transaction, decision process, or expression. Production remains subject to constitutional limits and ordinary law. If a claim targets protected speech, the court applies the appropriate substantive and procedural safeguards. If it targets a nonexpressive control, the presence of generated words does not automatically remove the design from review.

Anthropomorphic presentation has the same evidentiary status. Calling a system a companion, therapist, or agent cannot create a new legal person. Memory, proactive messages, affection, expertise claims, and relational objectives can nevertheless bear on intended use, foreseeable exposure, warning, and which evaluations are risk specific. The law evaluates the enterprise's own choices without treating the model as speaker, tortfeasor, or liability sink by default.

The boundary completes the proposal. The front door is classification neutral but fact enabling. It creates no preference for liability and no immunity through metaphor. It permits courts to choose the correct substantive category on a record that reveals the deployment rather than on a label selected before the relevant facts were available.

Conclusion

AI litigation can fail before a court reaches duty, defect, discrimination, causation, immunity, or defense. The affected person sees a result; the identity, version, configuration, evaluation, incident, and stop records remain divided among organizations. Ordinary procedure becomes powerful once the proper target and record holder are known. It cannot guarantee those predicates or create a deployment record after the event.

The deployment front door fills that gap. For a legislatively designated high-impact

deployment, a claimant with a legally cognizable injury or existing equitable entitlement and a plausible, risk-specific nexus may serve the visible operator. The operator identifies the system and relevant entities, preserves the record, and supplies a staged first answer. Only an enterprise with risk-specific, practical, and material control remains in the control lane; an enterprise without it completes the answer and terminates its special role.

The sequence is concrete. Legislatures designate the deployment and protected interest. Operators create a versioned minimum record and retrieval path before use. An affected person serves a verified, event-specific request on the entity already visible. The first answer maps identity, provenance, custody, and intervention authority. A cumulative control test confines participation to risk-specific, practically exercisable, material lanes. Federal procedure then governs any federal action. Each step solves one information problem and stops where the next body of law begins.

The rule works because record infrastructure precedes litigation. A versioned responsibility charter assigns risk and release ownership, feedback, intervention, provenance, and retention. The cascade specifies the facts relevant to authorization, notice, feasibility, causal sequence, and rebuttal. Privacy, privilege, trade secrets, security, proportionality, and federal procedure govern access. Compliance establishes no underlying breach, defect, causation, damages, or joint liability.

That sequence preserves the most valuable feature of ordinary law: its ability to distinguish cases. *Garcia* shows why an integrated operator's internal record matters even when its name is obvious. *Mobley* shows how a vendor-hosted, customer-controlled information structure can fracture identification and production across a divided stack. The front door makes the relevant facts reachable in both settings; the cause of action still determines who wins.

The payoff is larger than disclosure. A claimant can name the operative actor before limitations expire. A court can protect trade secrets while testing the relevant control. A supplier with no event or authority can leave early. An operator can retrieve a defined record

instead of reconstructing an organization after harm. And a defendant can prove an inactive feature, independent modification, effective intervention, or unavailable alternative on the same record. Observability strengthens adjudication in both directions.

Deployment law is uniquely positioned to create that observability. It operates when provenance, versioning, authority, feedback, and retention can still be designed into the system. Procedure is uniquely positioned to protect and test the resulting record once a dispute begins. The front door connects those institutional strengths: *ex ante* law keeps the facts knowable, and adjudication assigns their legal meaning after the parties can see them.

That division remains durable as models, vendors, and deployment architectures change. Sectoral lawmakers can revise the minimum fields and retention clock without reopening the merits standard. Courts can apply familiar pleading, discovery, evidence, and liability rules to a record that identifies the actual system before them.

Probabilistic output does not eliminate legal responsibility; it relocates the evidence of control. The deployment front door makes that evidence reachable before information asymmetry decides the case.

Authorities and Sources

K. S. Abraham and C. M. Sharkey, Untangling AI Liability, 115 Calif. L. Rev. (forthcoming 2027).

American Law Institute, *Restatement (Third) of Torts: Products Liability* (American Law Institute Publishers, 1998).

American Law Institute, *Restatement (Third) of Agency* (American Law Institute Publishers, 2006).

American Law Institute, *Restatement (Third) of Torts: Liability for Physical and Emotional Harm* (American Law Institute Publishers, 2010).

Ash v. Cort, 512 F.2d 909 (3d Cir. 1975).

Ashcroft v. Iqbal, 556 U.S. 662 (U.S. 2009).

J. Banks, *Deletion, Departure, Death: Experiences of AI Companion Loss*, 41 J. Soc. & Pers. Relationships 3547 (2024).

Bell Atlantic Corp. v. Twombly, 550 U.S. 544 (U.S. 2007).

Burlington Northern Railroad Co. v. Woods, 480 U.S. 1 (U.S. 1987).

Butler v. National Community Renaissance, 766 F.3d 1191 (9th Cir. 2014).

California Code of Civil Procedure, Cal. Civ. Proc. Code § 2035.010 (2026).

Columbia Insurance Co. v. Seescandy.com, 185 F.R.D. 573 (N.D. Cal. 1999).

Companion Chatbots, Cal. Bus. & Prof. Code §§ 22601–22606 (2025). 2025 Cal. Stat. ch. 677.

CVS Health Corp. v. Vividus, LLC, 878 F.3d 703 (9th Cir. 2017).

J. De Freitas et al. Lessons from an App Update at Replika AI: Identity Discontinuity in Human–AI Relationships. 25-018. May 21, 2025. URL: https://www.hbs.edu/ris/Publication%20Files/25-018_d256c104-e5ba-4911-b72c-874c5492166e.pdf (visited on 08/27/2026).

Energy Reorganization Act Employee Protection Provision, 42 U.S.C. § 5851 (2024).

Federal Arbitration Act, 9 U.S.C. § 7 (2024).

Federal Aviation Administration. Aviation Safety Reporting Program. Apr. 2, 2021. URL: https://www.faa.gov/documentLibrary/media/Advisory_Circular/AC_00-46F.pdf (visited on 08/27/2026).

Federal Aviation Administration. Safety Management System. June 17, 2026. URL: <https://www.faa.gov/faq-aviation-safety-outreach/safety-management-system> (visited on 08/27/2026).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 8(a)(2) (2025).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 11(b)(3) (2025).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 26 (2025).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 34(a)(1), 45 (2025).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 27 (2025).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 15(c) (2025).

Federal Rules of Civil Procedure, Fed. R. Civ. P. § 37(e) (2025).

Federal Trade Commission. FTC Launches Inquiry into AI Chatbots Acting as Companions. Sept. 11, 2025. URL: <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions> (visited on 08/27/2026).

Federal Trade Commission. Model Order to File a Special Report Regarding AI Companion Chatbots. Sept. 10, 2025. URL: [https://www.ftc.gov/system/files/ftc_gov/pdf/AICompanionChatbot6\(b\)Order.pdf](https://www.ftc.gov/system/files/ftc_gov/pdf/AICompanionChatbot6(b)Order.pdf) (visited on 08/28/2026).

Garcia v. Character Technologies, Inc. 785 F. Supp. 3d 1157 (M.D. Fla. May 21, 2025).

Garcia v. Character Technologies, Inc. No. 6:24-cv-01903-ACC-DCI, Dkt. 244 (M.D. Fla. Jan. 7, 2026). order dismissing settled action.

Garcia v. Character Technologies, Inc. No. 6:24-cv-01903-ACC-DCI, Dkt. 273 (M.D. Fla. June 11, 2026). order resolving charging lien.

M. A. Geistfeld, *Insurability and Liability for AI-Caused Harms*, 32 Ins. L. Rev. 17 (2026).

Hanna v. Plumer, 380 U.S. 460 (U.S. 1965).

Z. Henderson, *Rethinking Robot Liability*, 67 B.C. L. Rev. 479 (2026).

Houston Federation of Teachers, Local 2415 v. Houston Independent School District, 251 F. Supp. 3d 1168 (S.D. Tex. May 4, 2017).

Illinois Supreme Court Rules, Ill. Sup. Ct. R. § 224 (2026).

In re Citric Acid Litigation, 191 F.3d 1090 (9th Cir. 1999).

H. R. Kirk et al. Neural Steering Vectors Reveal Dose and Exposure-Dependent Impacts of Human–AI Relationships. 2512.01991v2. Feb. 18, 2026. URL: <https://arxiv.org/abs/2512.01991> (visited on 08/27/2026).

Krupski v. Costa Crociere S.p.A. 560 U.S. 538 (U.S. 2010).

L. Laestadius et al., *Too Human and Not Human Enough: A Grounded Theory Analysis of Mental Health Harms from Emotional Dependence on the Social Chatbot Replika*, 26 New Media & Soc’y 5923 (2024).

N. G. Leveson, *Engineering a Safer World: Systems Thinking Applied to Safety* (Cambridge, Mass.: MIT Press, 2012).

- N. G. Leveson and J. P. Thomas. STPA Handbook. Mar. 2018. URL: https://psas.scripts.mit.edu/home/get_file.php?name=STPA_handbook.pdf (visited on 08/27/2026).
- Mobley v. Workday, Inc. 740 F. Supp. 3d 796 (N.D. Cal. July 12, 2024).
- Mobley v. Workday, Inc. No. 3:23-cv-00770-RFL, Dkt. 128 (N.D. Cal. May 16, 2025). order granting preliminary collective certification.
- Mobley v. Workday, Inc. No. 3:23-cv-00770-RFL, Dkt. 158 (N.D. Cal. July 29, 2025). order concerning HiredScore features and customer identification.
- Mobley v. Workday, Inc. No. 3:23-cv-00770-RFL, Dkt. 340 (N.D. Cal. May 29, 2026). discovery order.
- Mobley v. Workday, Inc. No. 3:23-cv-00770-RFL, Dkt. 376 (N.D. Cal. July 2, 2026). order denying certification of an interlocutory appeal.
- National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework (AI RMF 1.0). Jan. 2023. URL: <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf> (visited on 08/27/2026).
- New York Civil Practice Law and Rules, N.Y. C.P.L.R. § 3102(c) (2026).
- S. Rismani et al., *From Plane Crashes to Algorithmic Harm: Applicability of Safety Engineering Frameworks for Responsible ML* CHI '23 Proc. 2:1–2:18 (Apr. 19, 2023).
- Rotella v. Wood, 528 U.S. 549 (U.S. 2000).
- Rules Enabling Act, 28 U.S.C. § 2072 (2024).
- Safety Management Systems, 14 C.F.R. pt. 5 (2025).
- Sergeeva v. Tripleton International Ltd. 834 F.3d 1194 (11th Cir. 2016).
- Speller v. Sears, Roebuck & Co. 100 N.Y.2d 38 (N.Y. May 6, 2003).
- Strike 3 Holdings, LLC v. Doe, 964 F.3d 1203 (D.C. Cir. 2020).
- Teamsters Local 404 Health Services & Insurance Plan v. King Pharmaceuticals, Inc. 906 F.3d 260 (2d Cir. 2018).

Texas Rules of Civil Procedure, Tex. R. Civ. P. § 202 (2026).

U.S. Sentencing Commission. Guidelines Manual, Section 8B2.1: Effective Compliance and Ethics Program. 2025. URL: <https://www.ussc.gov/guidelines/2025-guidelines-manual/annotated-2025-chapter-8> (visited on 08/27/2026).

Uniform Law Commission. Uniform Electronic Transactions Act. 1999. URL: <https://www.uniformlaws.org/committees/community-home?CommunityKey=2c04b76c-2b7d-4399-977e-d5876ba7e034> (visited on 08/27/2026).

United States v. United States Gypsum Co. 438 U.S. 422 (U.S. 1978).

Walker v. Armco Steel Corp. 446 U.S. 740 (U.S. 1980).

Whirlpool Corp. v. Marshall, 445 U.S. 1 (U.S. Feb. 26, 1980).

Winzer v. Kaufman County, 916 F.3d 464 (5th Cir. 2019).

Workman v. United States Postal Service, 127 F.4th 237 (10th Cir. 2025).

R. Zhang et al., *The Dark Side of AI Companionship: A Taxonomy of Harmful Algorithmic Behaviors in Human–AI Relationships* CHI '25 Proc. 1 (2025).